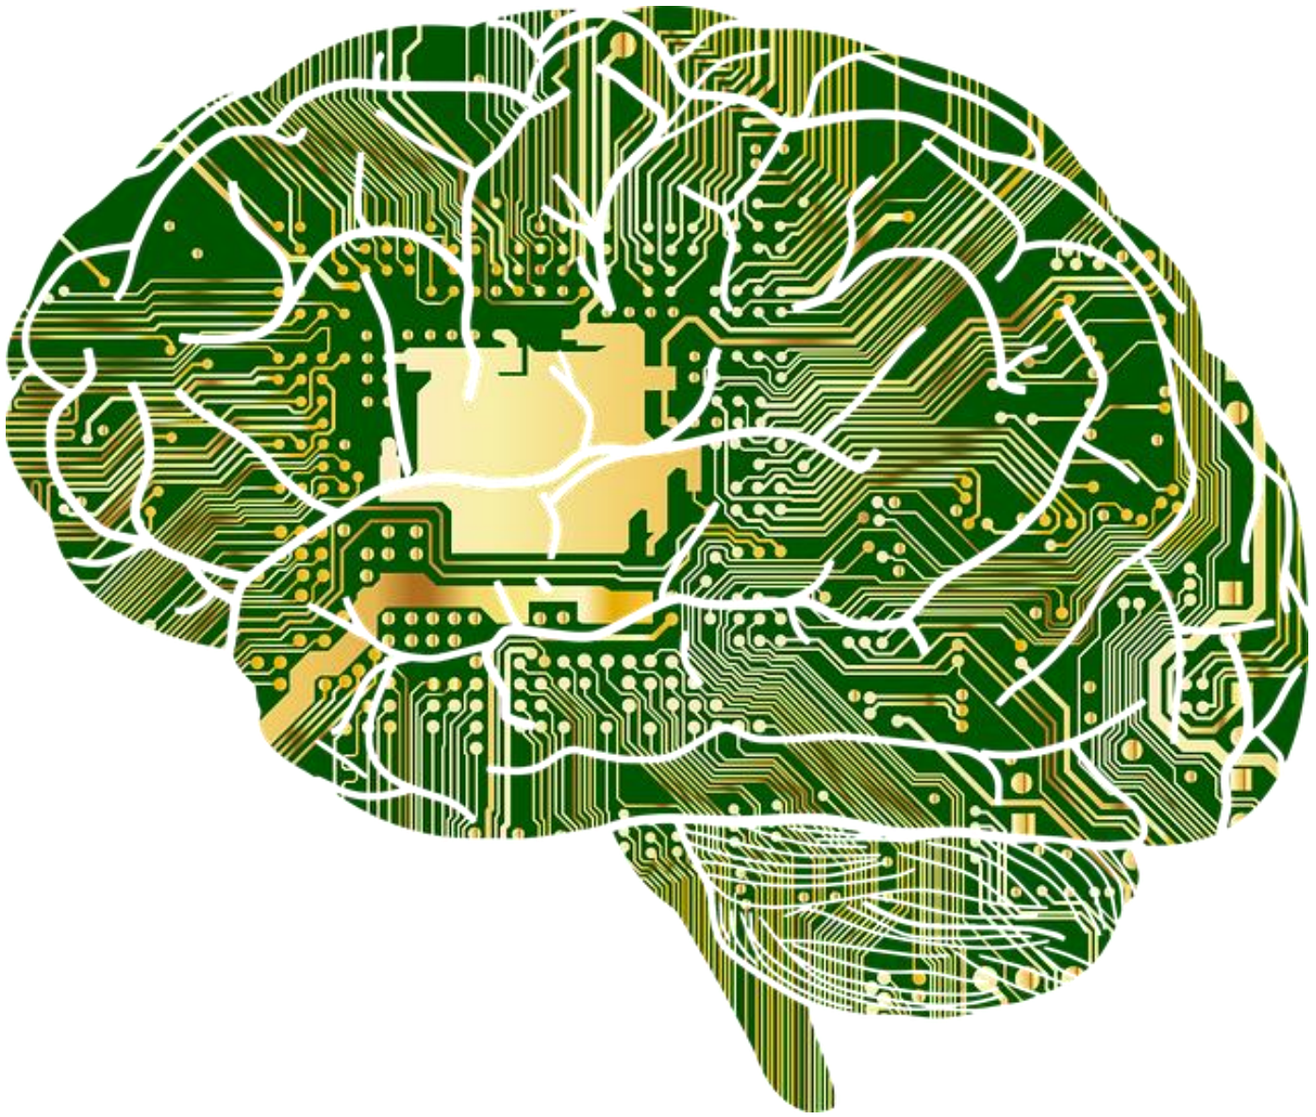

Deep learning

Introduction

Caspar van Leeuwen, David Ruhe

SURF SARA

SURFsara

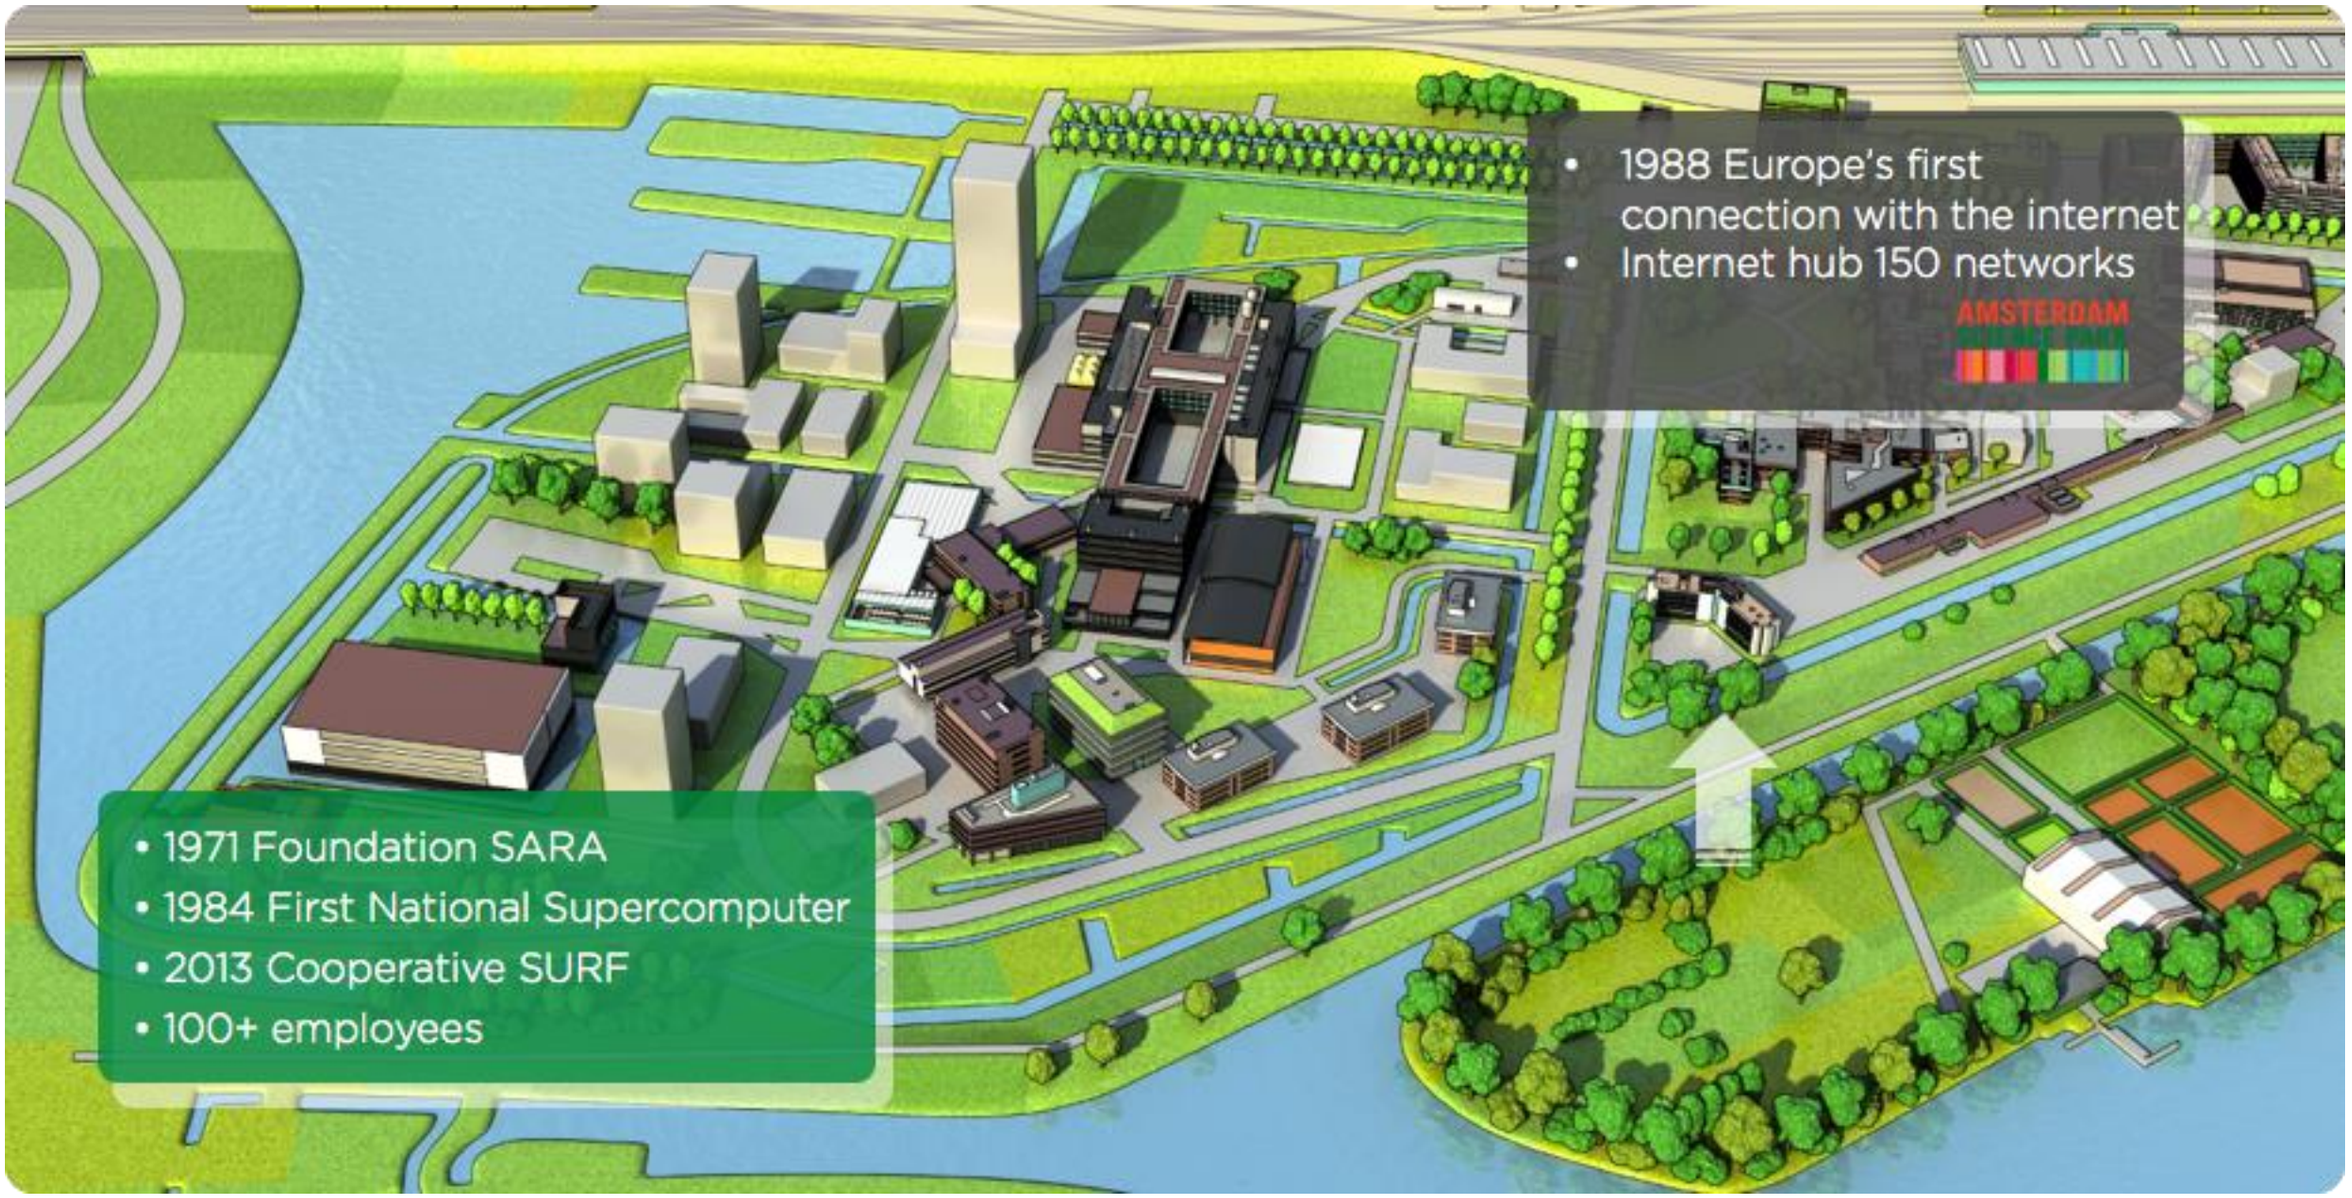
• 1988 Europe's first connection with the internet
• Internet hub 150 networks
AMSTERDAM
• 1971 Foundation SARA
• 1984 First National Supercomputer
• 2013 Cooperative SURF
• 100+ employees

SURF

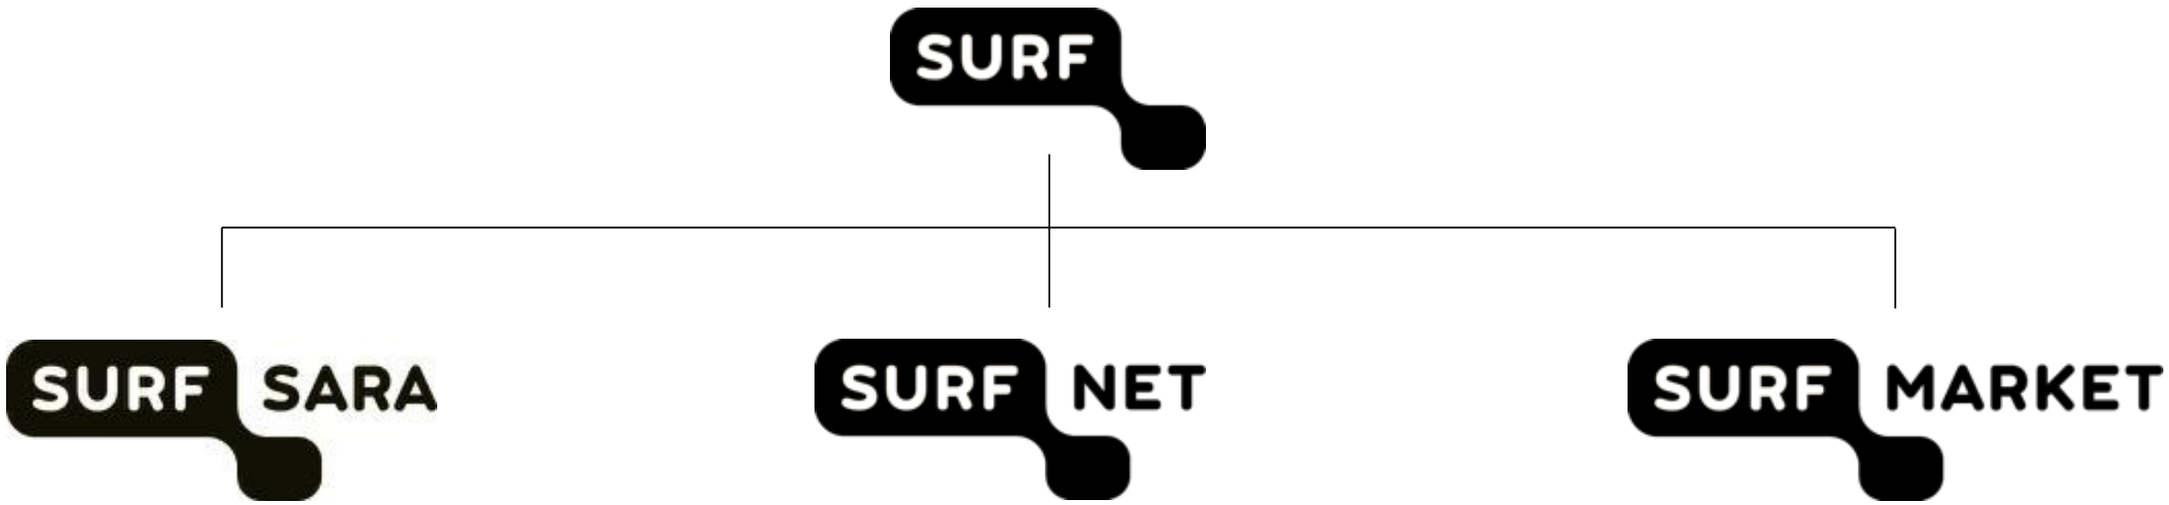

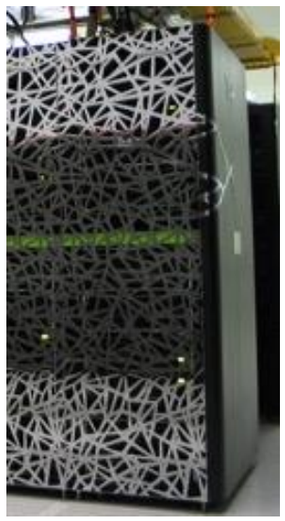

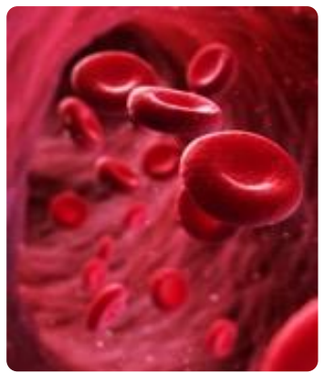

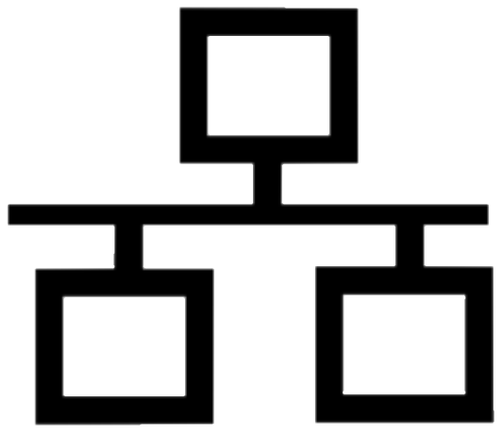

SURFsara

High performance computing	**Data processing**	**Data services**
Supercomputing	Data analytics	Mass online storage
Clustercomputing	Grid services	PID service
Machine learning	Visualization	Data management
HPC cloud	HPC cloud	Data preservation

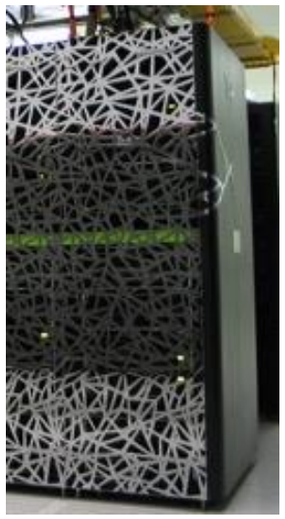

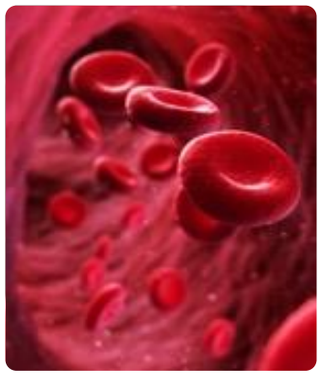

Who are we?

Caspar van Leeuwen

- Advisor Cluster/Supercomputing
- Support users in using Cartesius / Lisa
- Strong interest in Deep Learning
- Other activities, e.g. Training

David Ruhe

- Advisor Machine Learning
- Support ML workloads on HPC infrastructure
- Other activities, e.g. joint (ML) research projects

Course information

- The course will be “hands-on".
- We will use **R**, the programming language.
  - Disclaimer: R is *not* a common HPC language, but was thought to fit best with your backgrounds. We are sure we can manage, but are not fluent in R syntax!
- We will use **Keras**, a popular deep learning framework.
- Online Jupyter notebook environment - later today
- Syllabus, announcements & discussion board @ Canvas

Your backgrounds

Any experience in

- Programming
- R
- Statistics
- Machine learning

Course plan

1. **Introduction and ML & DL fundamentals**
2. **Optimisation & network training**
3. ML tasks (regression vs classification) & regularisation
4. Sequential data analysis with recurrent neural networks
5. Image processing with convolutional neural networks
6. Unsupervised learning
7. Complex architectures and models

Day 1

Day 2

Day 3

Today’s program

14:00-14:30 What is ML / DL? What is a neuron?

14:30-15:30 Hands on: building a neuron from scratch

15:30-15:45 How do neural networks work?

15:45-16:15 Break

16:15-16:30 How do neural networks work? (continued)

16:30-17:15 Hands on: building a network from scratch

17:15-18:00 Loss function & updating weights

18:00-19:00 Diner

19:00-19:45 Hands on: the XOR problem

19:45-20:00 Dataset splitting & Performance evaluation

20:00-21:00 Hands on: Keras on fashion Mnist

If we finish early: Machine learning tasks (classification vs regression)

Overview

Today we will cover

- Where deep learning has made an impact.
- The basic components of a neural network.
- Build a basic neural network from scratch
- Perform binary classification using a neural network
- Get to know the online Jupyter notebook environment.

Mach rning

```
DecisionTreeClassif epth 1 with 3 nodes
  If (DEPARTURE_TIM
   Predict: NOT_DEL
  Else (DEPARTURE_T
   Predict: DELAYED
```

P s

Machine le

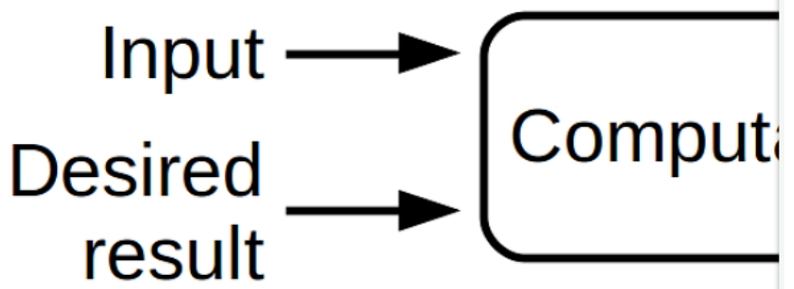

Features and result of ‘Decision’ system

(delivery, full-term).

Limited human input

Source: Machine Learning for an Expert Sys oolery et al. (1994)

Machine learning

- Learning tasks are often broken into two broad categories, **supervised** or **unsupervised**.
- In **supervised** learning we train using the expected output.
  - **Regression**
  - **Classification**.
- In **unsupervised** learning the expected output is not provided (or does not exist) during training. Instead we rely on the *similarity* of data points.
  - **Clustering**
  - **Dimensionality reduction**
- Machine learning has many good models but most require some manual work (feature engineering) to perform.

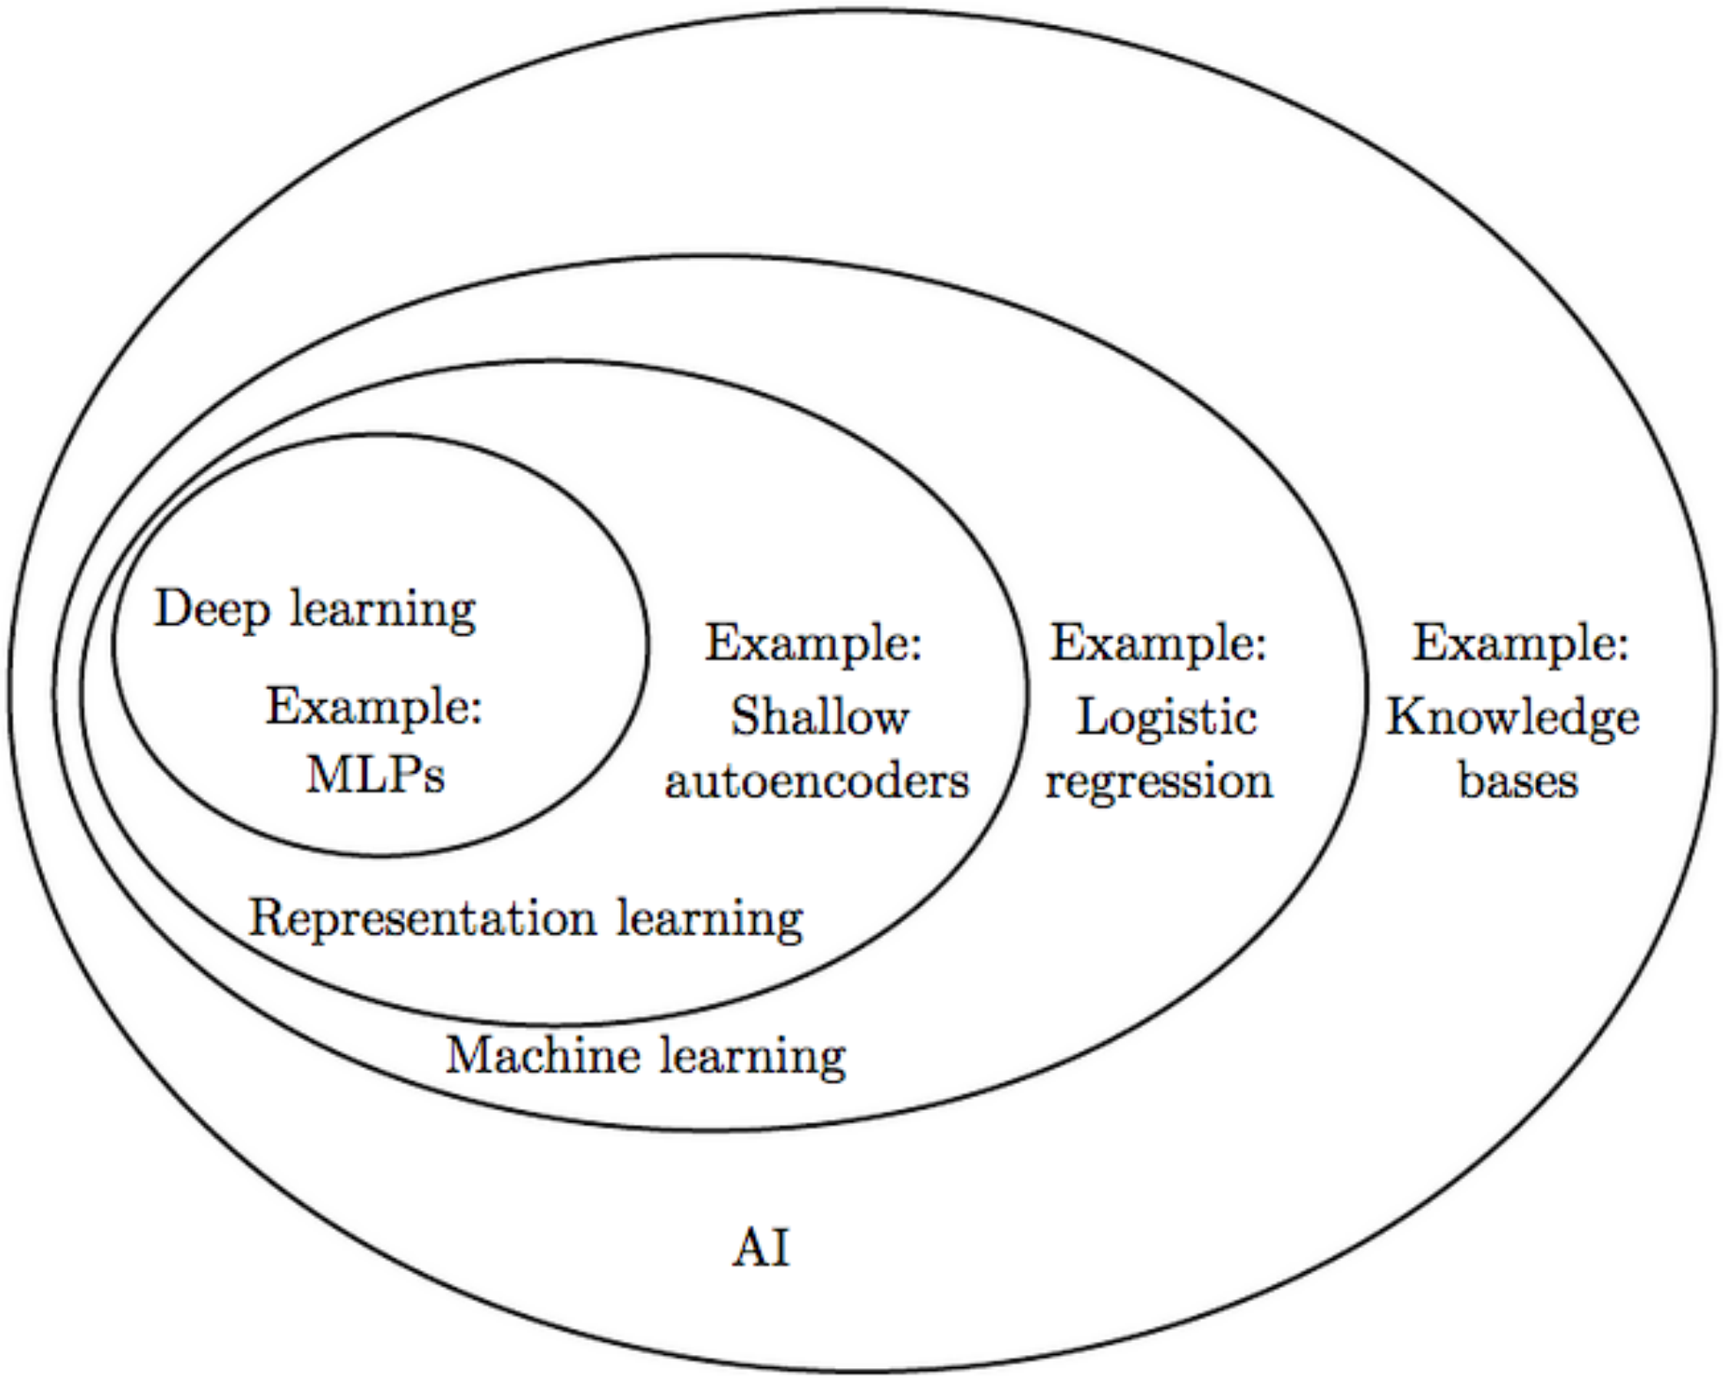

Source: [Deep Learning, Ian Goodfellow and Yoshua Bengio and Aaron Courville]

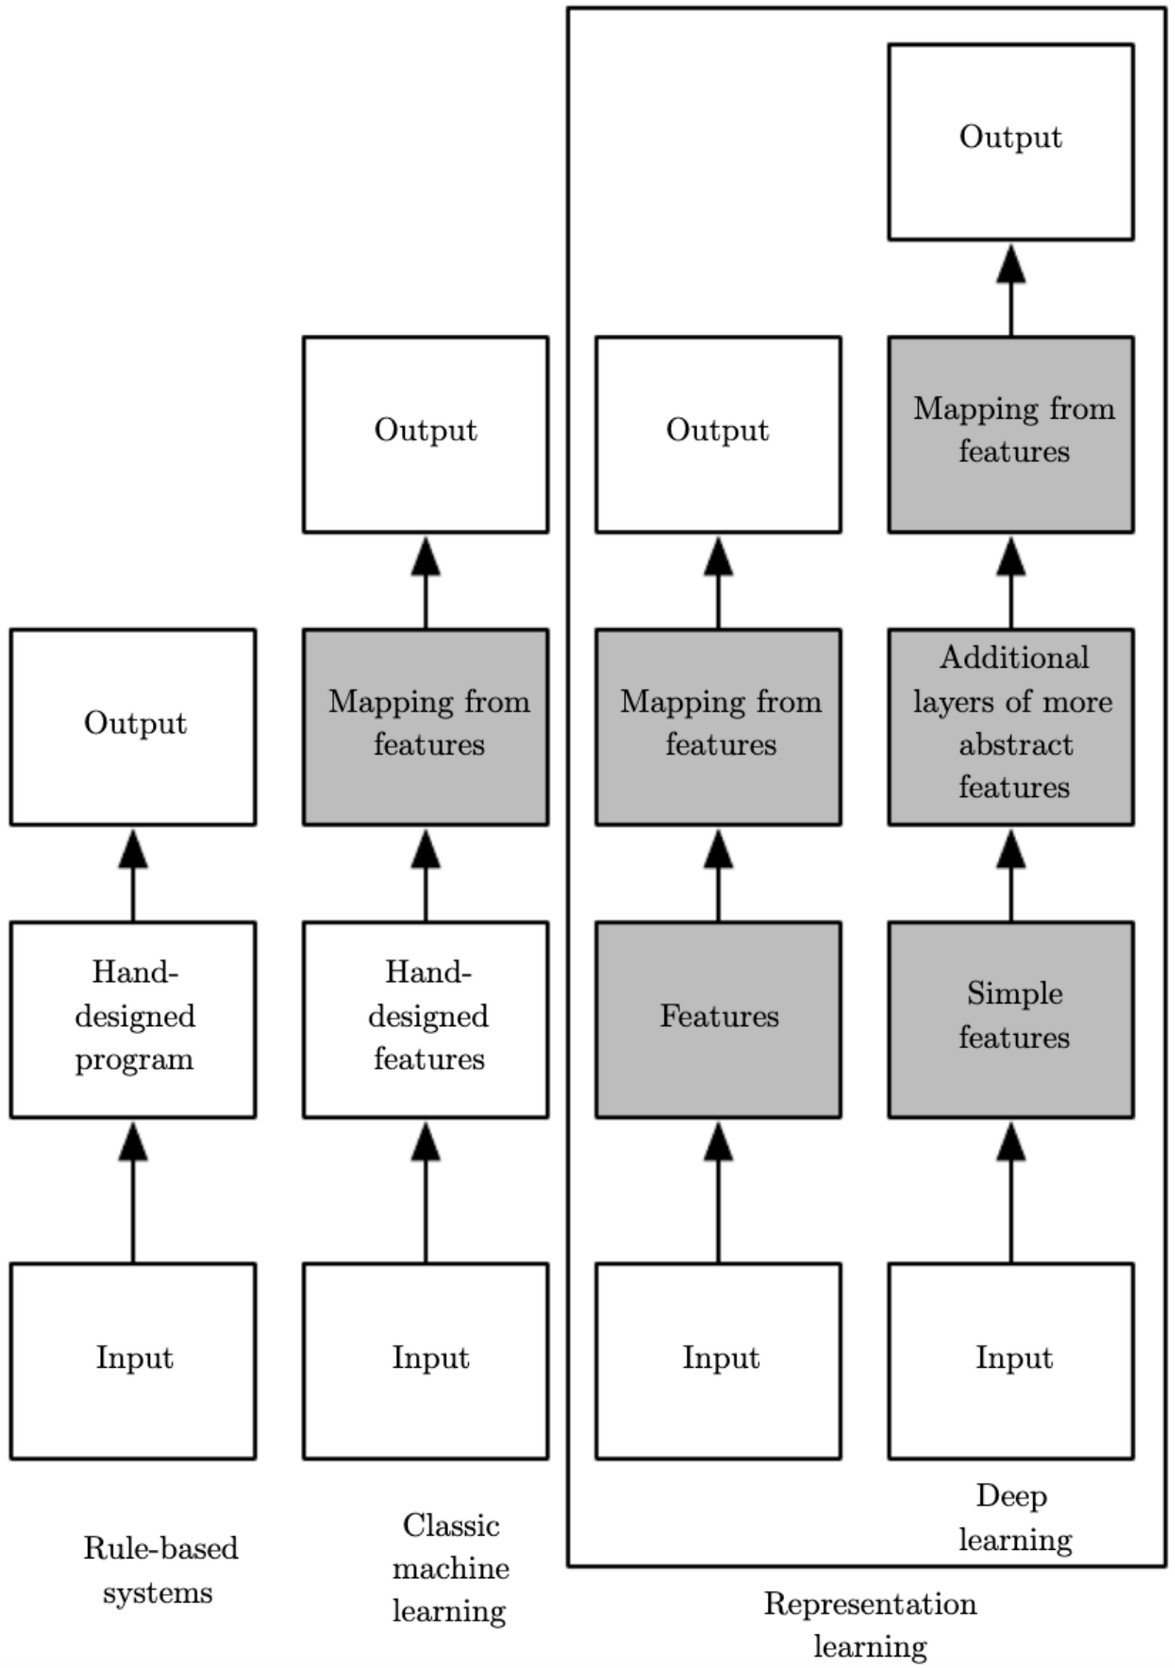

Source: [Deep Learning, Ian Goodfellow and Yoshua Bengio and Aaron Courville]

What is deep learning?

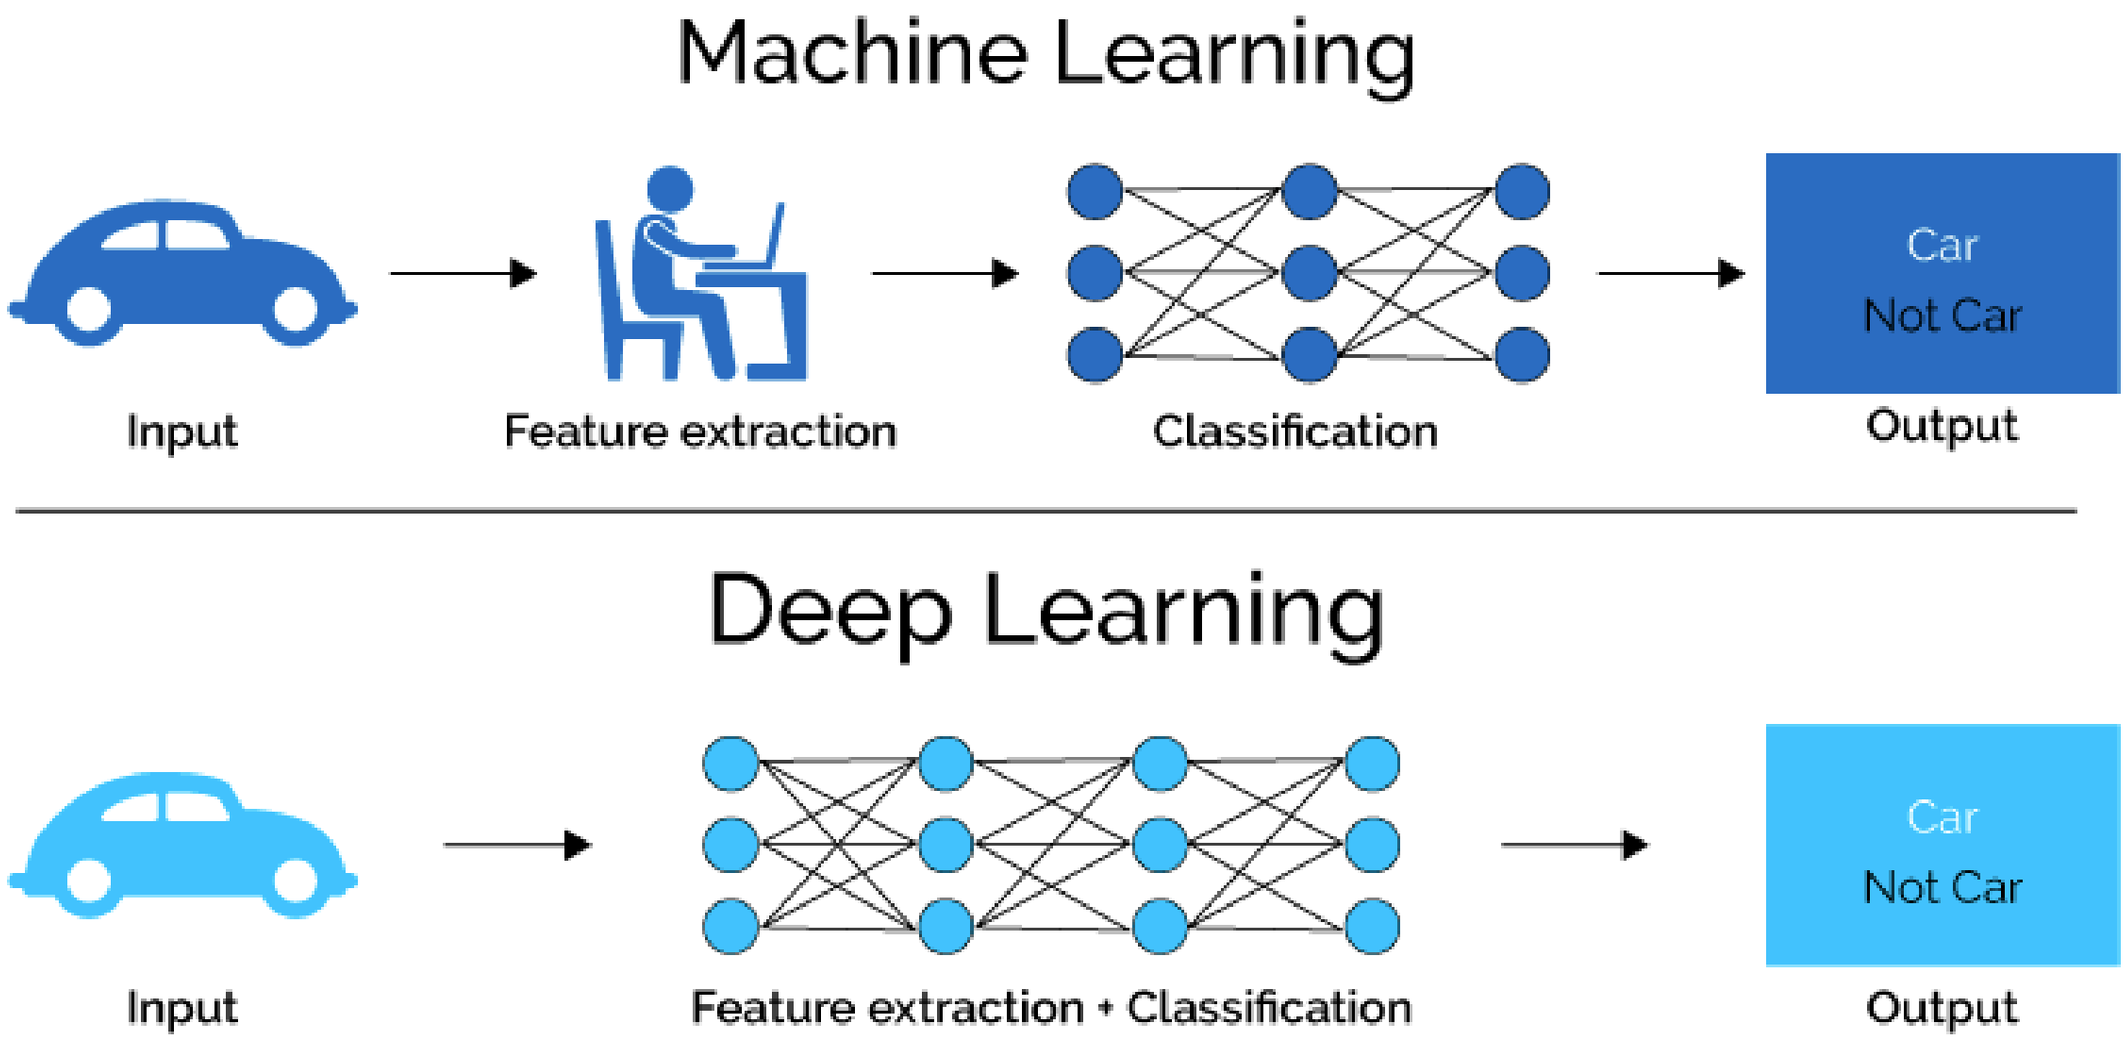

Source: Alex, D., 2017: Difference between Machine Learning and Deep Learning.

Success of deep learning

- Machine translation
- Speech recognition
- Language recognition based on characters
- Image classification
- Image screen caption generation
- Recommender systems
- Drug design
- ...

Success of deep learning

Style transfer

Source: https://ai.googleblog.com/2016/10/supercharging-style-transfer.html

Success of deep learning

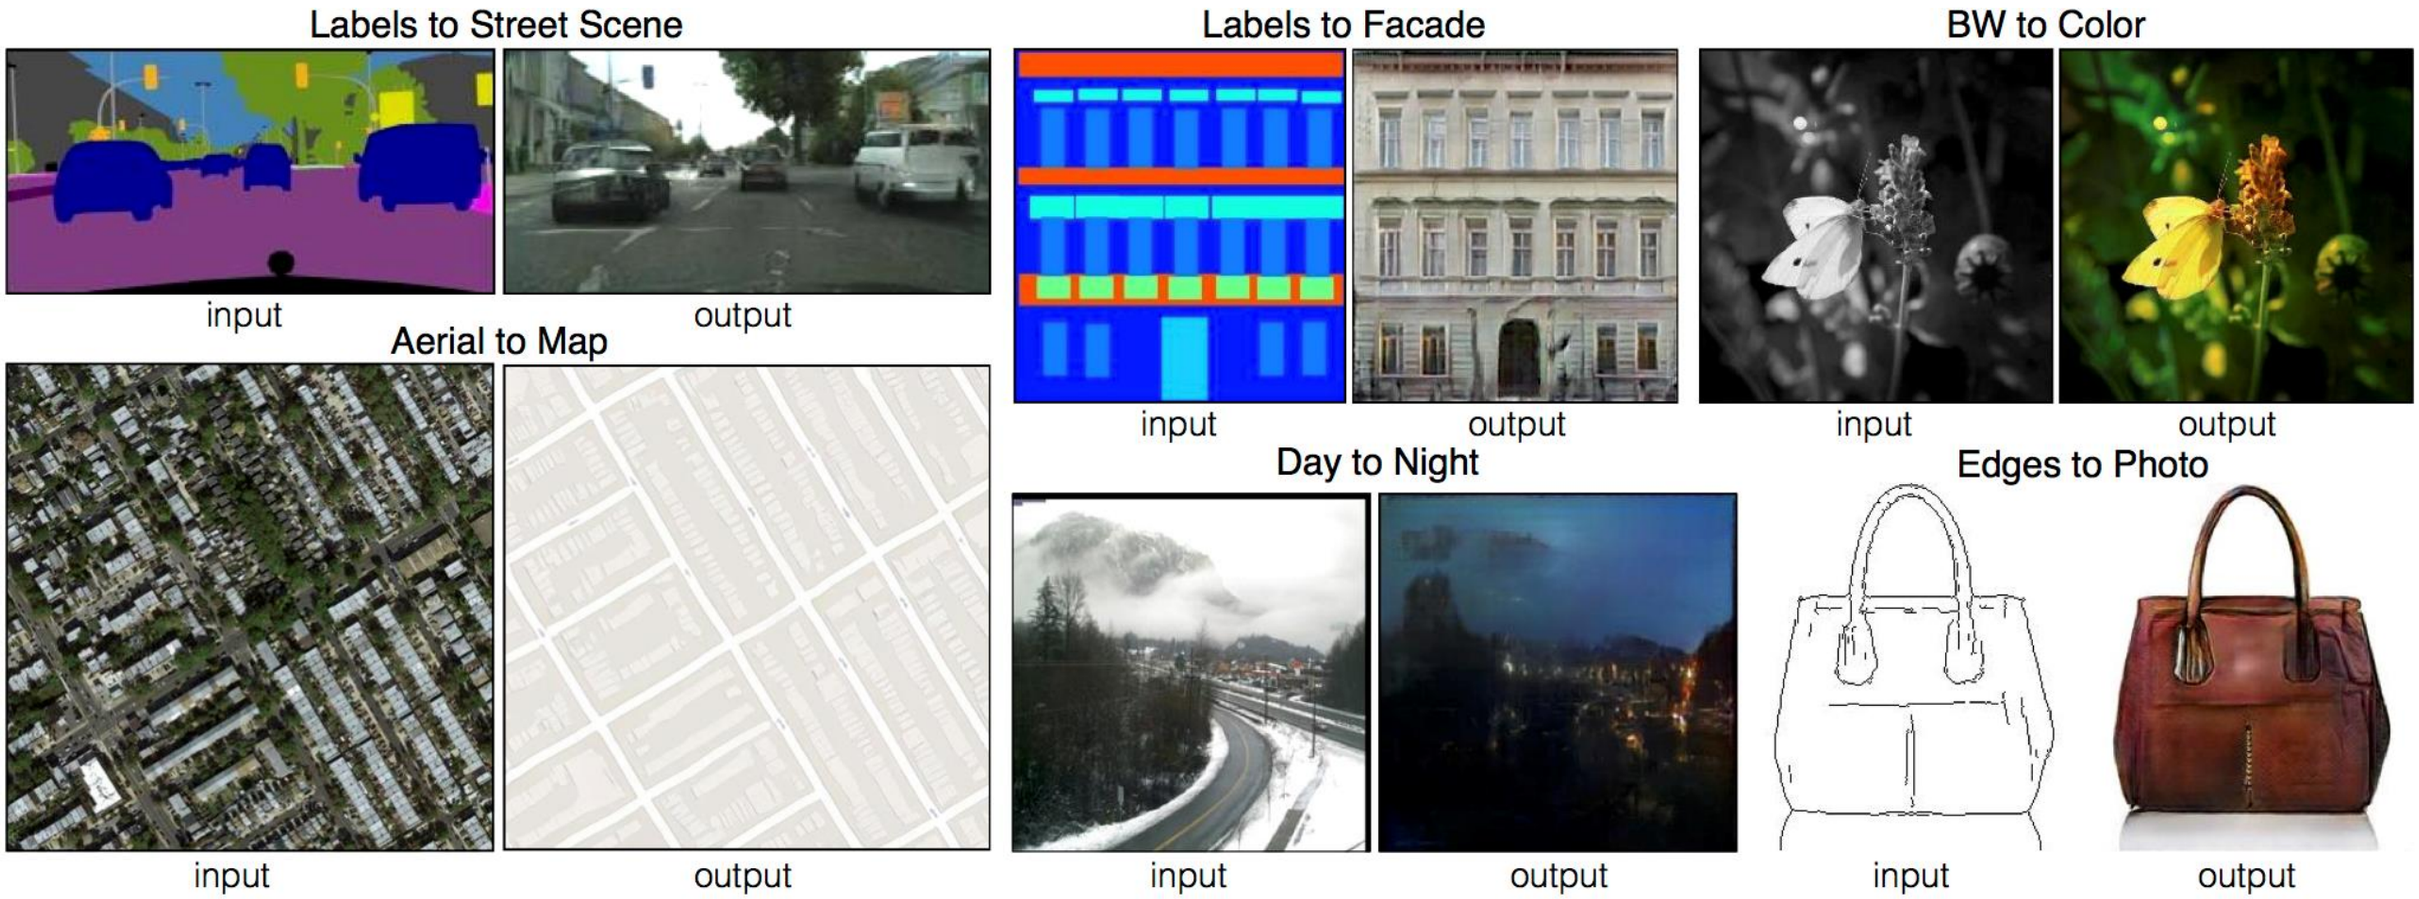

Image to image
Source: https://phillipi.github.io/pix2pix/

Success of deep learning

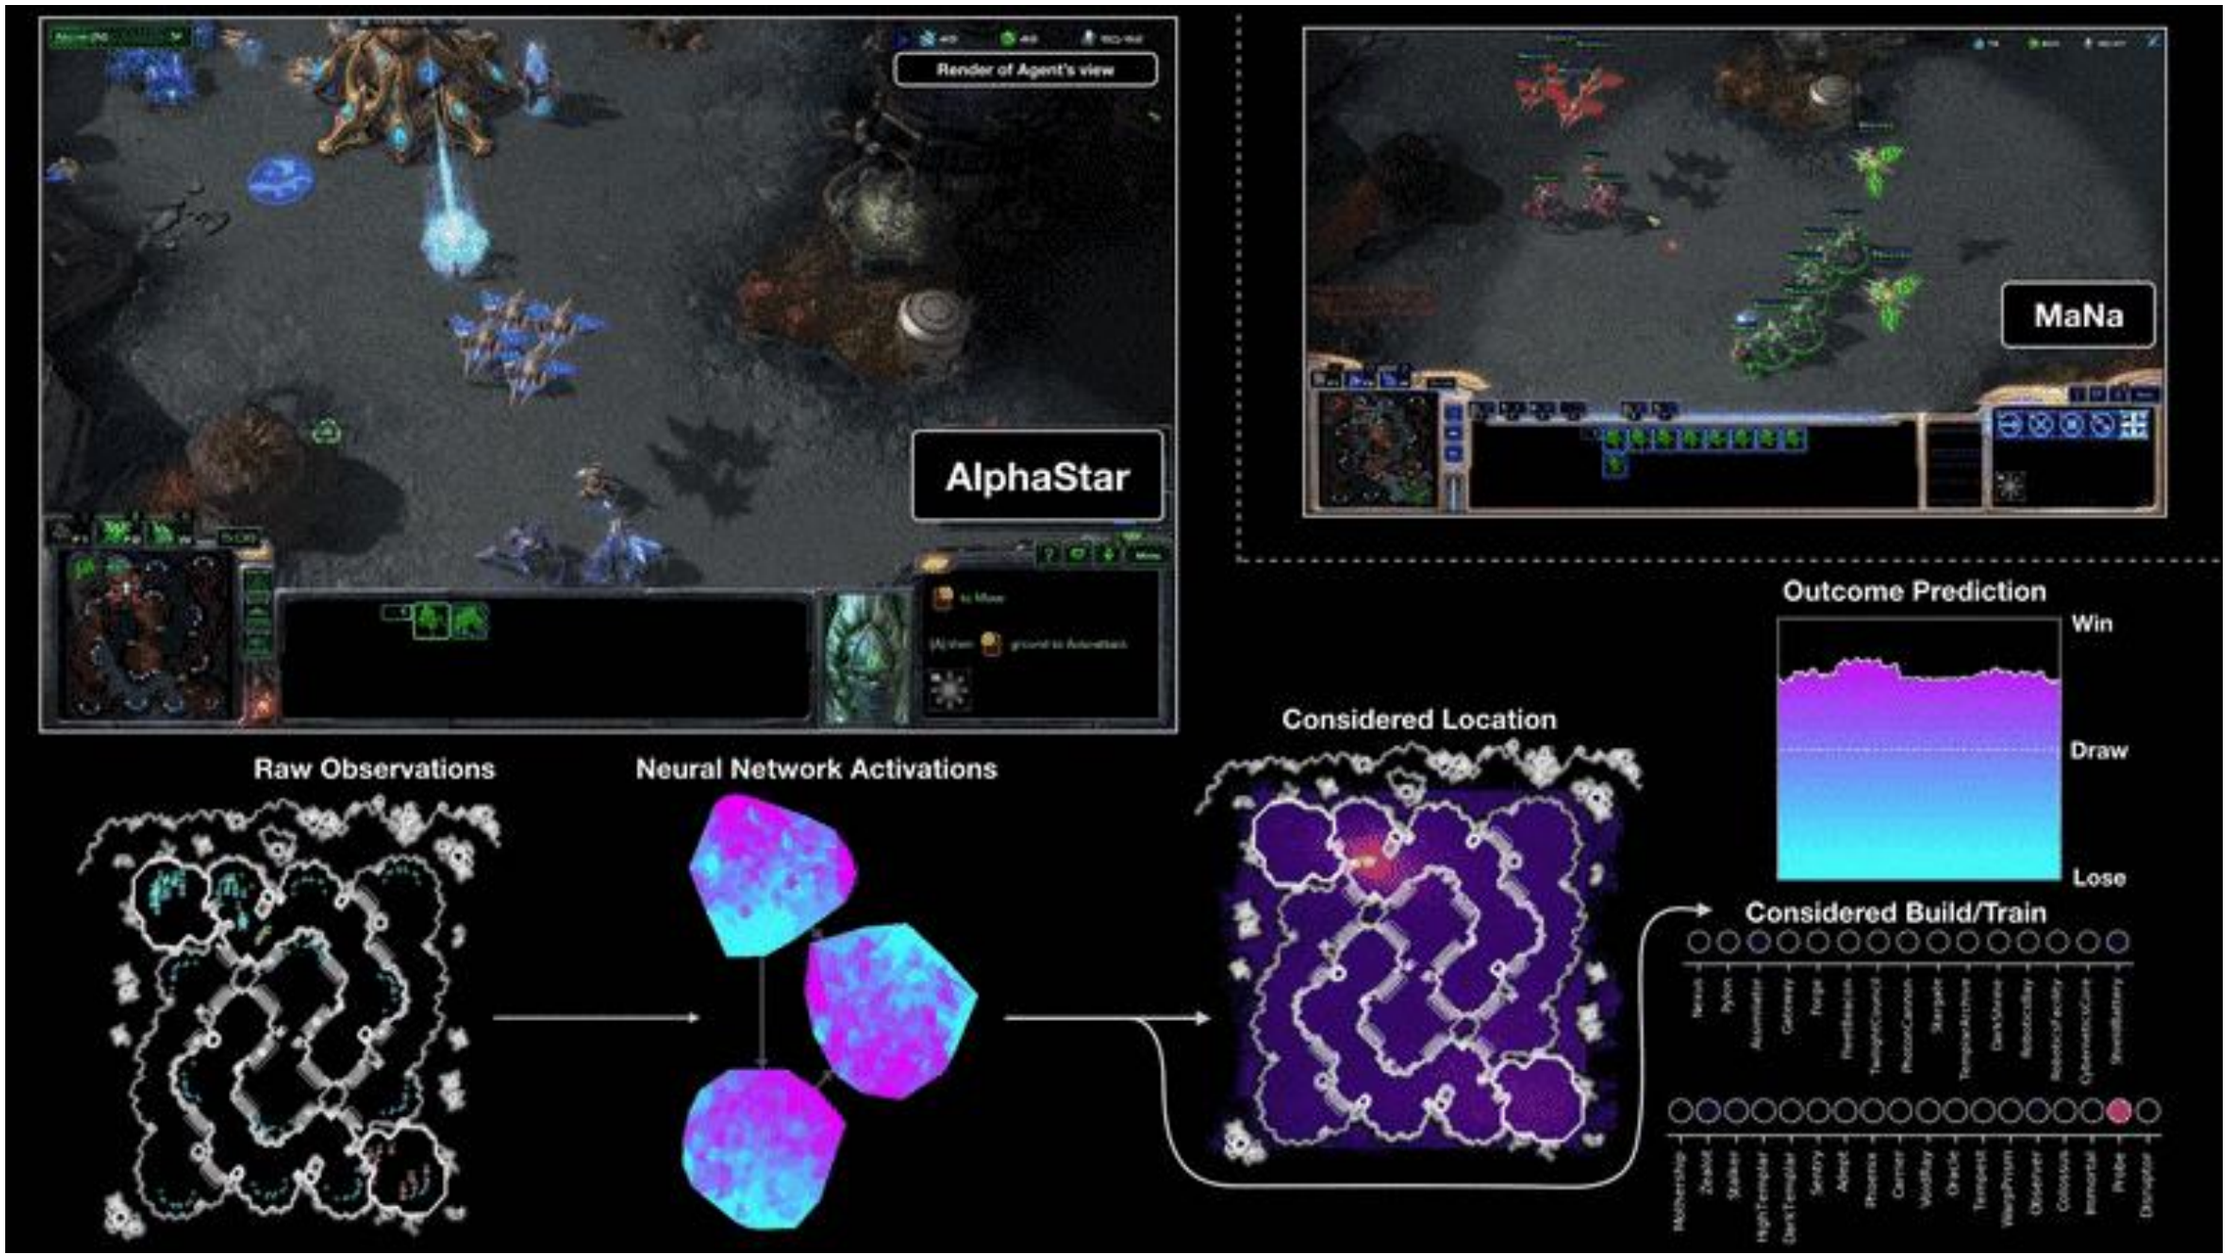

Deep reinforcement learning

Source: https://deepmind.com/blog/alphastar-mastering-real-time-strategy-game-starcraft-ii/#gif-264

Success of deep learning

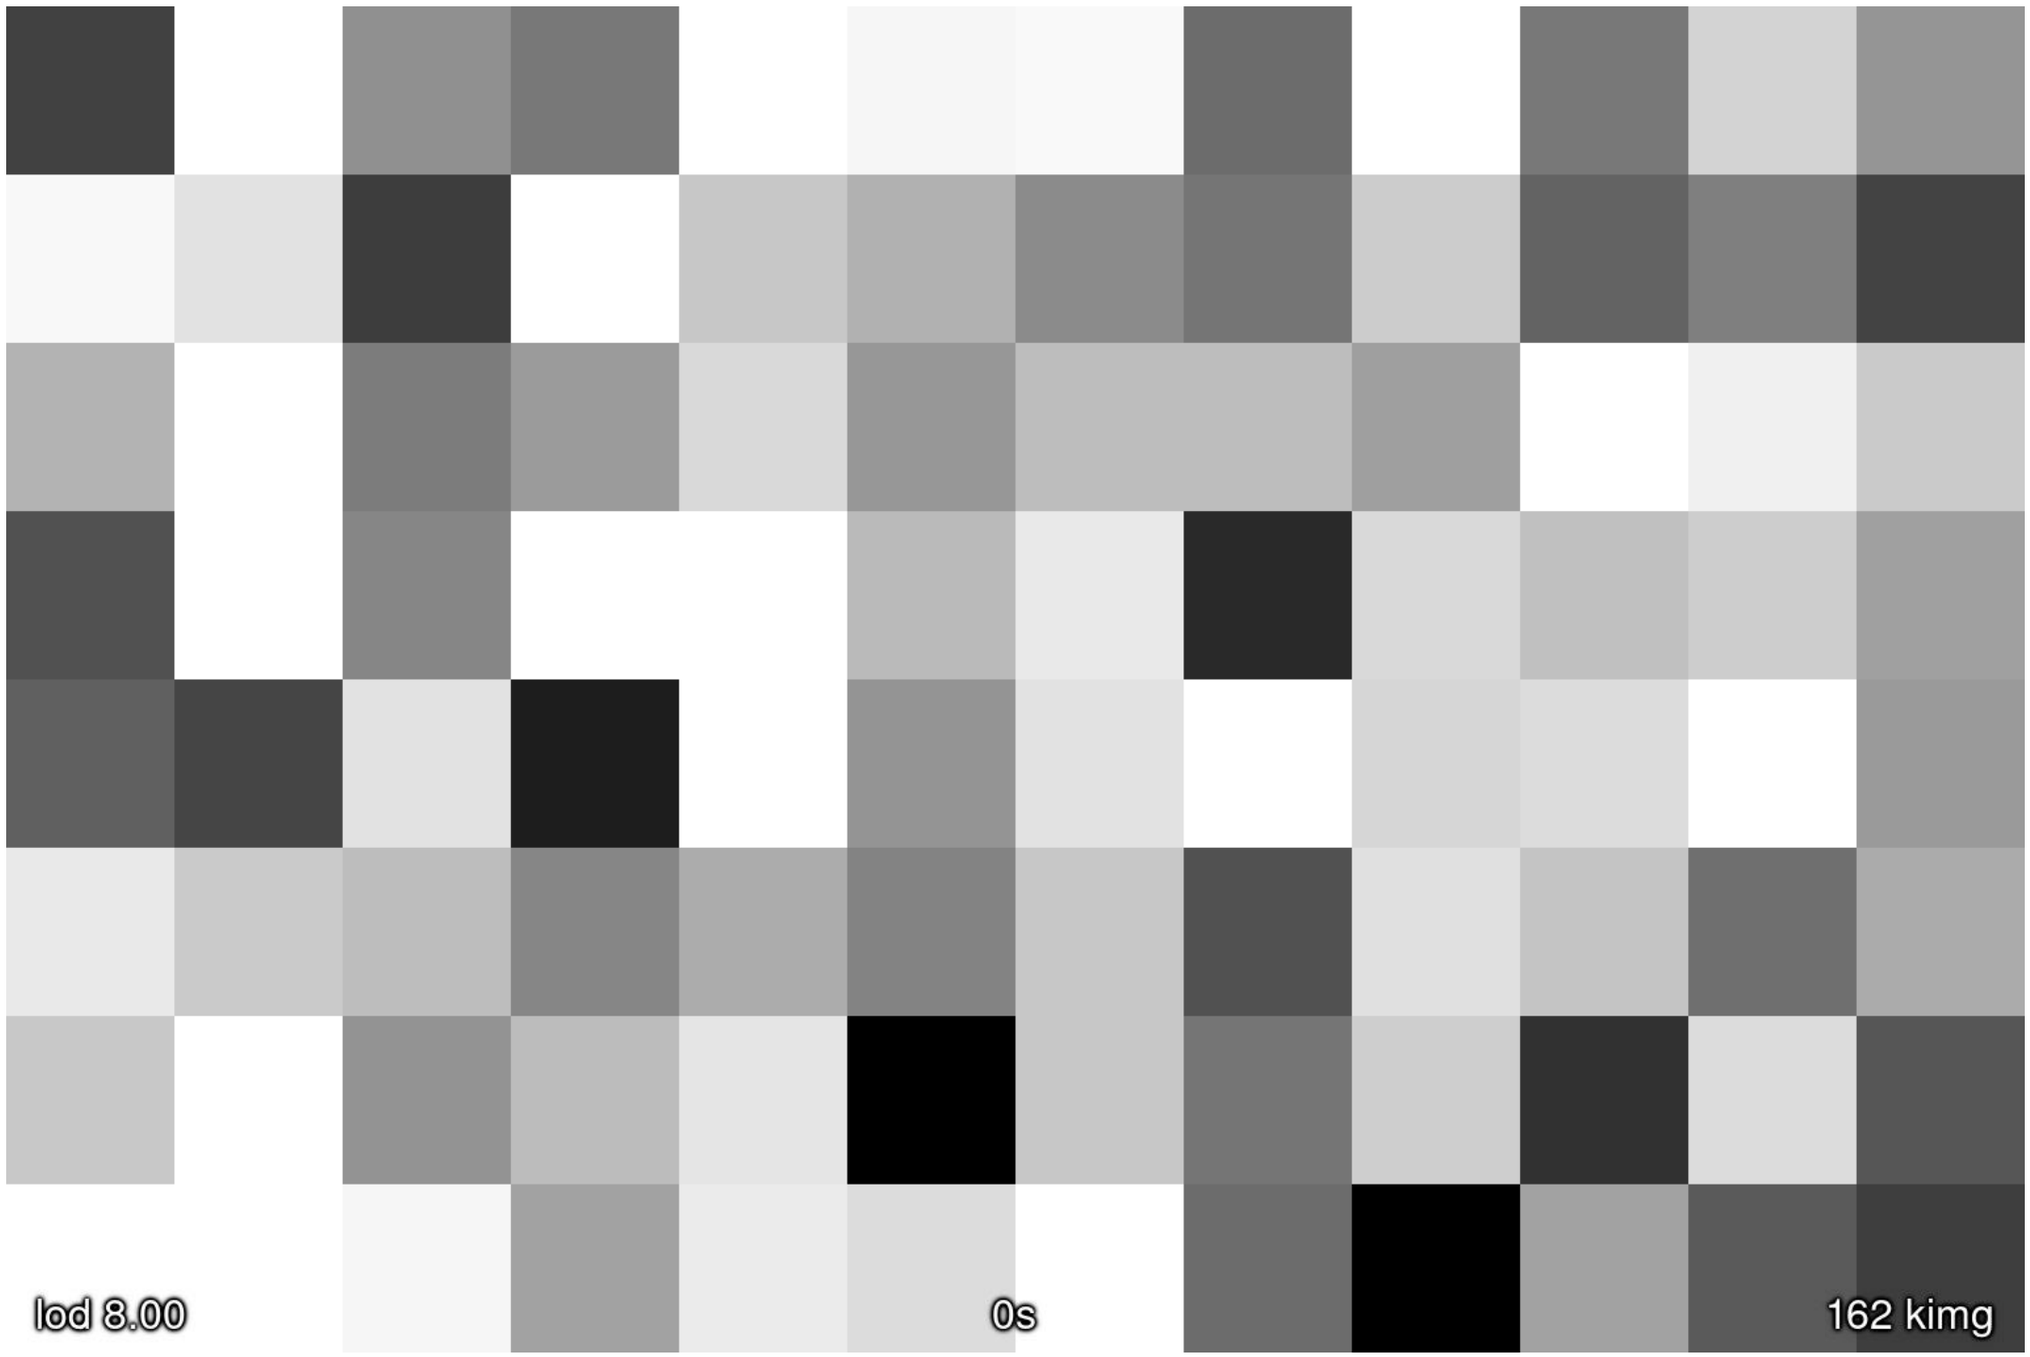

Deep learning

Deep learning is based on neural networks

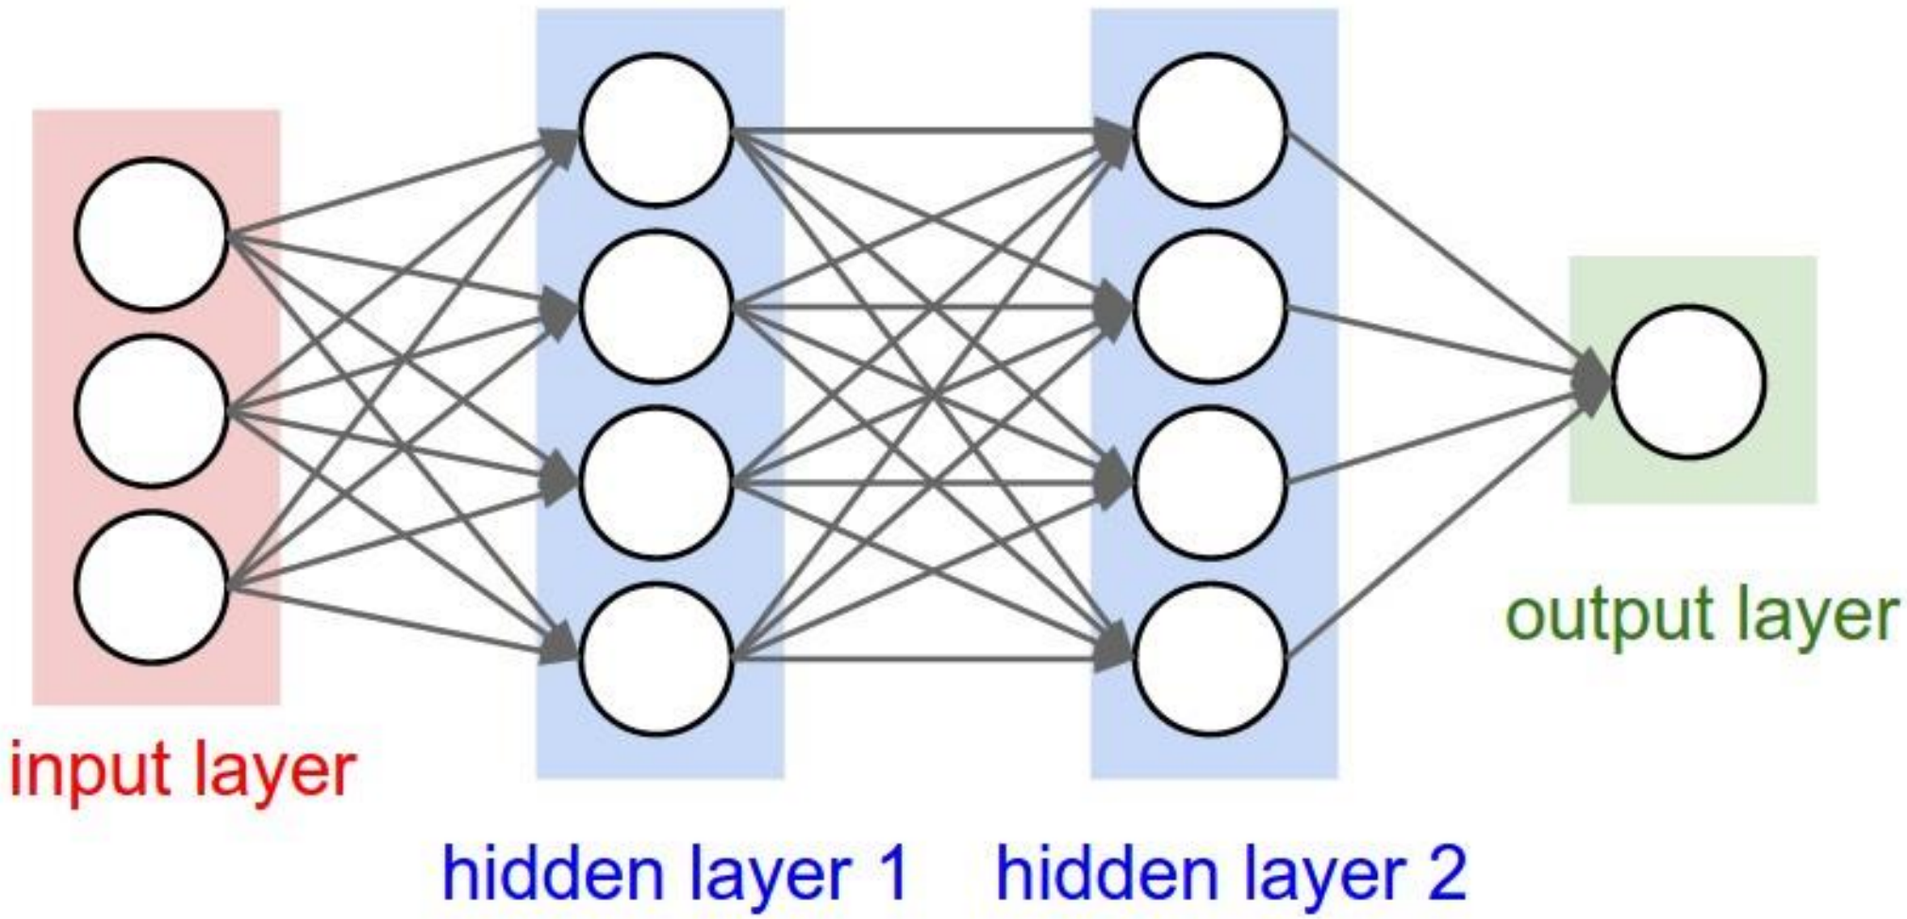

What is a neural network?

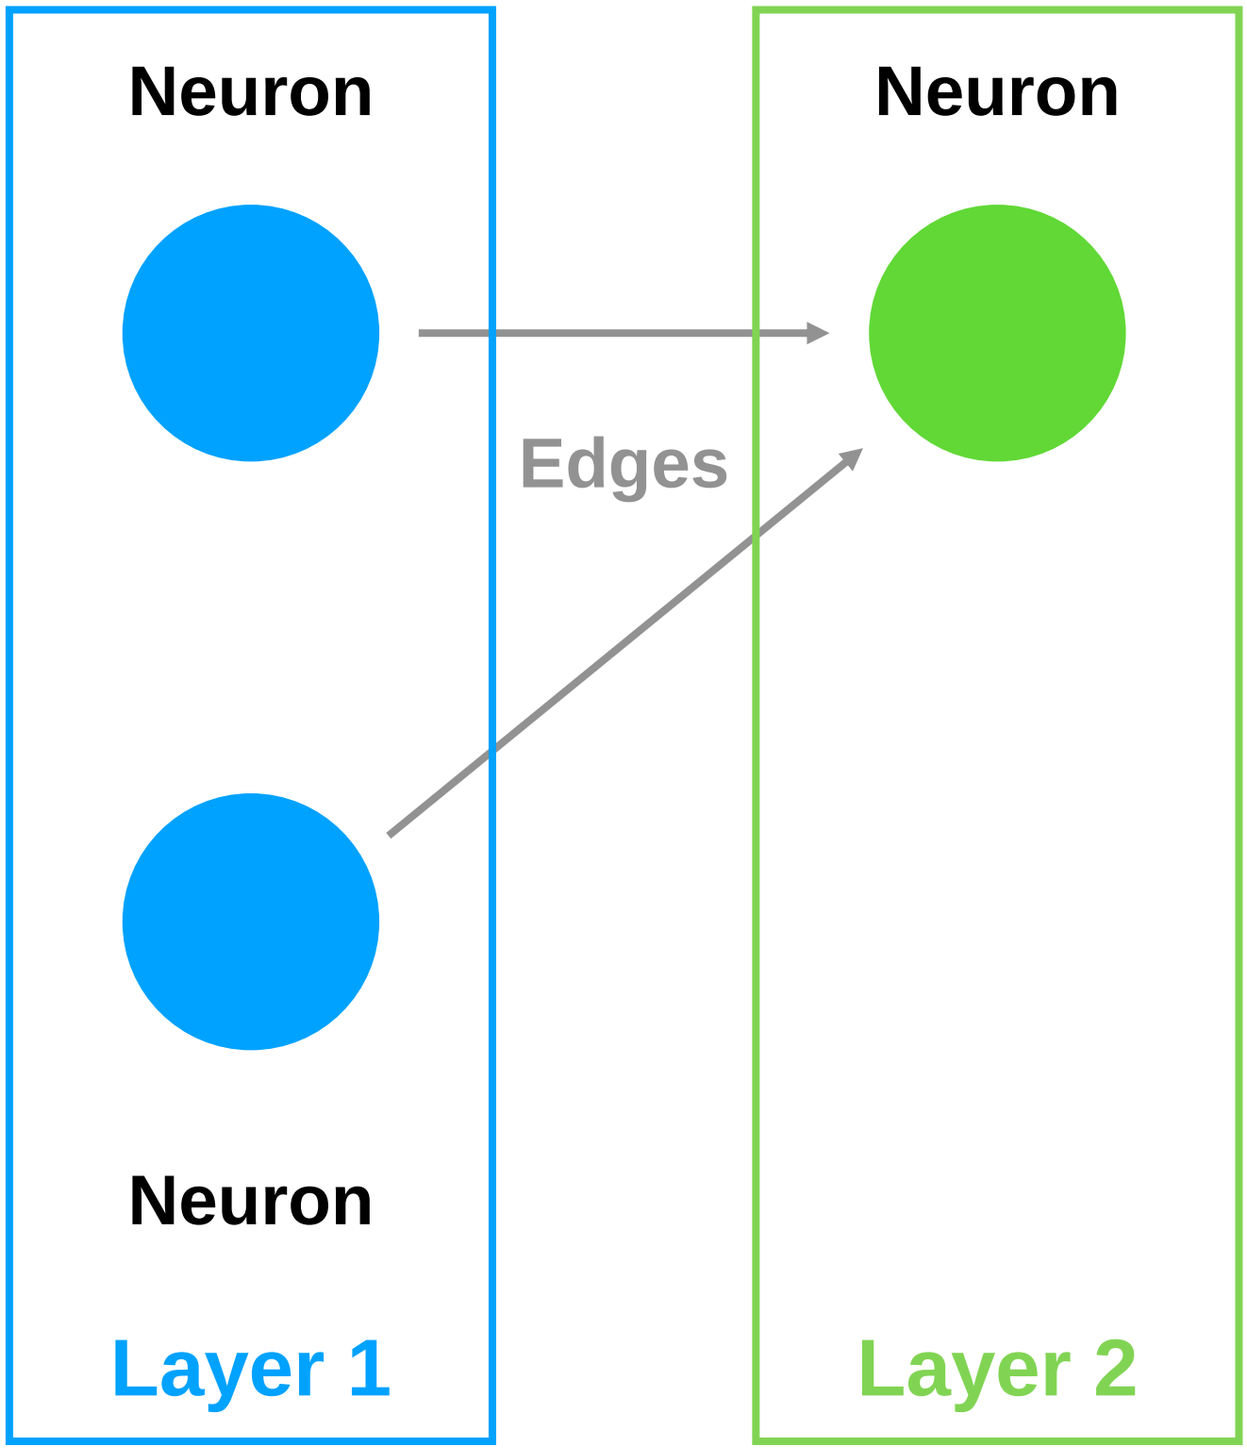

A single neuron

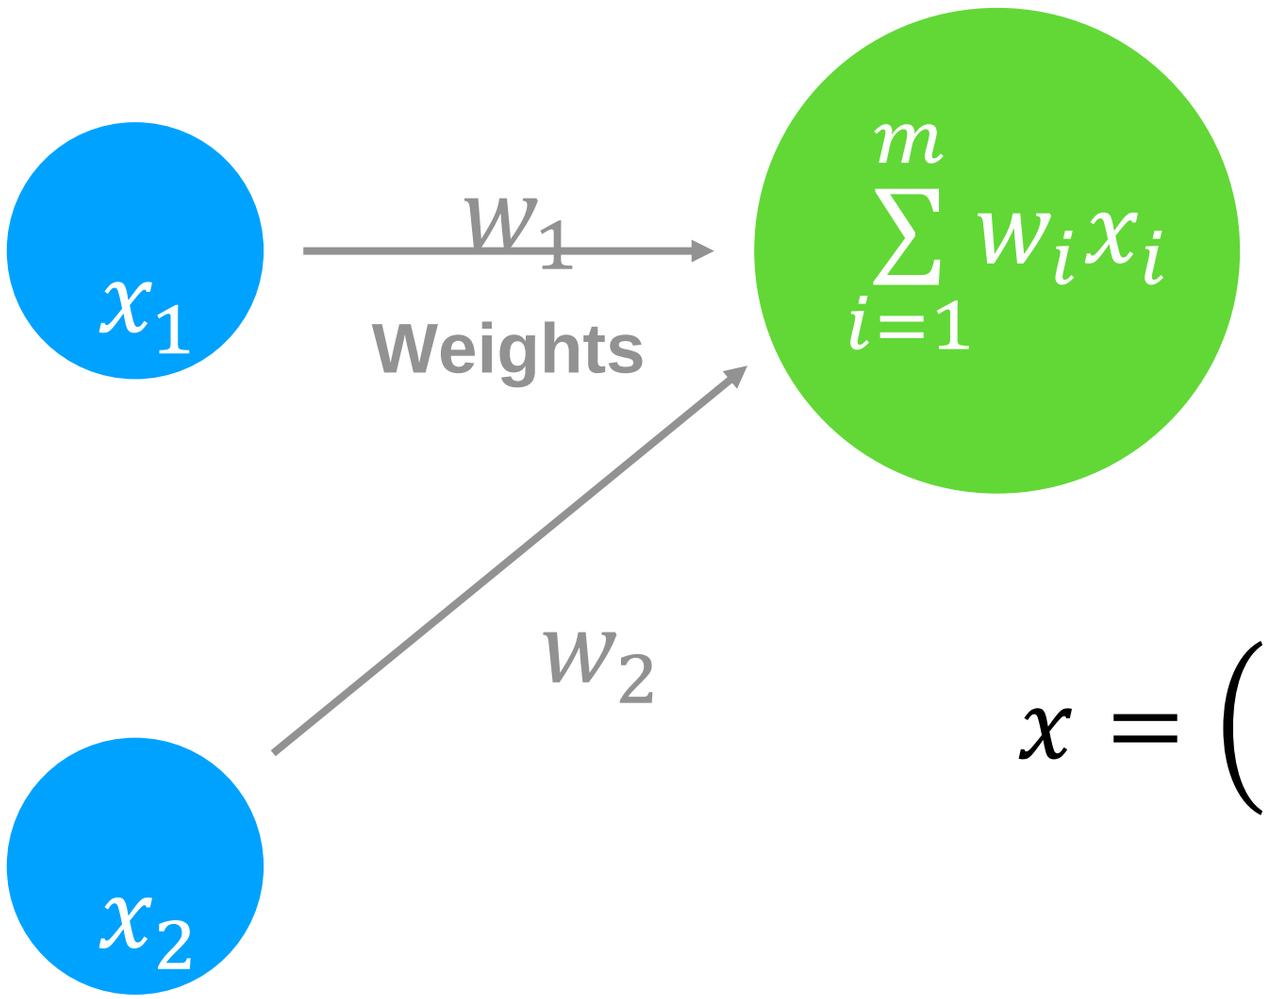

$$w_1 x_1 + w_2 x_2$$

$$x = \begin{pmatrix} 1 \\ 2 \end{pmatrix} \quad w = \begin{pmatrix} 0.5 \\ 1 \end{pmatrix}$$

$$\mathbf{w}^\intercal \mathbf{x} = \begin{pmatrix} 0.5 & 1 \end{pmatrix} \begin{pmatrix} 1 \\ 2 \end{pmatrix} = 2.5$$

A single neuron

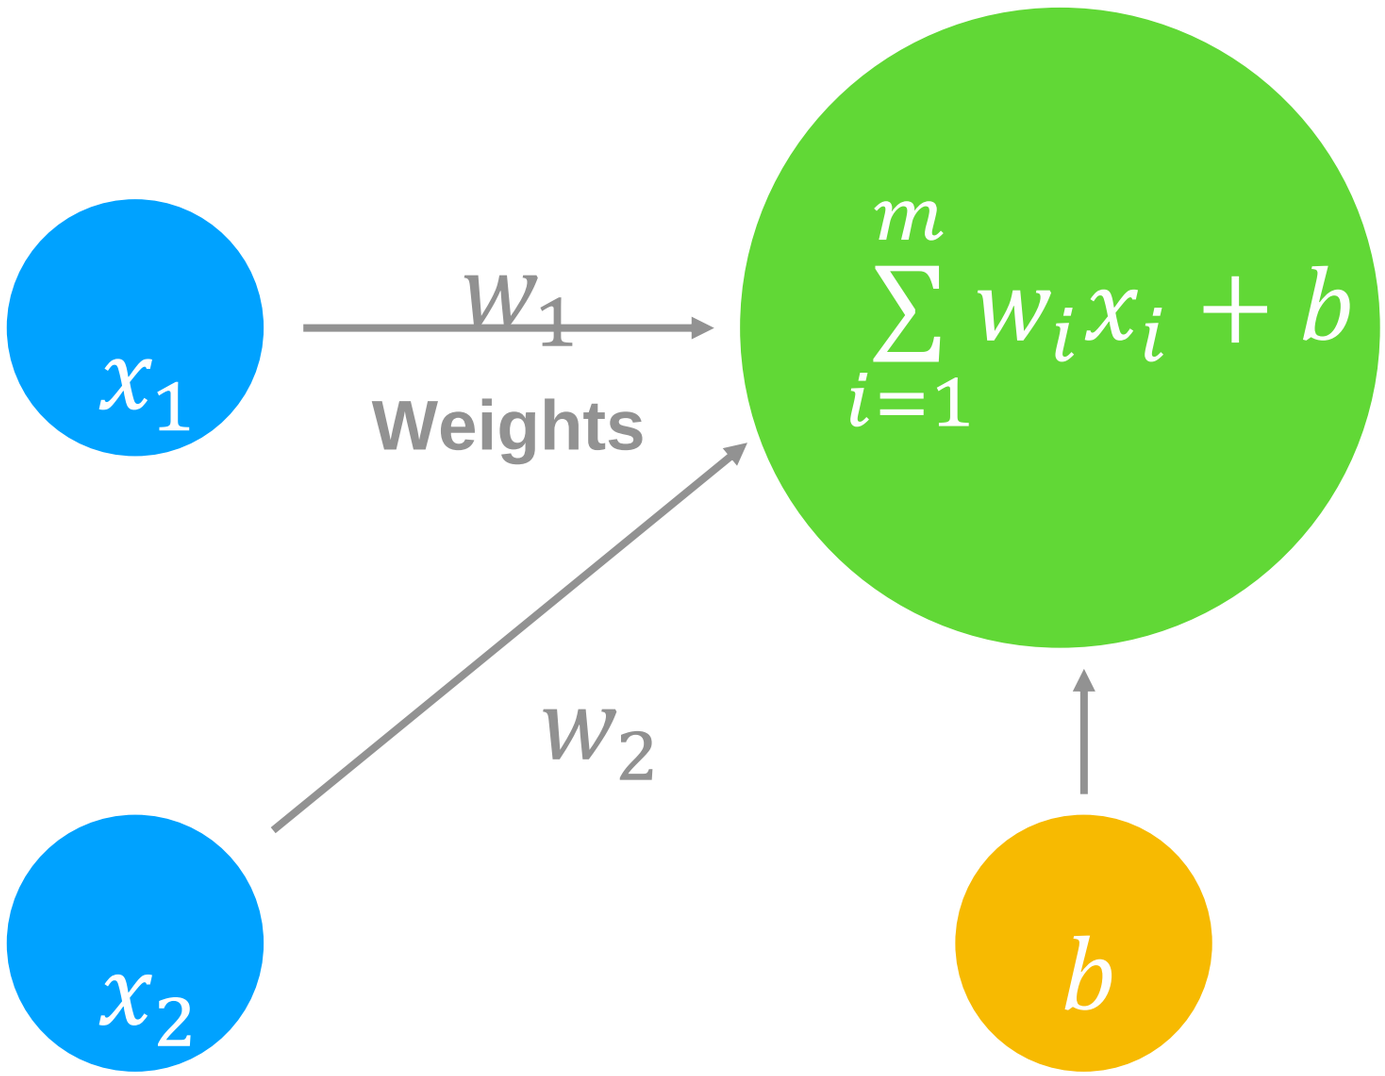

$$w_1 x_1 + w_2 x_2 + b$$

$$b = -1$$

$$\mathbf{w}^\intercal \mathbf{x} + b = \begin{pmatrix} 0.5 & 1 \end{pmatrix} \begin{pmatrix} 1 \\ 2 \end{pmatrix} - 1 = 1.5$$

A single neuron

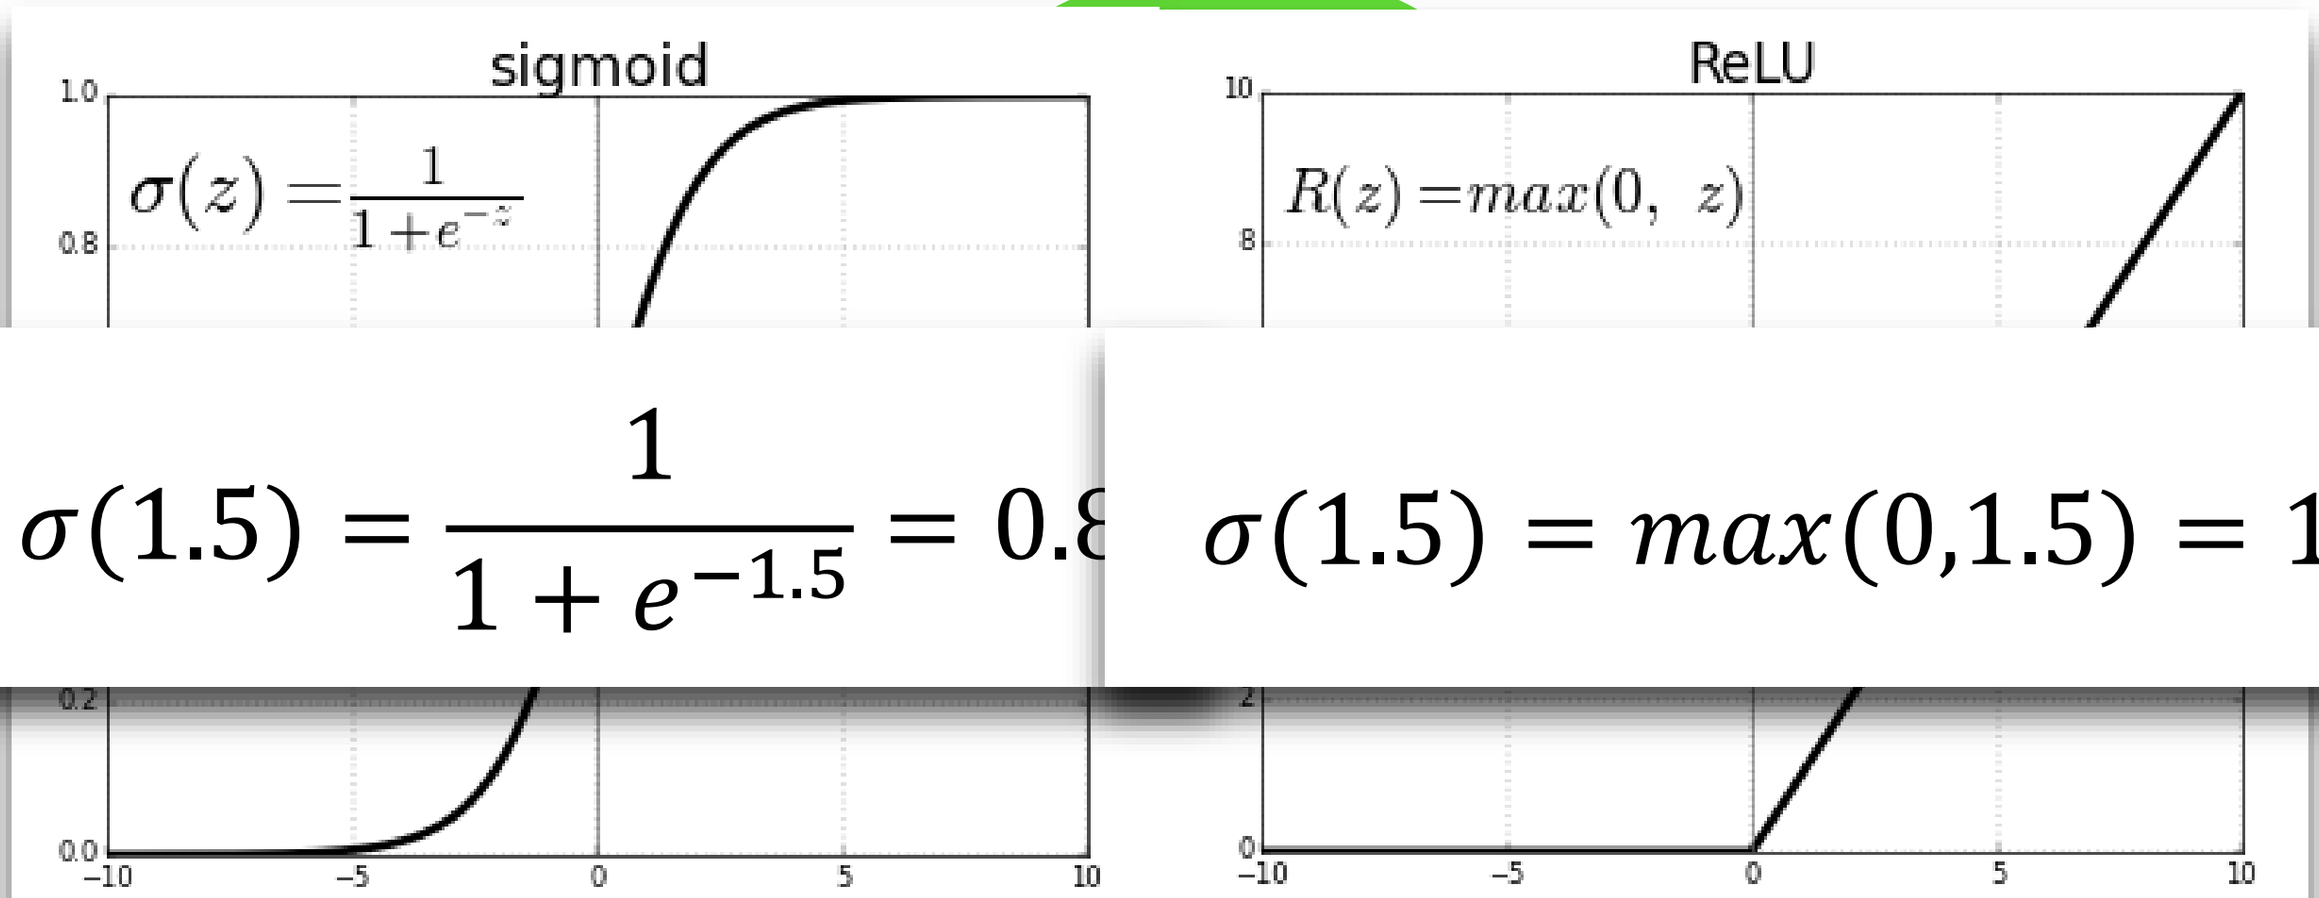

$$\sigma(1.5) = \frac{1}{1 + e^{-1.5}} = 0.8$$

$$\sigma(1.5) = max(0,1.5) = 1.5$$

x_2

b

Source: [Activation Functions in Neural Networks]

Prediction

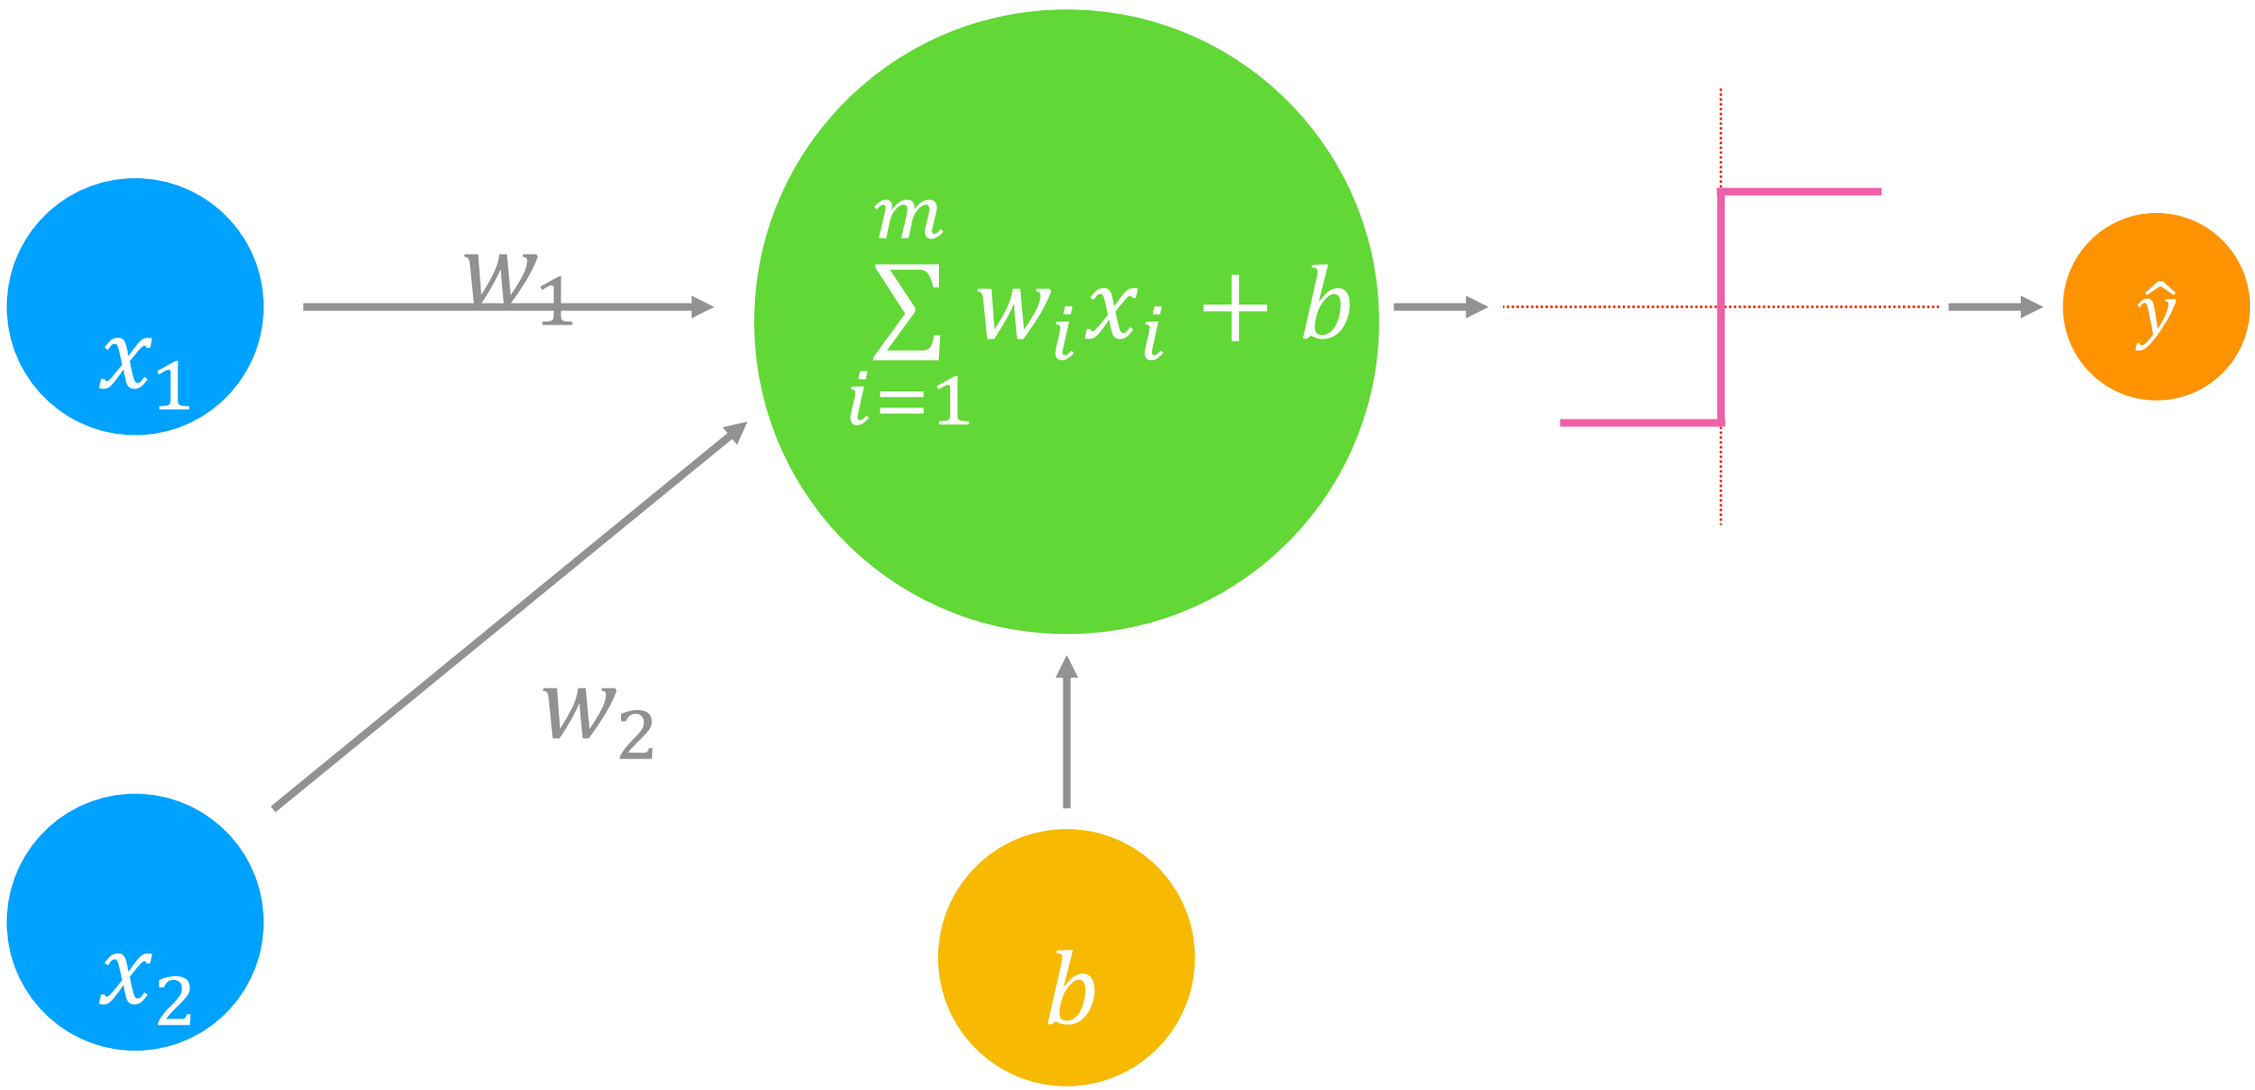
x_1
w_1
$\sum_{i=1}^{m} w_i x_i + b$
$\hat{y}$
w_2
x_2
b
inputs
weights
summation + bias
activation
output

Before we begin

- Refresh some R
- Tensor manipulation
- Implement a neuron from scratch

Hands-on

- Exercises in Jupyter notebooks
- Notebooks run on Lisa, the national compute cluster
  - Lisa has approximately 300 nodes, 4800 cores
  - During course hours: 16 full nodes reserved
  - Outside course hours: 1 node / capacity for 8 notebooks
- First users of Jupyter @ Lisa!

Hands-on

- Notebook session will start when you login (may take up to a minute or so).
- Closing the browser does *not* shut down the notebook session.
- To manually start / shut down a notebook session:
  - ‘Home’ (top left) => Start / Stop My Server.
  - Please shut down sessions when working outside of course hours, since capacity is limited!

Hands-on

First time: login & change password @ https://portal.surfsara.nl

Go to https://https://jupyter.lisa.surfsara.nl:8000/

Notebook: 01a-neuron-from-scratch.ipynb

14:30-15:30

Notebook recap

- Data preparation / restructuring is necessary
- Weights change the ‘slope’
- Bias shifts decision boundary

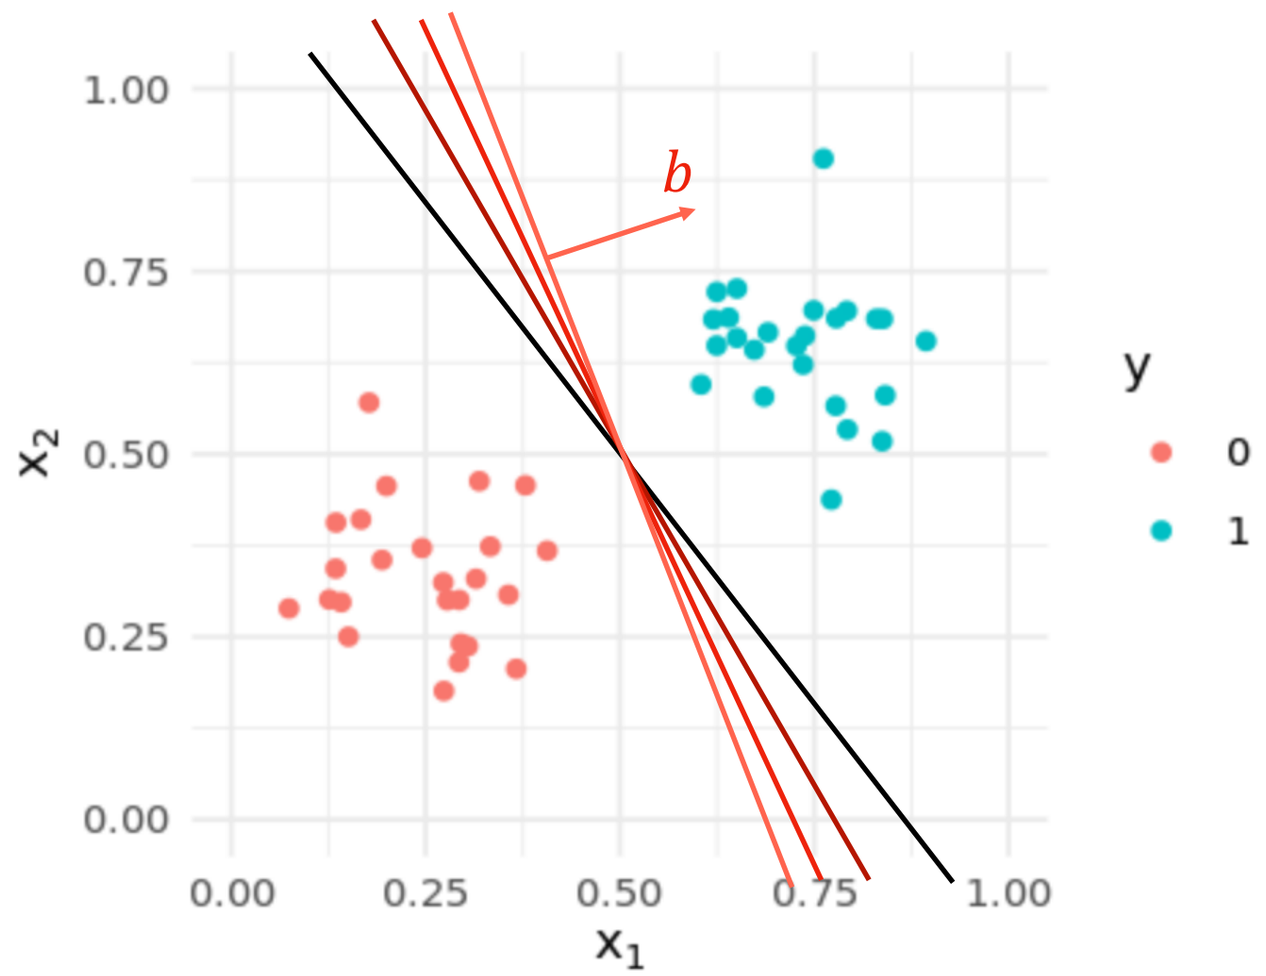

Today’s program

14:00-14:30 What is ML / DL? What is a neuron?

14:30-15:30 Hands on: building a neuron from scratch

15:30-15:45 How do neural networks work?

15:45-16:15 Break

16:15-16:30 How do neural networks work? (continued)

16:30-17:15 Hands on: building a network from scratch

17:15-18:00 Loss function & updating weights

18:00-19:00 Diner

19:00-19:45 Hands on: the XOR problem

19:45-20:00 Dataset splitting & Performance evaluation

20:00-21:00 Hands on: Keras on fashion Mnist

If we finish early: Machine learning tasks (classification vs regression)

Why the success?

No single reason, much of the theory is old

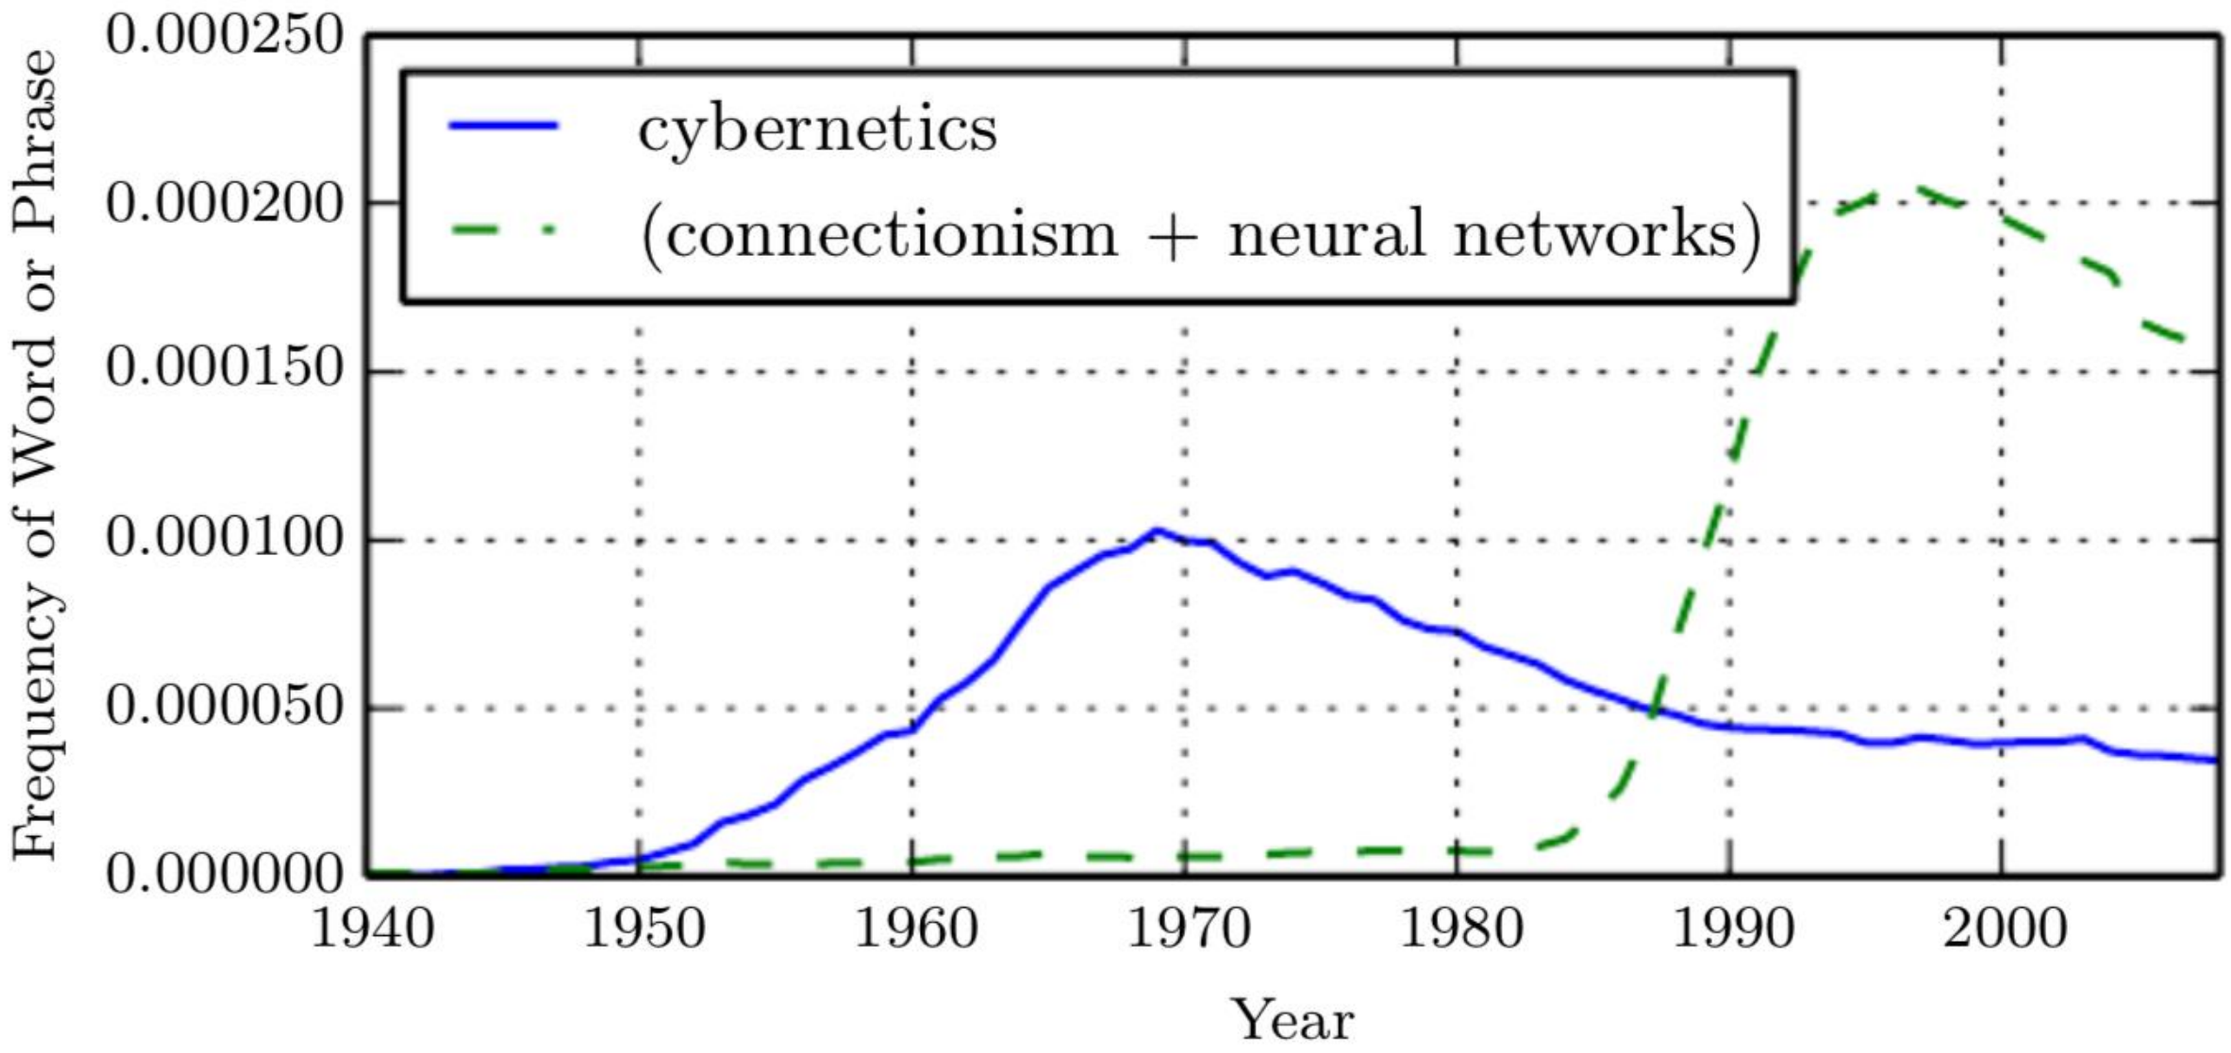

Why the success?

More data available

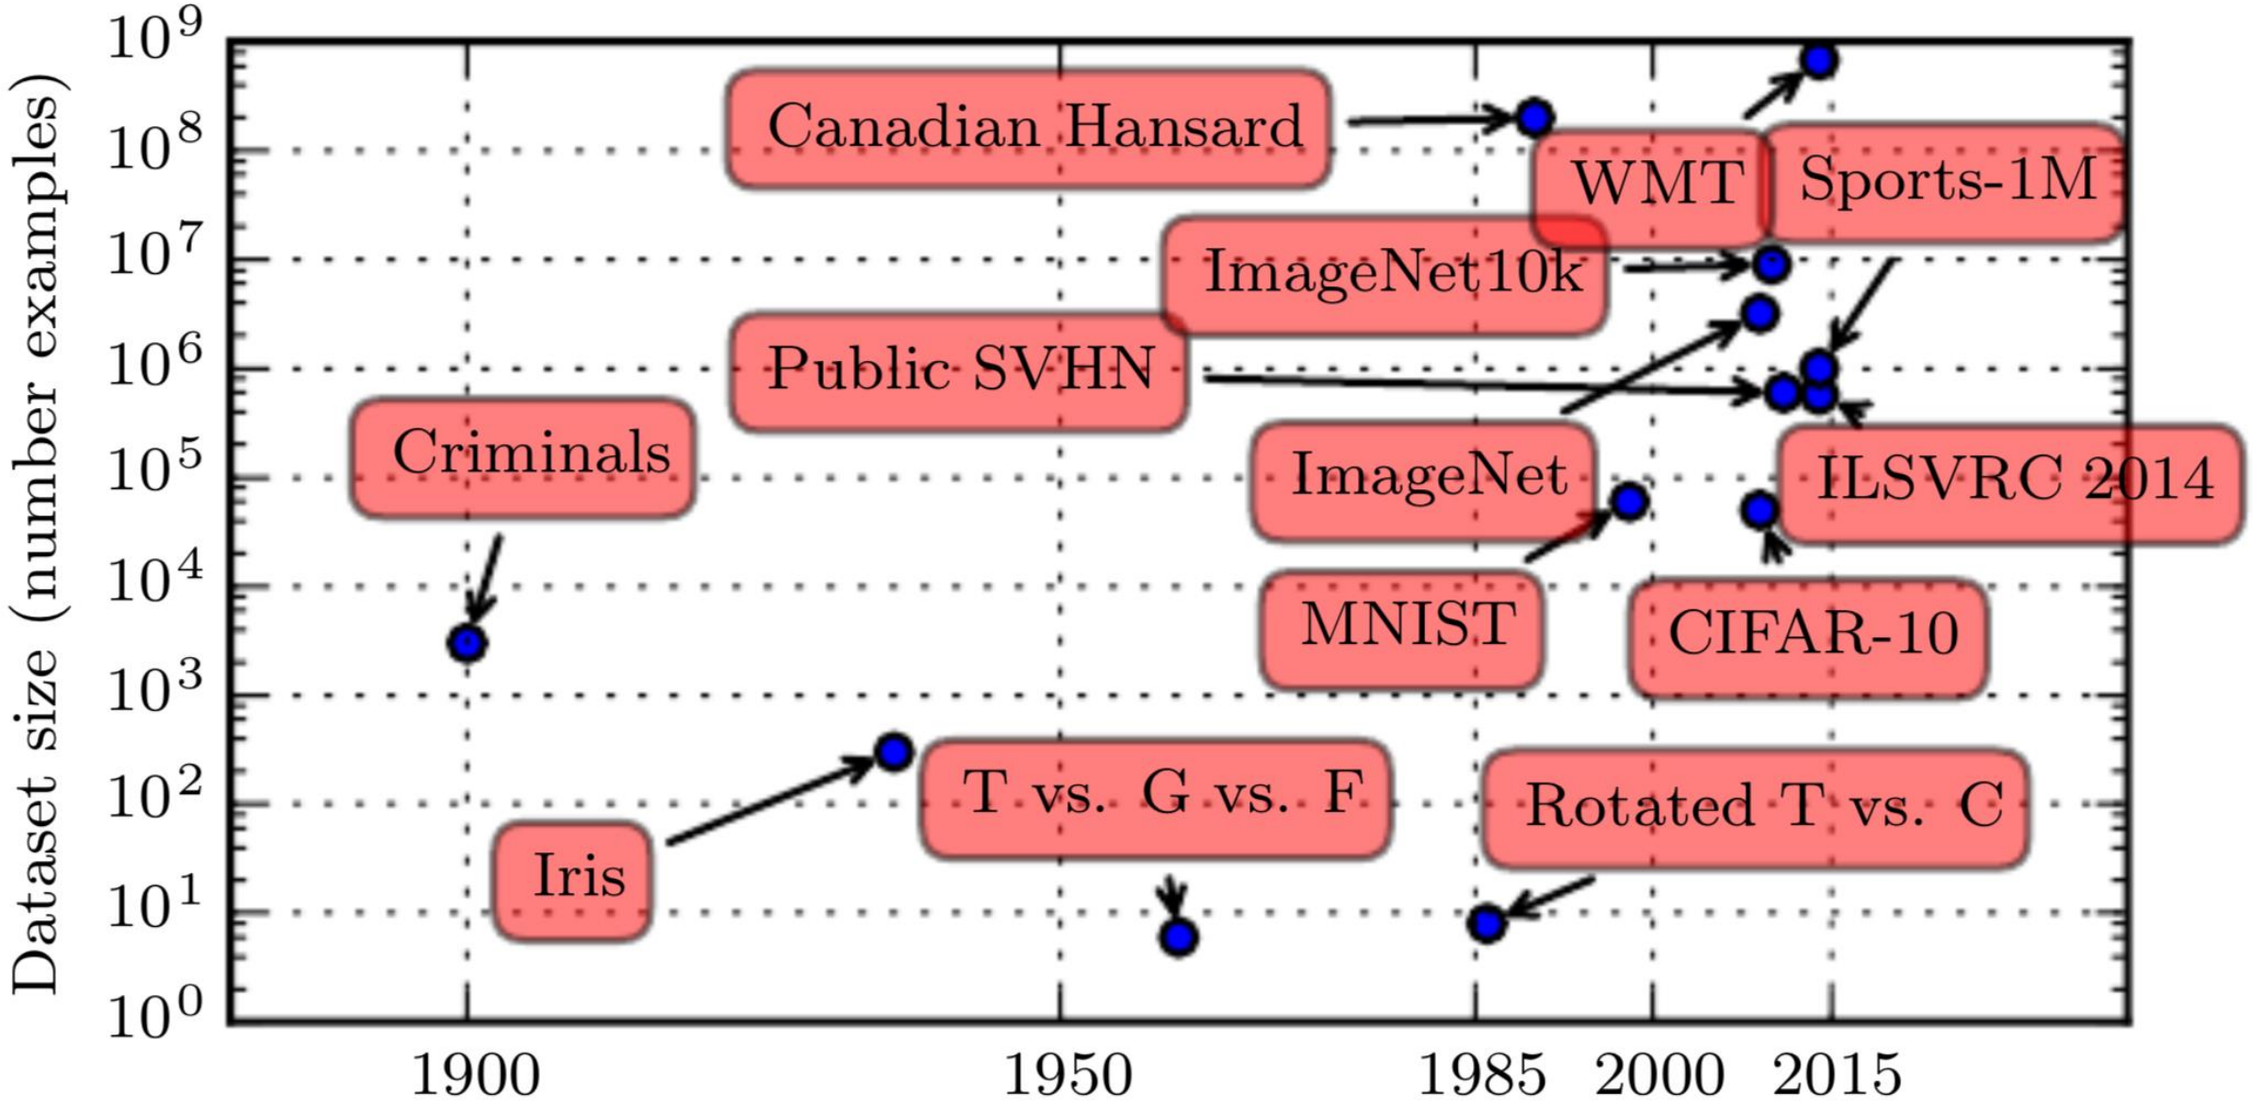

Source: [Deep Learning, Ian Goodfellow and Yoshua Bengio and Aaron Courville]

Why the success?

Bigger models

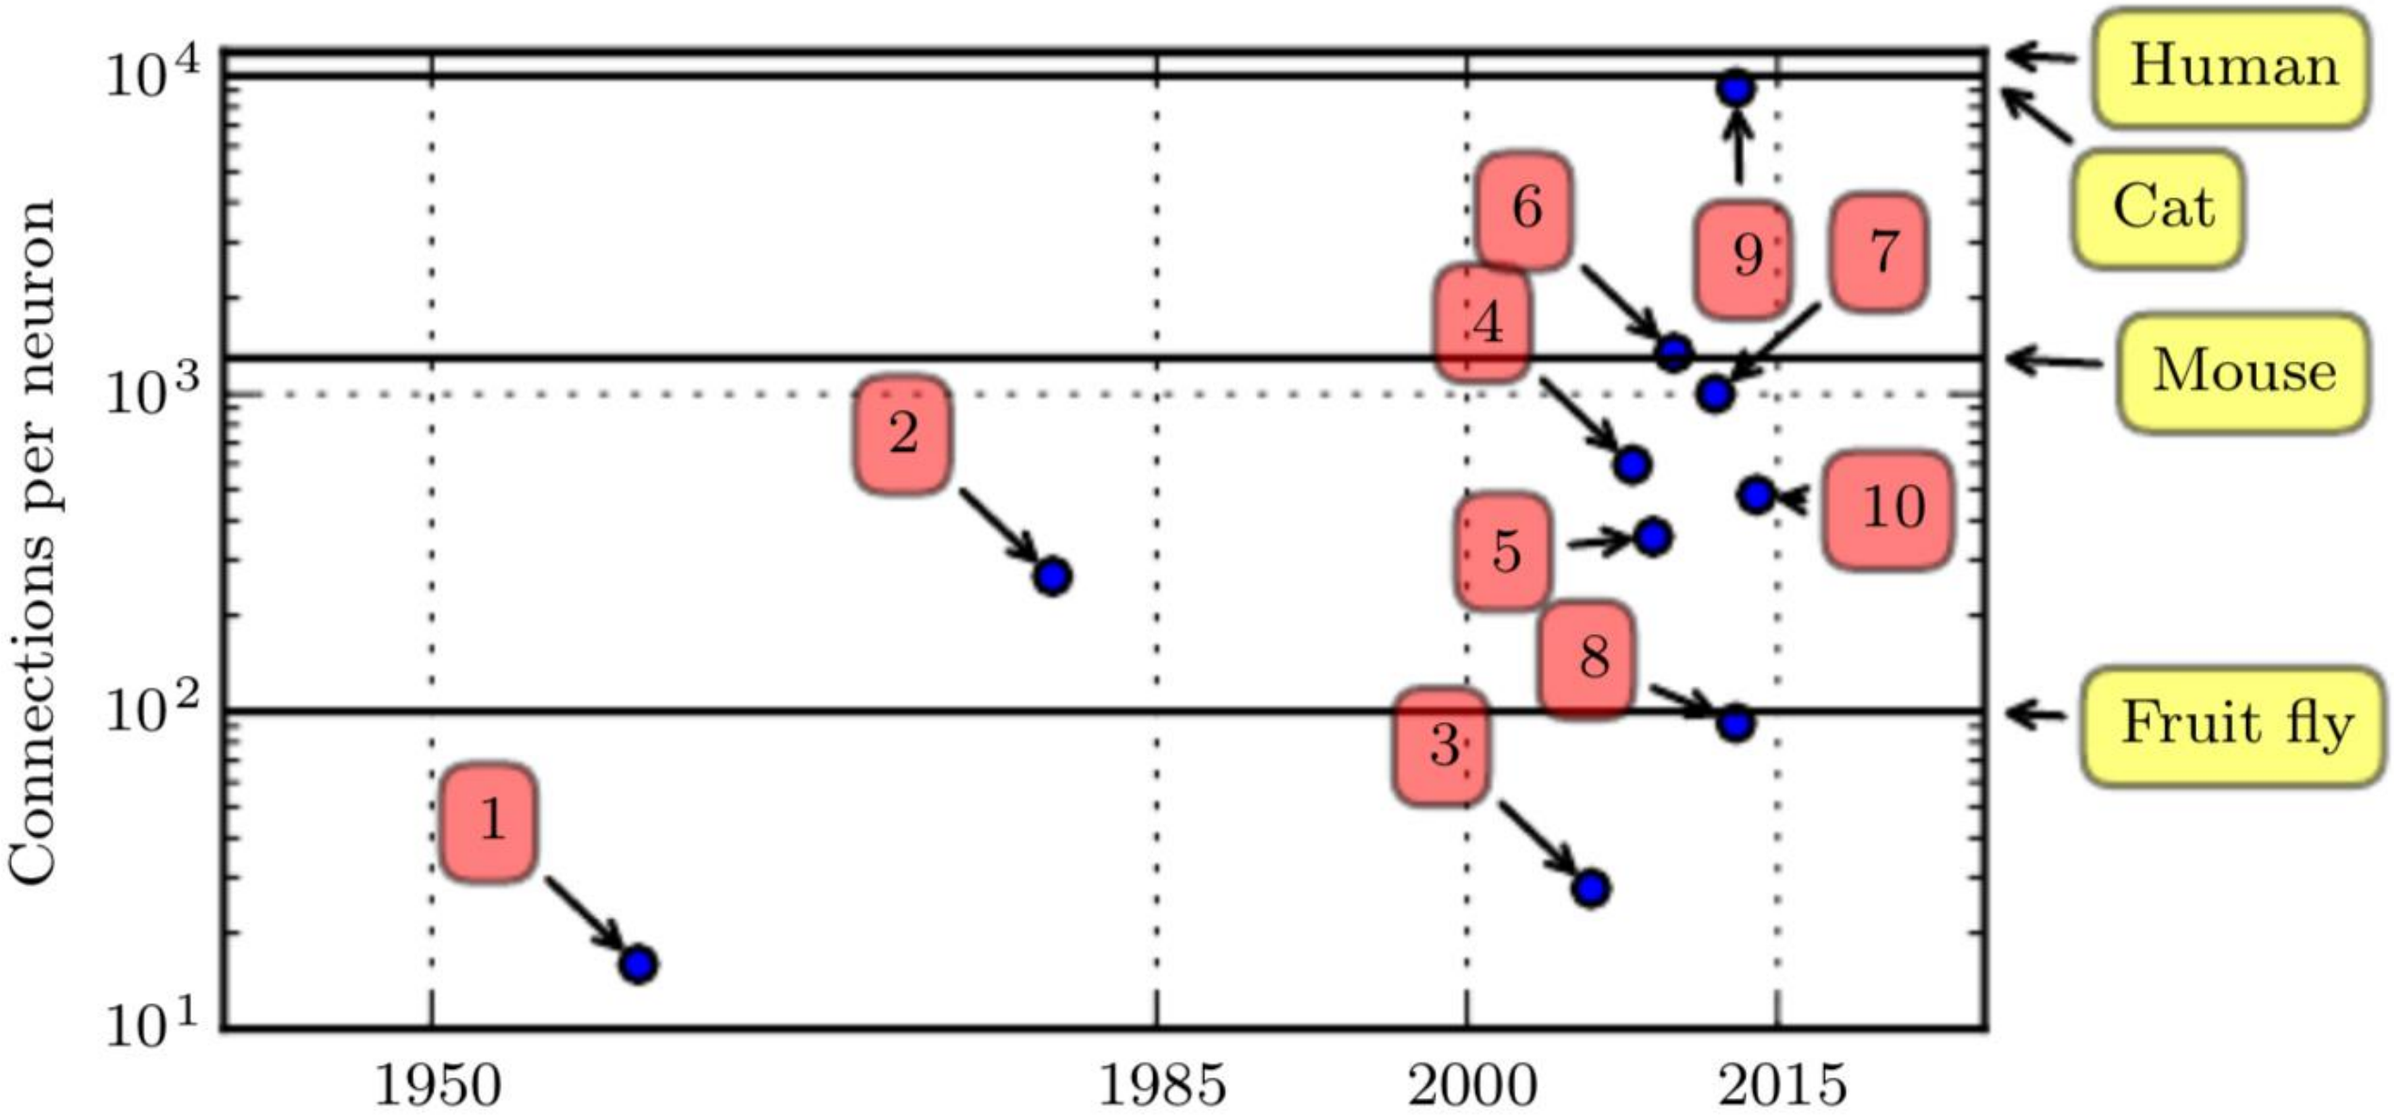

Source: [Deep Learning, Ian Goodfellow and Yoshua Bengio and Aaron Courville]

Why the success?

More compute power

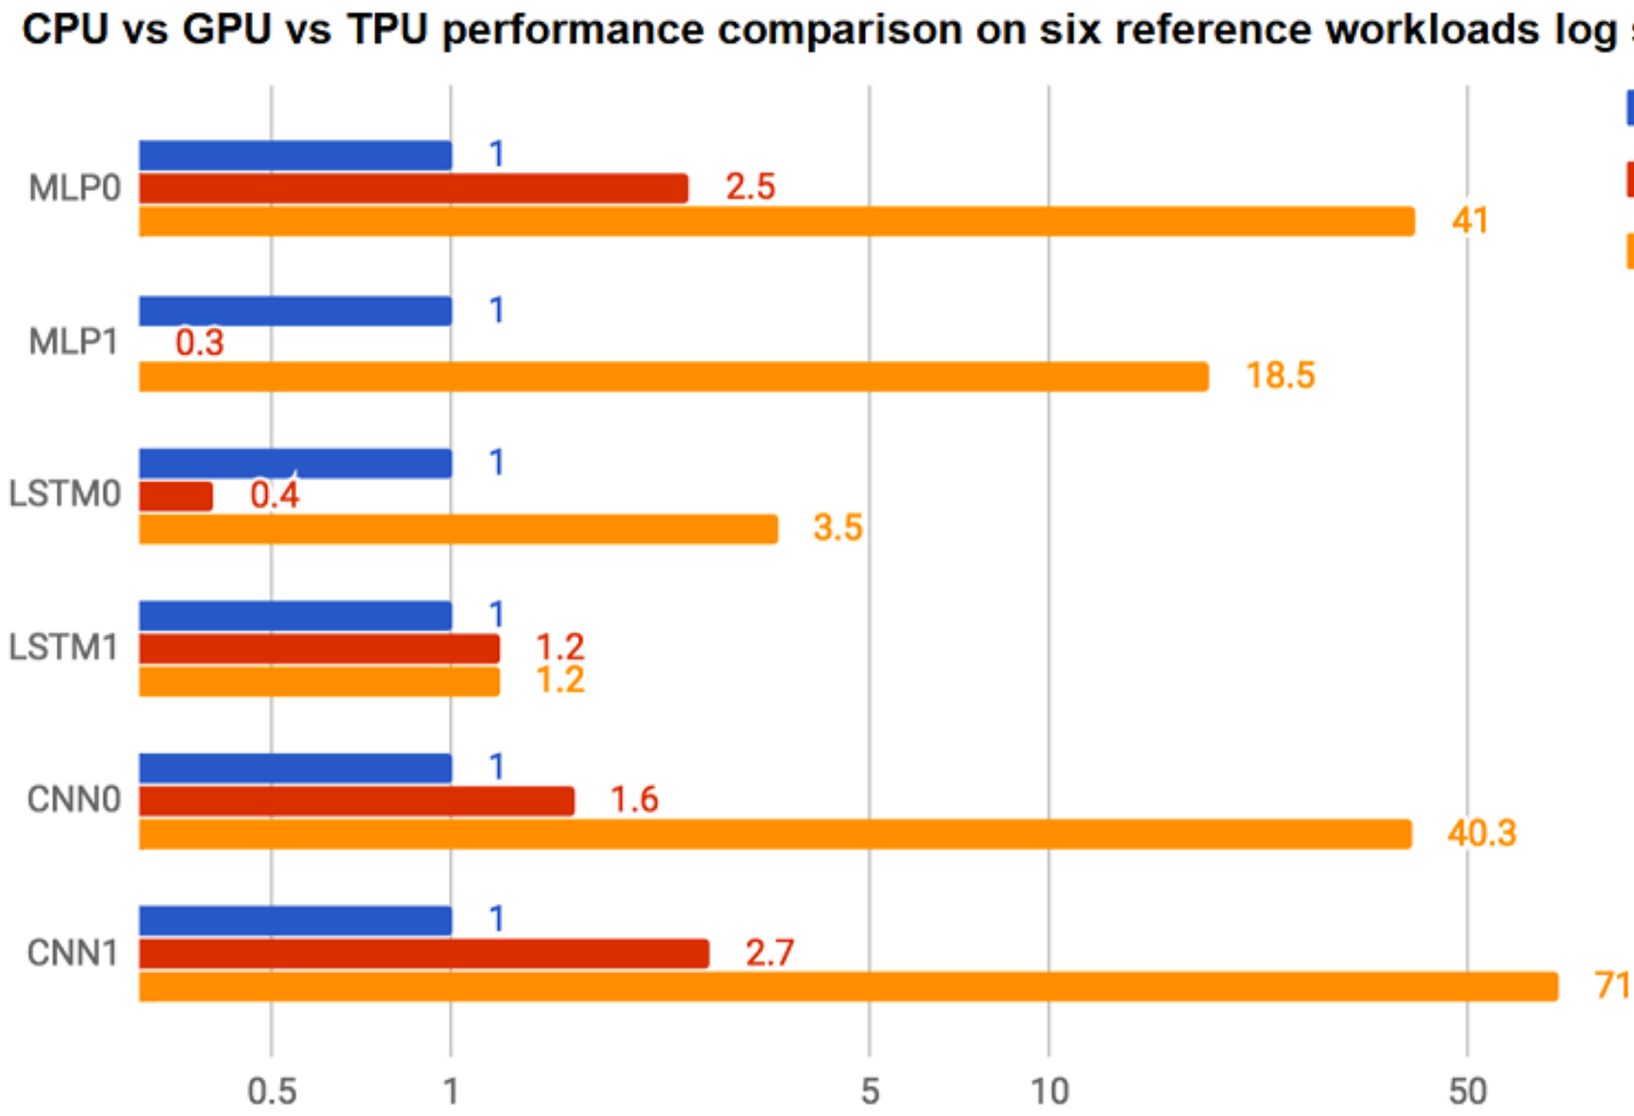

Source: https://www.androidauthority.com/google-powering-ai-816146/

How do NNs work?

- We want to make a model to predict housing prices

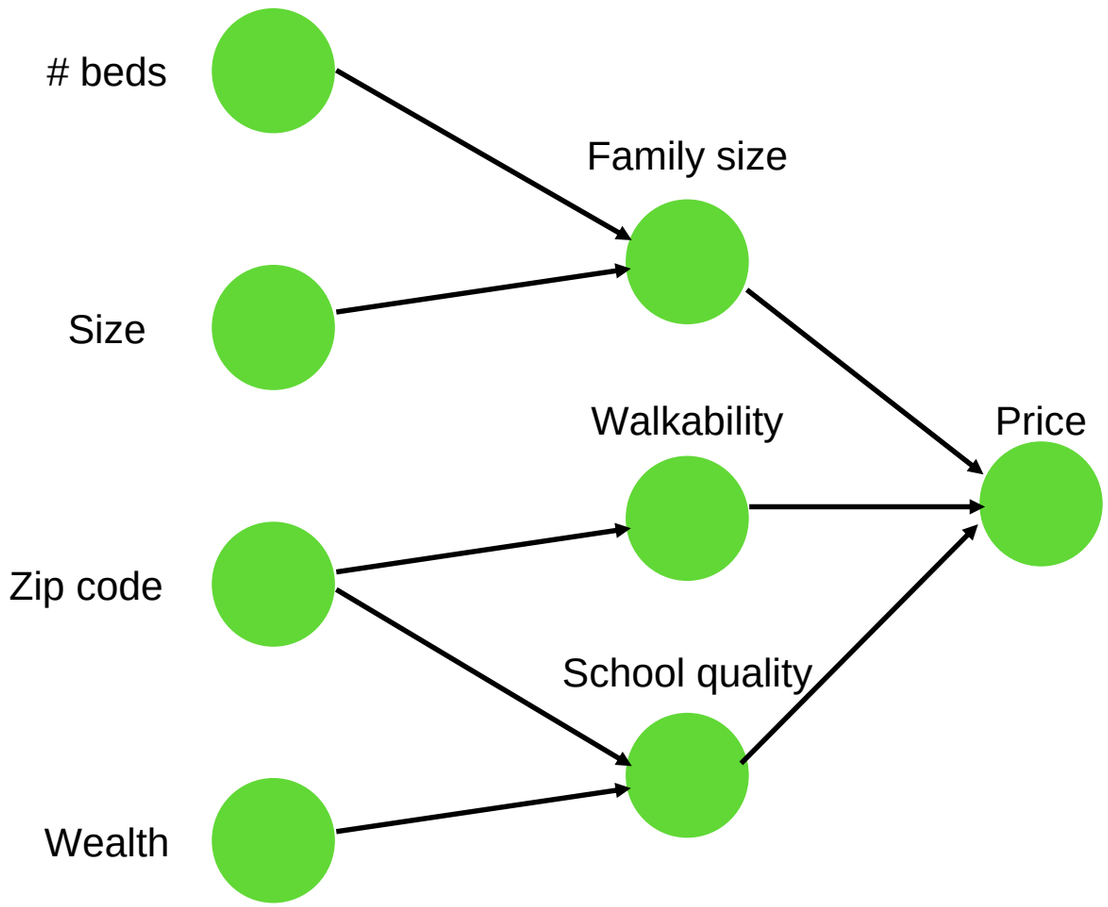

Hidden layers

- Or just hit it with a hammer and let the model learn the features

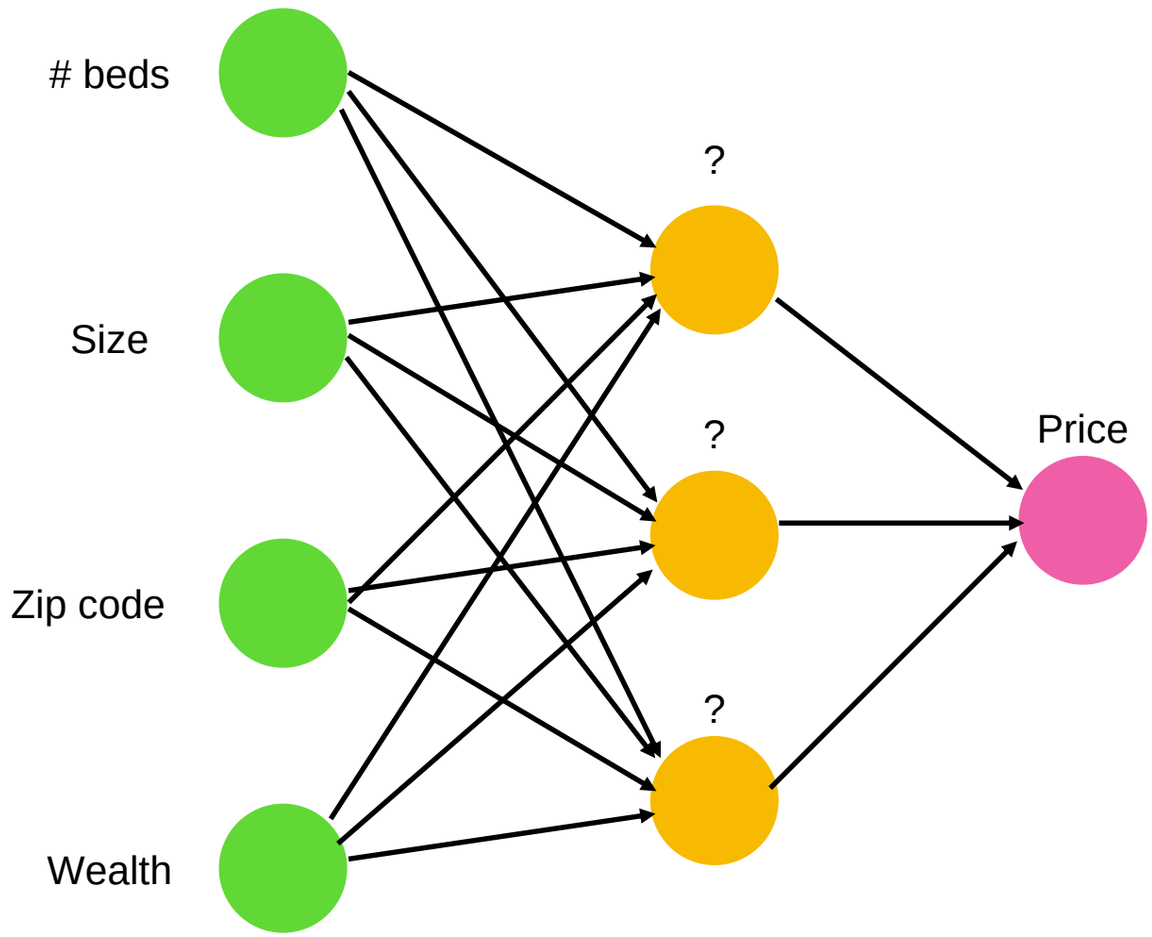

How does it work?

Deep neural networks learn hierarchical feature representations

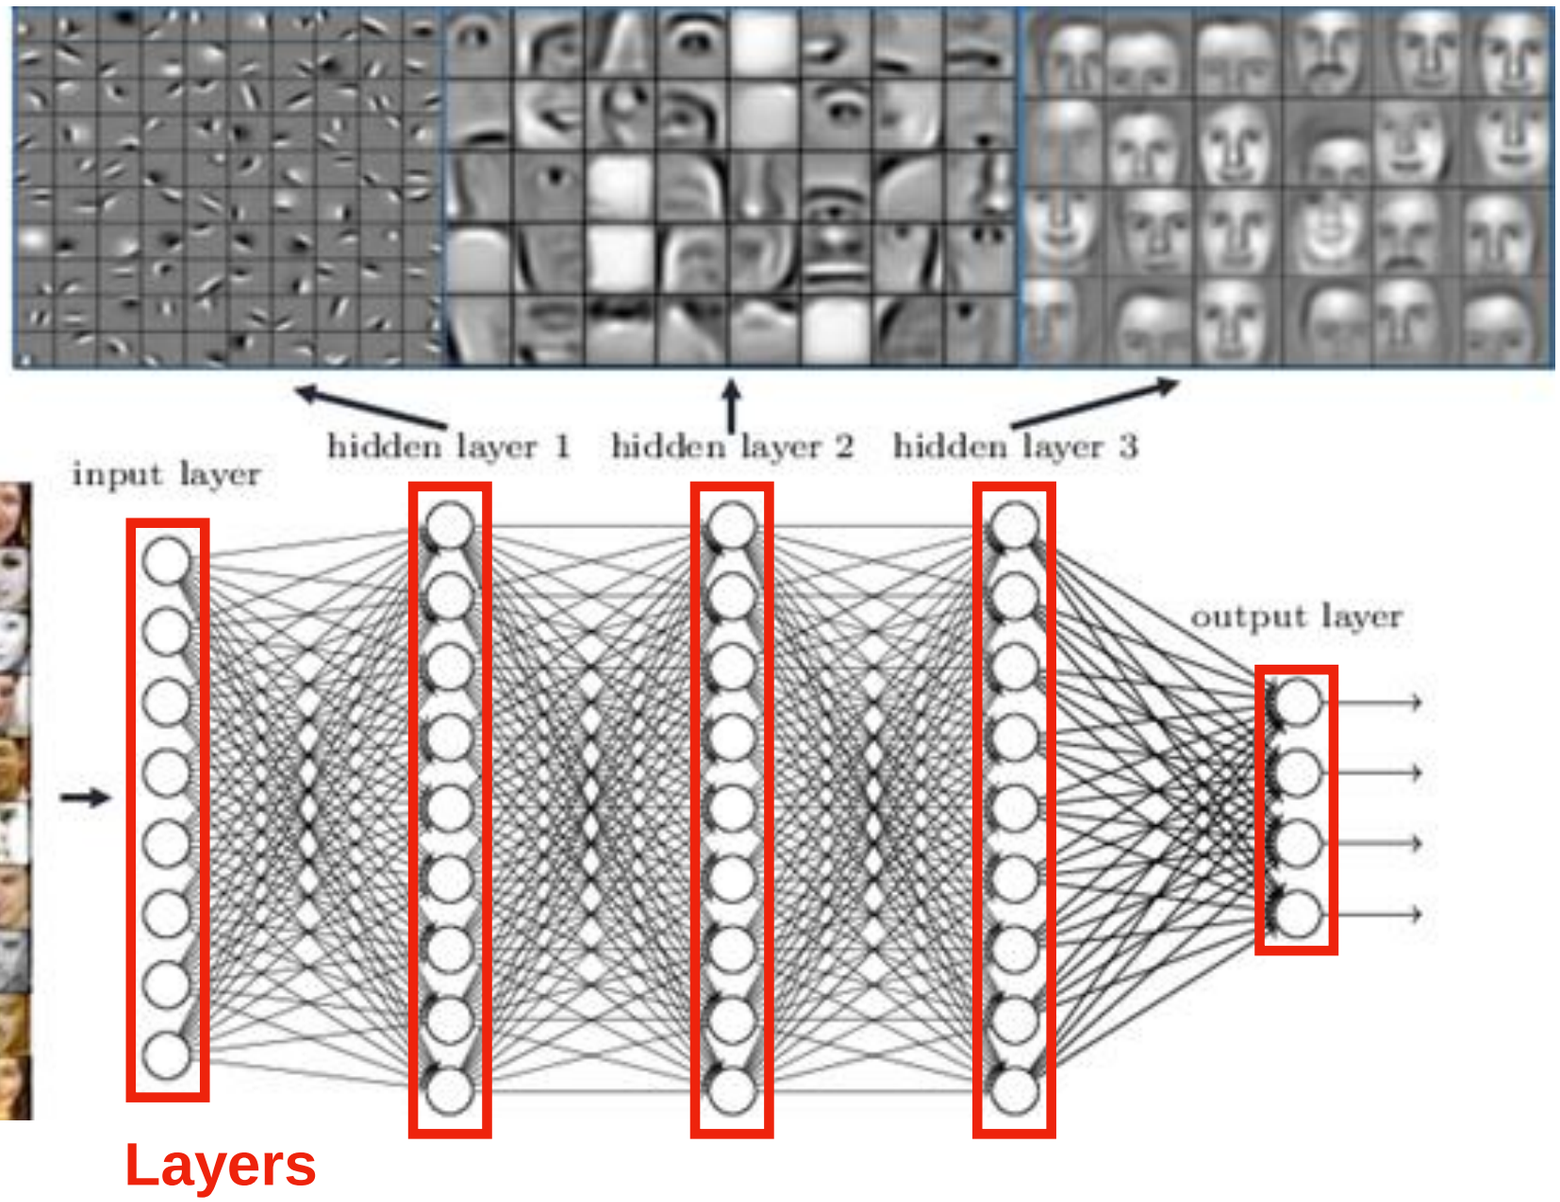

Source: [A Comprehensive Hands-on Guide to Transfer Learning with Real-World Applications in Deep Learning]

Hidden layers

- We can
- Interpret

Universal approximation theorem

"Given a neural network with a single hidden layer of sufficient size, the network can model can approximate any function in $\mathbb{R}^p$ "

In other words:

- There is a 'true' function f relating the inputs X to outputs Y
- NN can approximate this function to arbitrary precision given sufficient layer size
- Required layer size can be extremely large and grow rapidly (e.g. exponential) with the dimensionality of the problem (p)

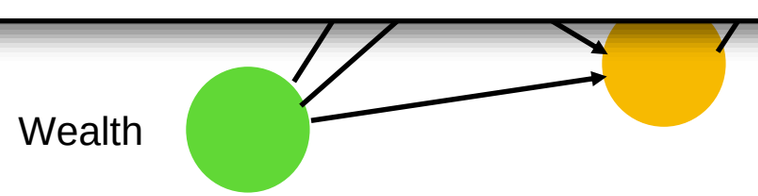

The components

- Neural networks are made from **neurons** and **edges**.
- A collection of neurons is a **layer**.
- The output of a previous layer is used as an input to the next layer.
- The input layer is x, the raw data, and the output layer is $\hat{y}$, the prediction.
- Anything in between is "**hidden**".
- Layers are usually represented as **vectors**.
- Edges are usually represented as **matrices - the weights**.
- The **edges** are **parameters** to be trained.

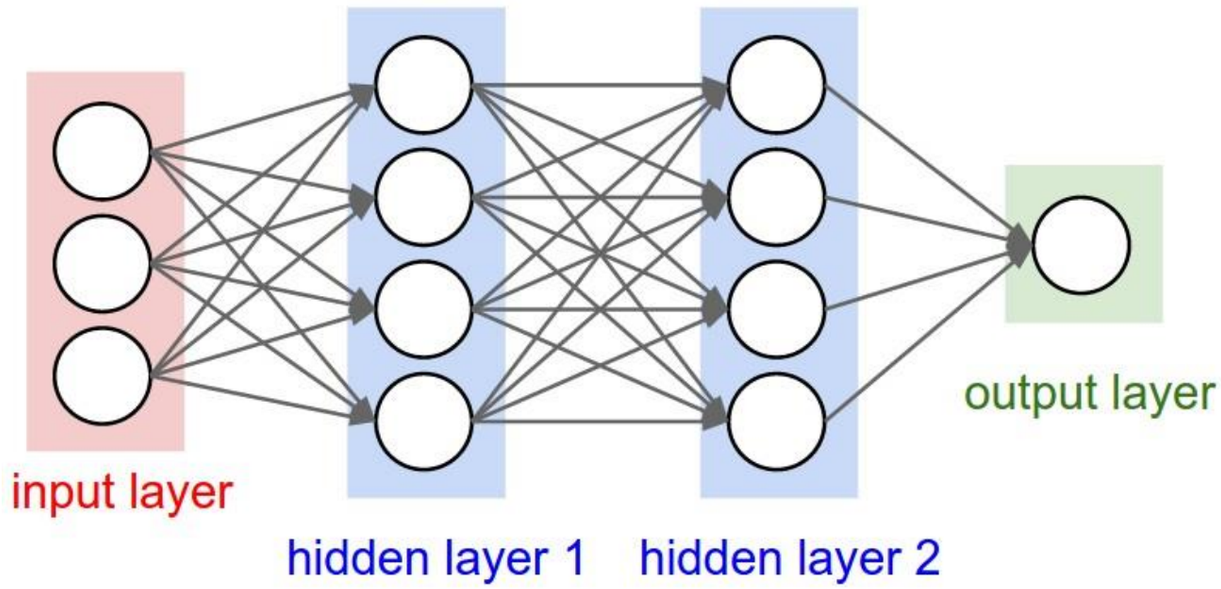

The components

- There are many **types** of layers, the previous example is a **dense** layer: all neurons in one layer are connected to all neurons in the next layer.

- Piling many dense layers, creates a Multi Layer Perceptron (MLP) and is very powerful, but expensive.

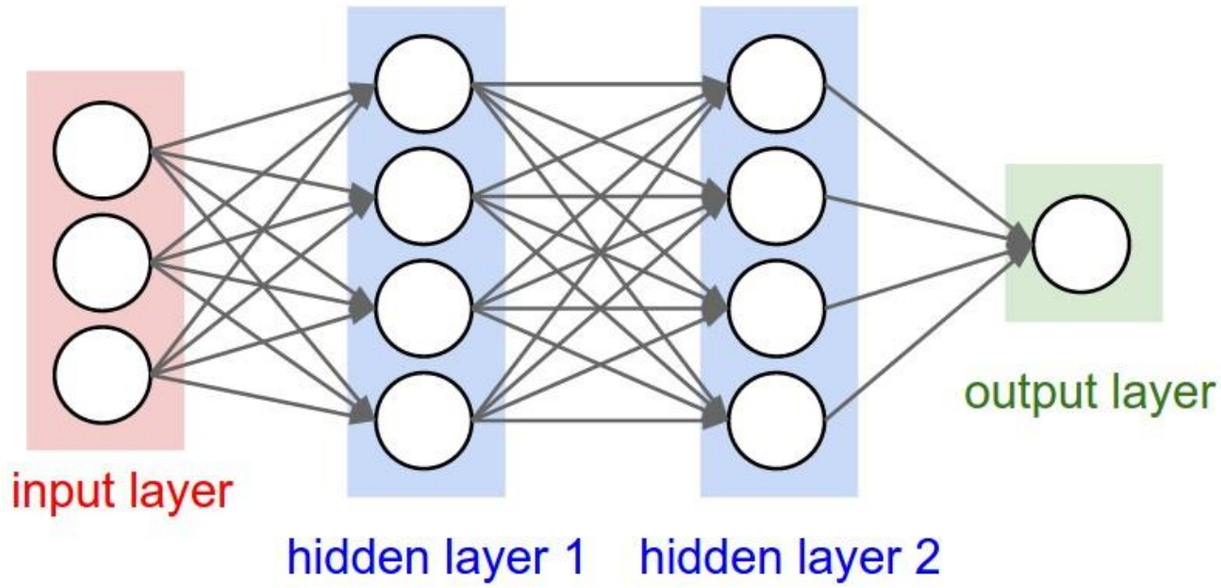

Neura[illegible]overs

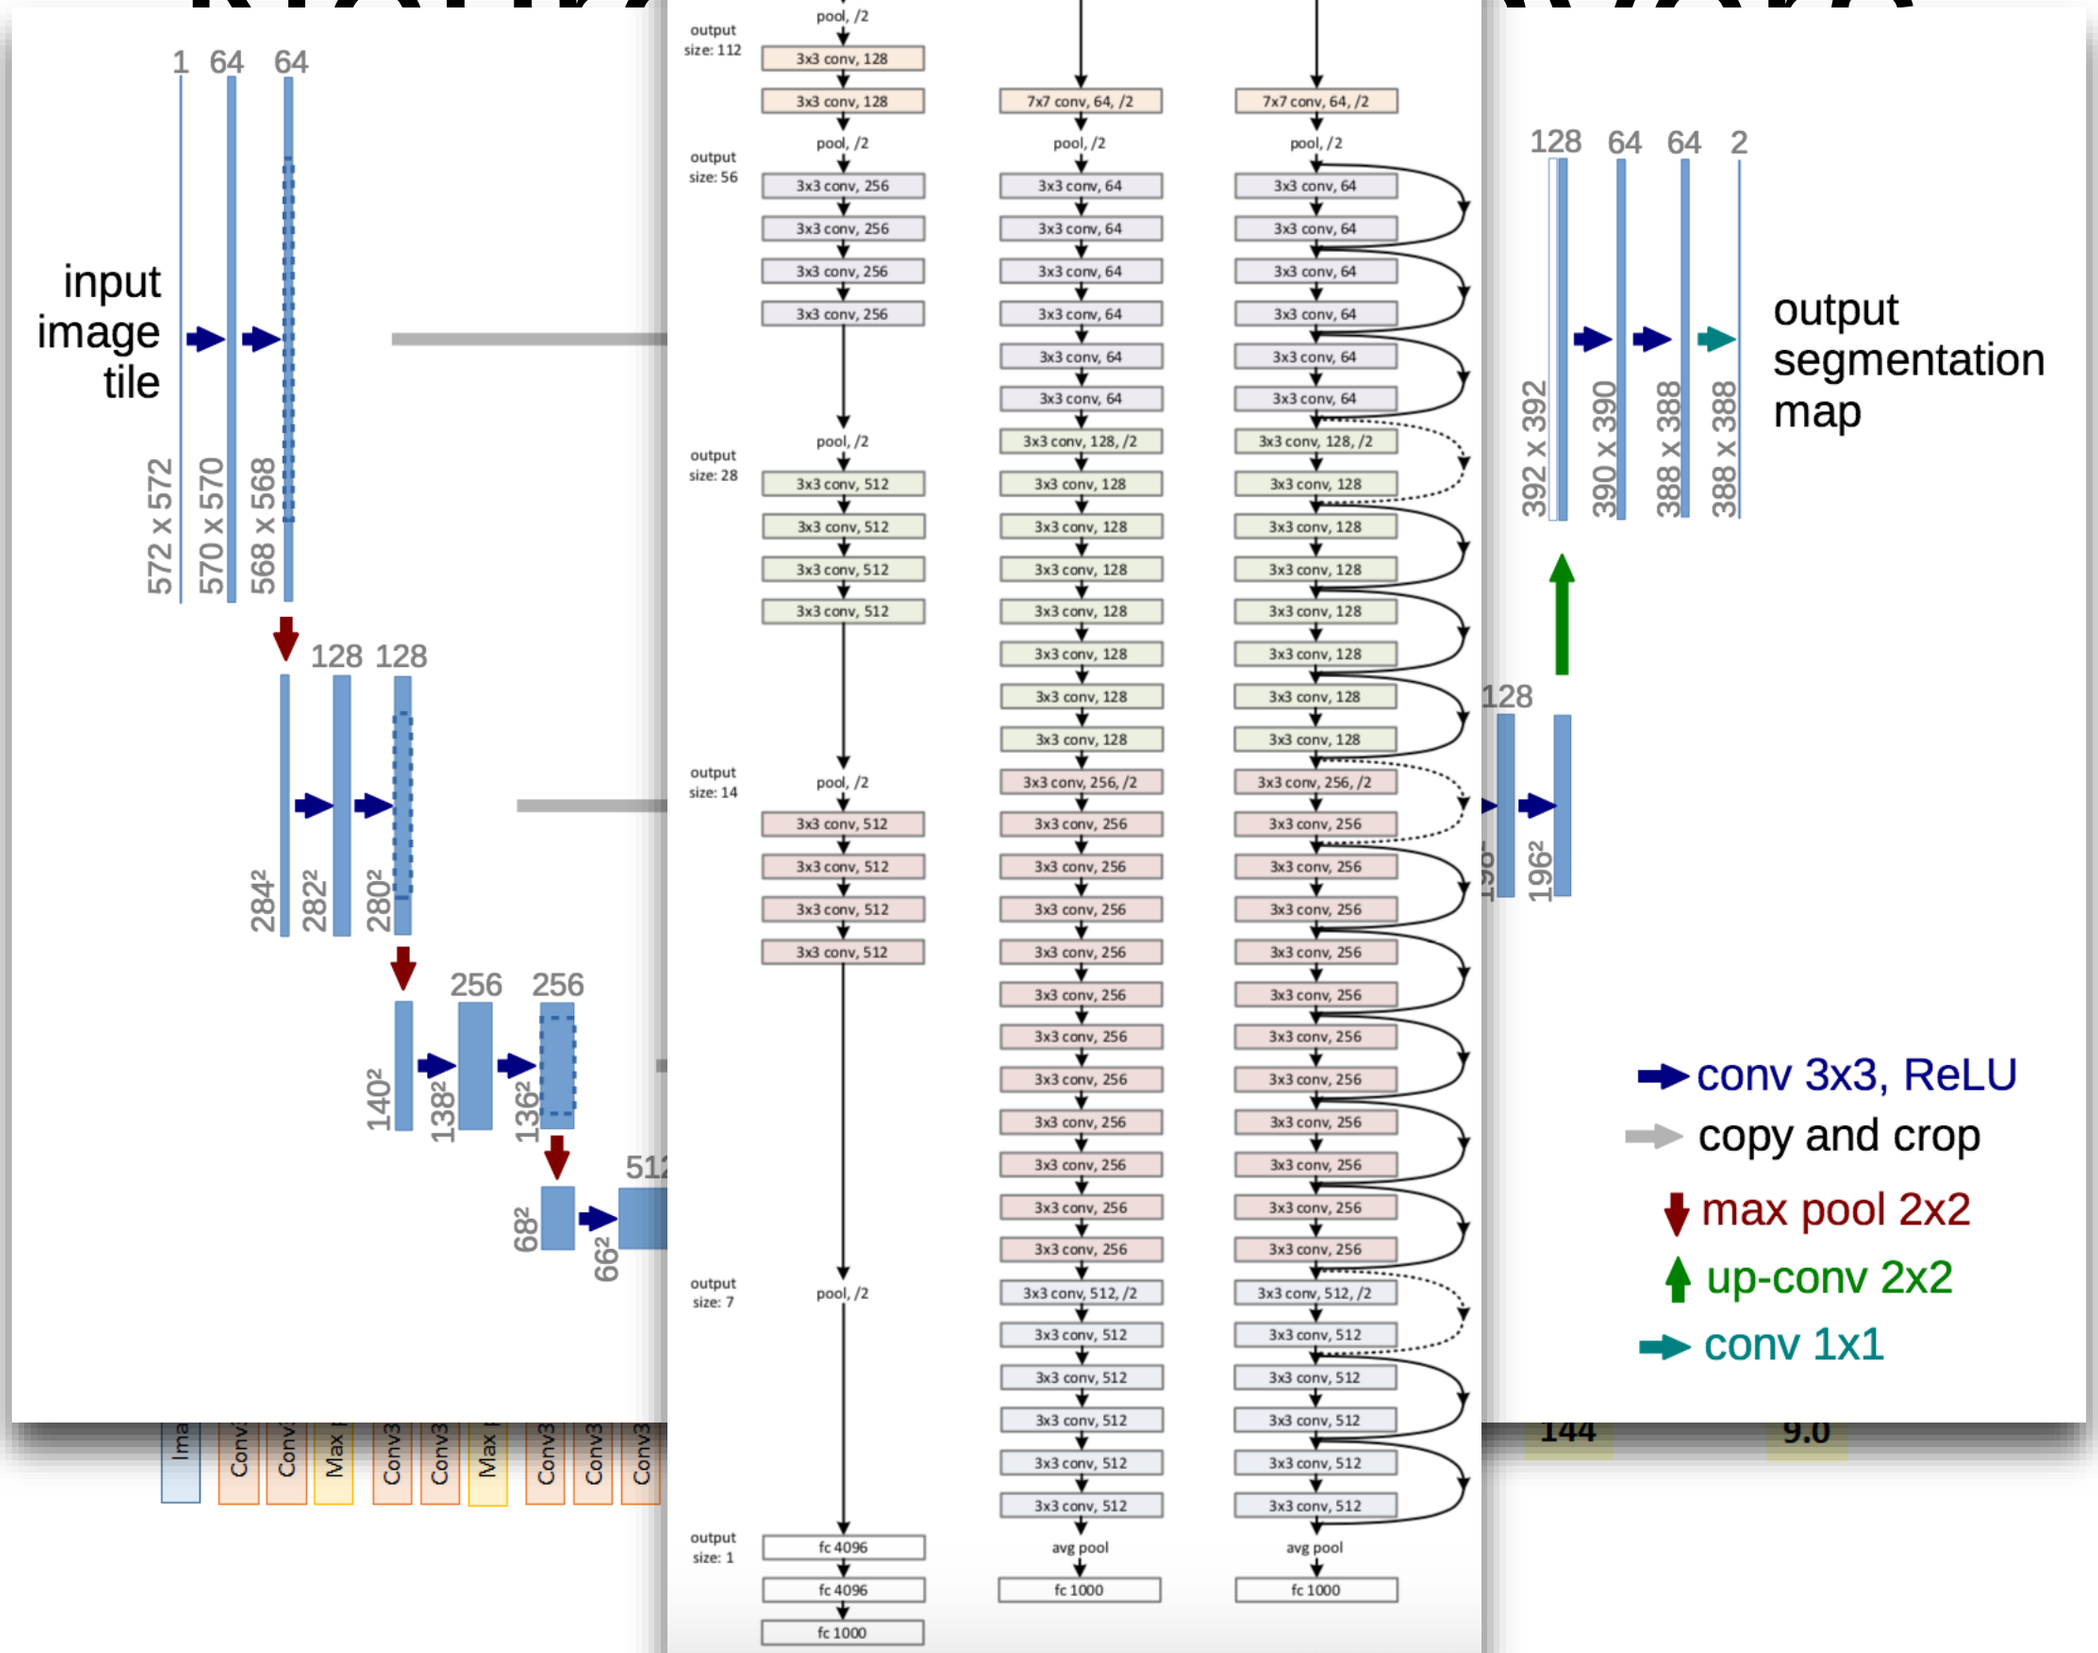
input image tile
1 64 64
572 x 572
570 x 570
568 x 568
128 128
284²
282²
280²
256 256
140²
138²
136²
512
68²
66²
128 64 64 2
392 x 392
390 x 390
388 x 388
388 x 388
output segmentation map
128
196²
conv 3x3, ReLU
copy and crop
max pool 2x2
up-conv 2x2
conv 1x1
VGG-19
34-layer plain
34-layer residual
image
output size: 224
3x3 conv, 64
pool, /2
output size: 112
3x3 conv, 128
output size: 56
3x3 conv, 256
output size: 28
3x3 conv, 512
output size: 14
output size: 7
output size: 1
fc 4096
fc 1000
7x7 conv, 64, /2
3x3 conv, 128, /2
3x3 conv, 256, /2
3x3 conv, 512, /2
avg pool

The XOR problem

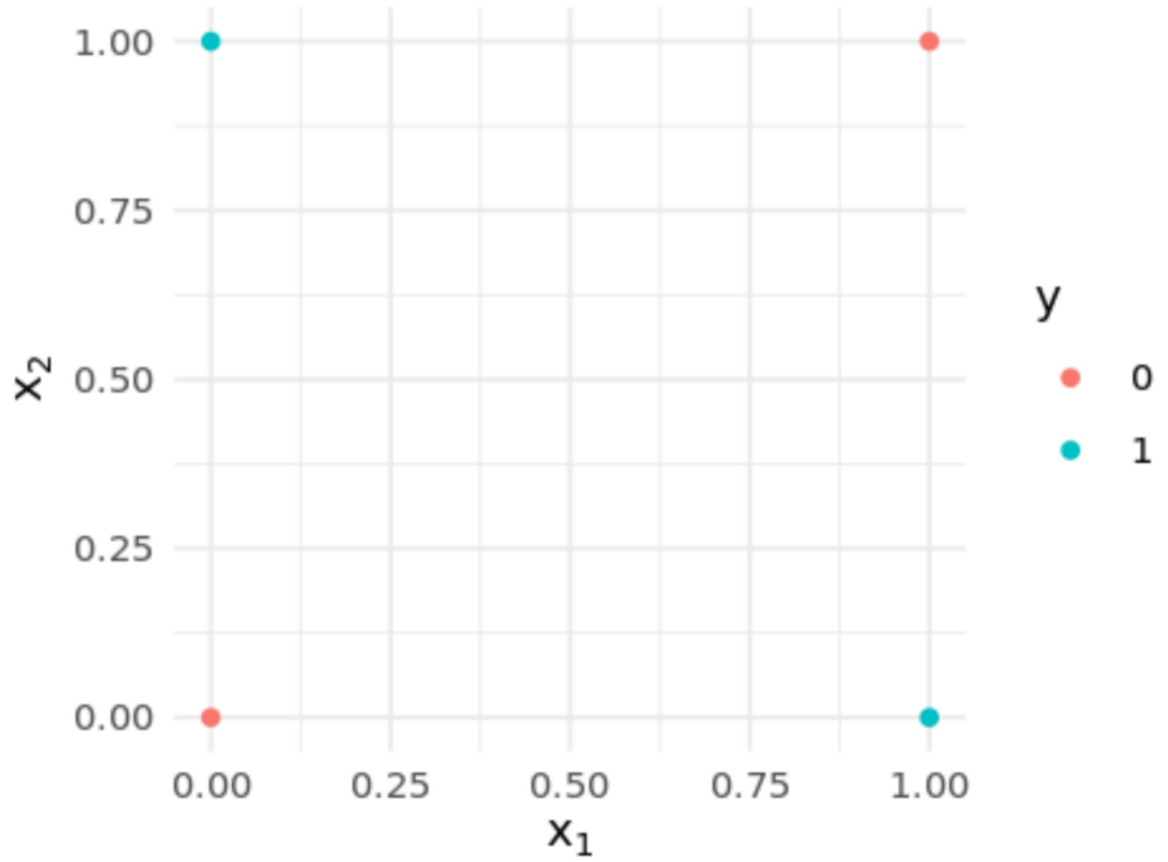

- The XOR problem is a famous problem (1960s) in deep learning.
- Early networks only used a single layer, and cannot solve XOR.
- By adding a single hidden layer we can implement XOR.

Prediction with a hidden layer

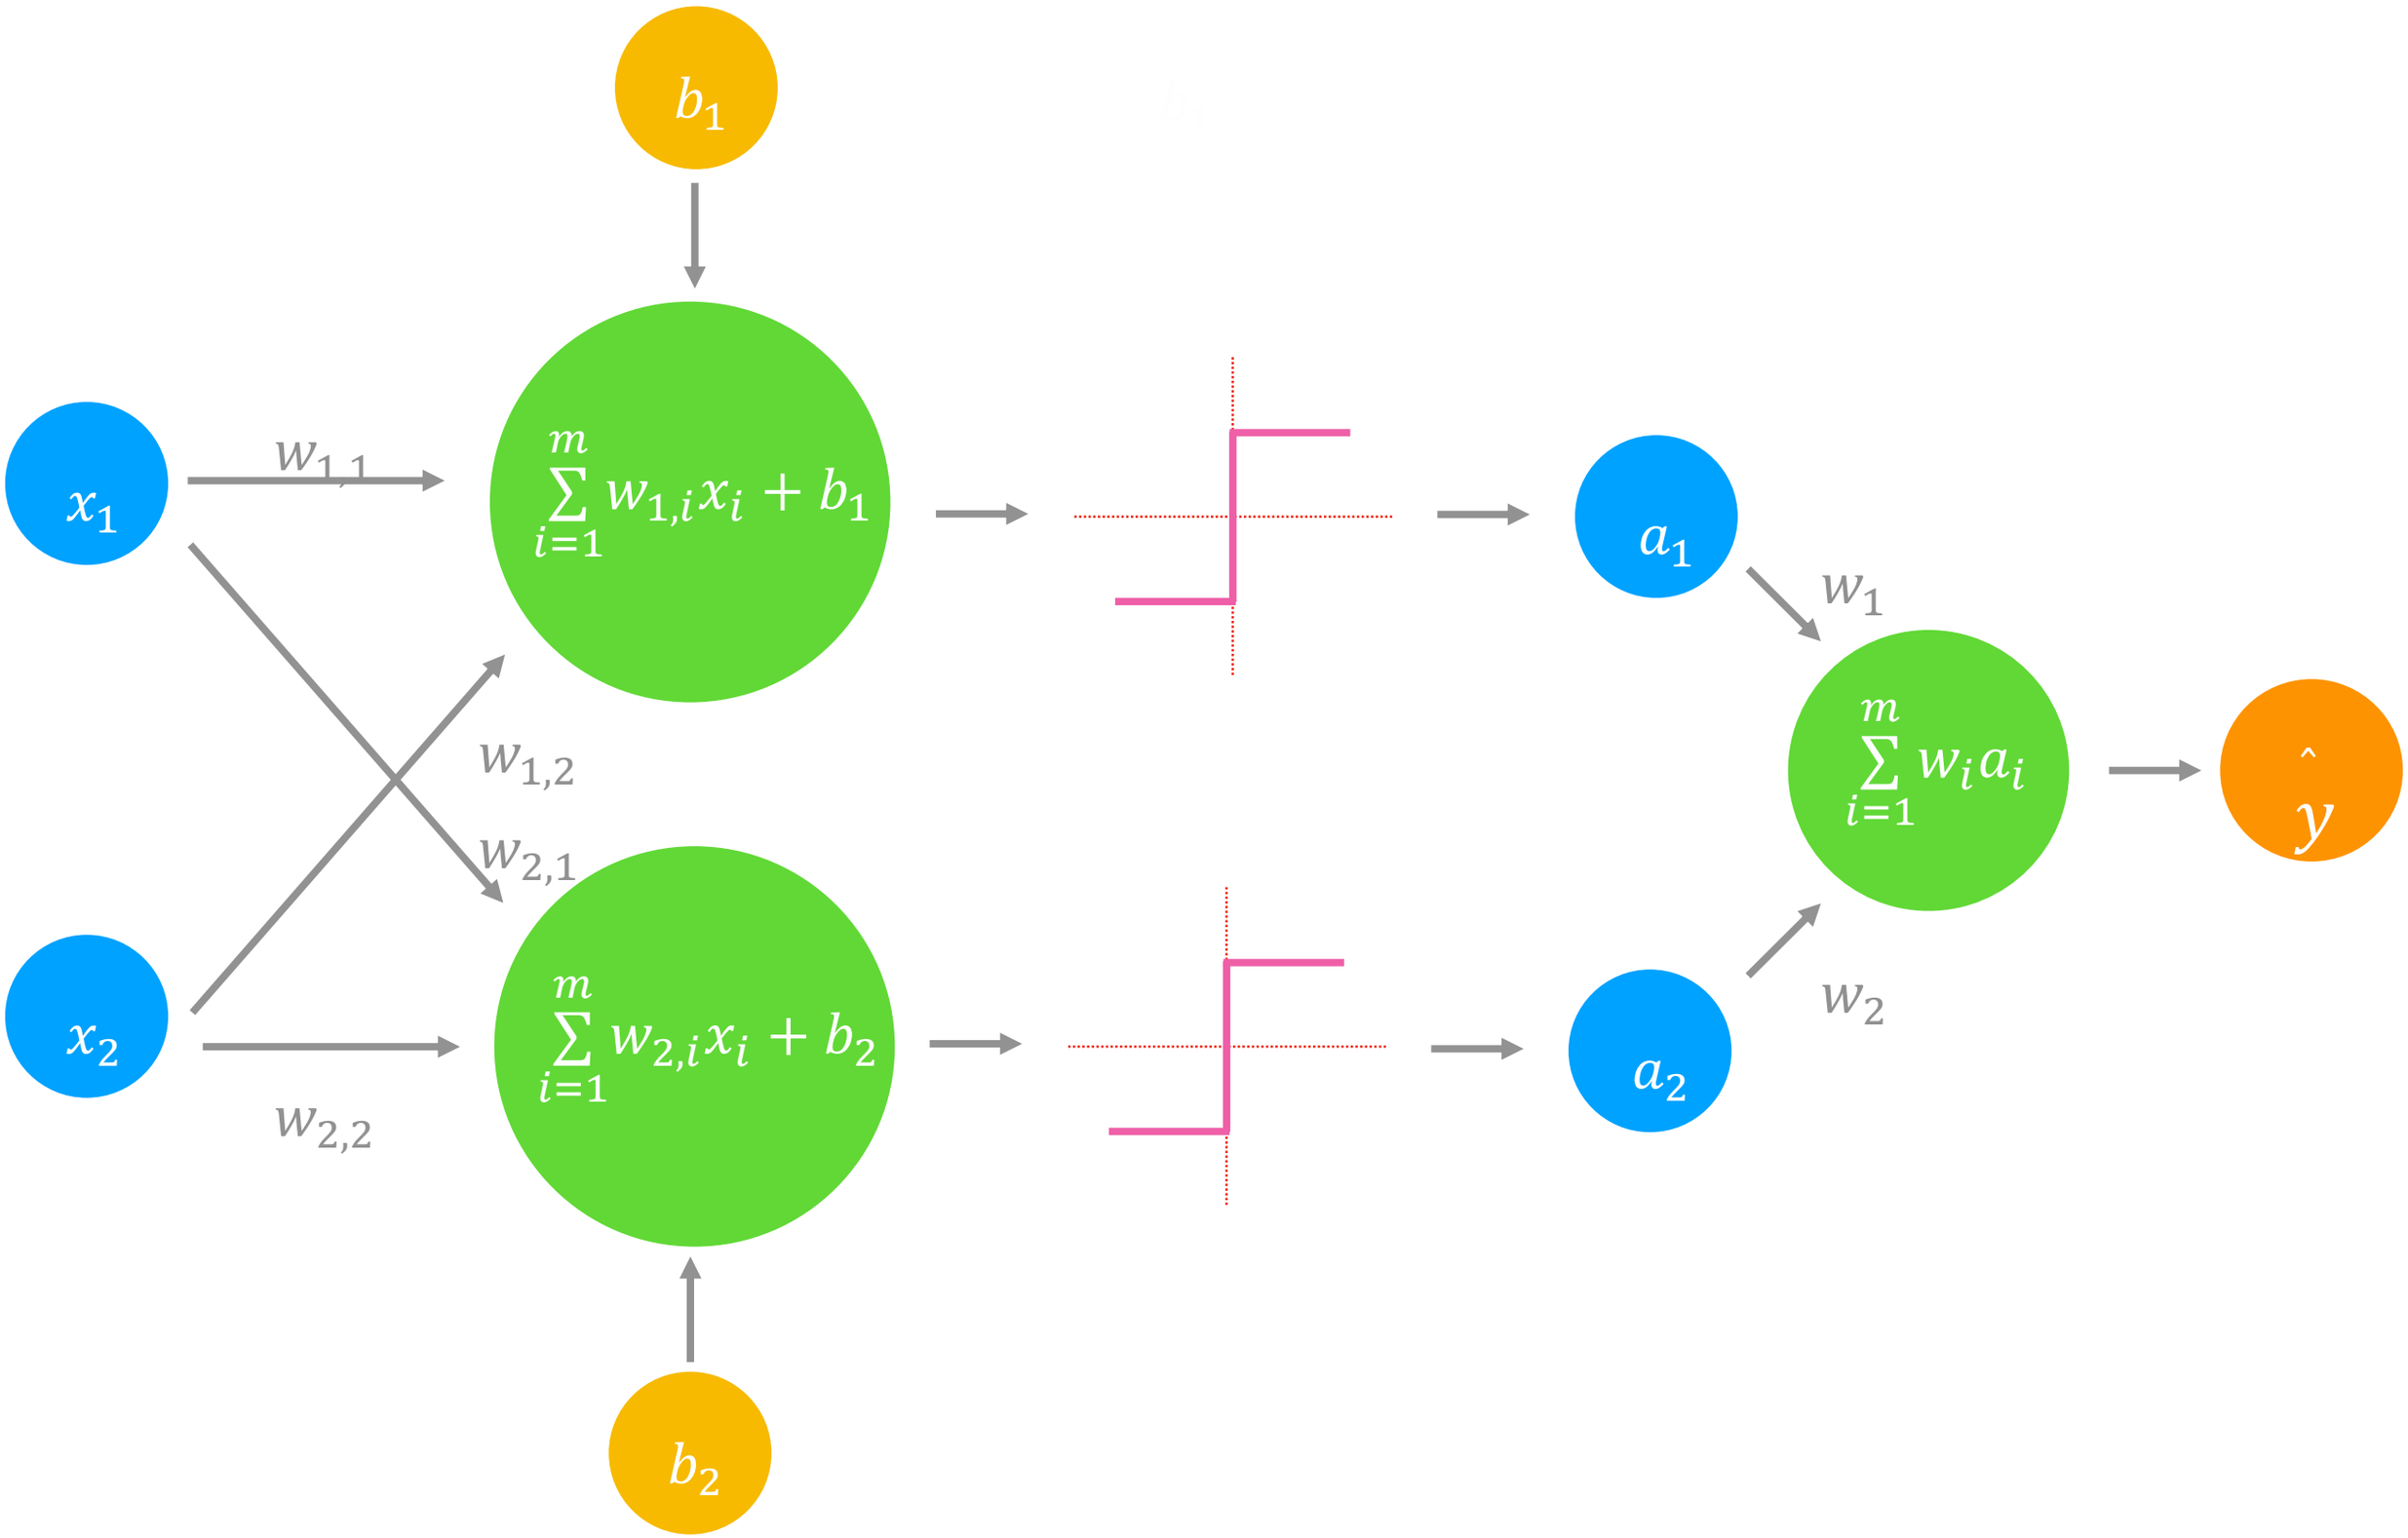

Hands-on

Go to https://https://jupyter.lisa.surfsara.nl:8000/

Notebook: 01b-network-from-scratch.ipynb

16:30-17:15

Summary

- Deep learning is a subfield of machine learning in which the model learns to represent the data using many levels of concept hierarchies.
- Deep learning has greatly improved performance in many tasks involving complex data.
- Neural networks are made from connected **neurons** which combine into **layers**. The network maps the **input x** to an **output y**.
- The environment @ https://jupyter.lisa.surfsara.nl:8000/
- If you have any questions, post them @ Canvas