

# Deep learning


Sequential modelling / RNNs

# Today's program

- 14:00-14:15 Recap
- 14:15-15:00 Machine learning tasks: regression / classification
- 15:00-15:45 Hands-on: multiclass Fashion MNIST
- 15:45-16:15 Break
- 16:15-16:45 Optimizers, regularization techniques
- 16:45-17:30 Hands-on: Regularization techniques on F-MNIST
- **17:30-18:00 Analyzing sequential data, RNNs**
- 18:00-19:00 Diner
- 19:00-19:45 Hands-on: Predicting future temperatures with an RNN
- 19:45-20:15 Types of RNNs: LSTM, GRU
- 20:15-21:00 Hands-on: creating sequences, temperature prediction with GRU-based RNN
- Time left: Improving RNNs: regularization, stacking, stateful and bi-directional RNNs
- Time left: Hands-on: Improved RNNs on temperature prediction

# Sequential data

- Data is sequential when the data has some order.
- The whole dataset can consist of a single order (sunspots) or many individual orders (sentences).



the cat sat on the mat .  
the book is open .



# Sequential data

in deep learning

- Machine translation
- Speech recognition
- Music generation
- Sentiment classification
- Video activity recognition
- ...

"Hoi, hoe gaat het?" → "Hi, how are you?"




→ "Hi, how are you?"

Some style / nothing →



"Hoi, hoe gaat het?" →




→ Running


# Sequential data

- In previous lectures our data have been made up from a single example.
- A single example can have many features.
  - Temperature, air pressure, etc.
- Now each example is made from a single sequence.
  - "Hi, hoe gaat het?"
- Each sequence has many examples.
  - "Hi", "hoe", "gaat", "het"
- That is, in each iteration we process a single sequence, many examples.


Classic example




Sequence example



# Sequential data

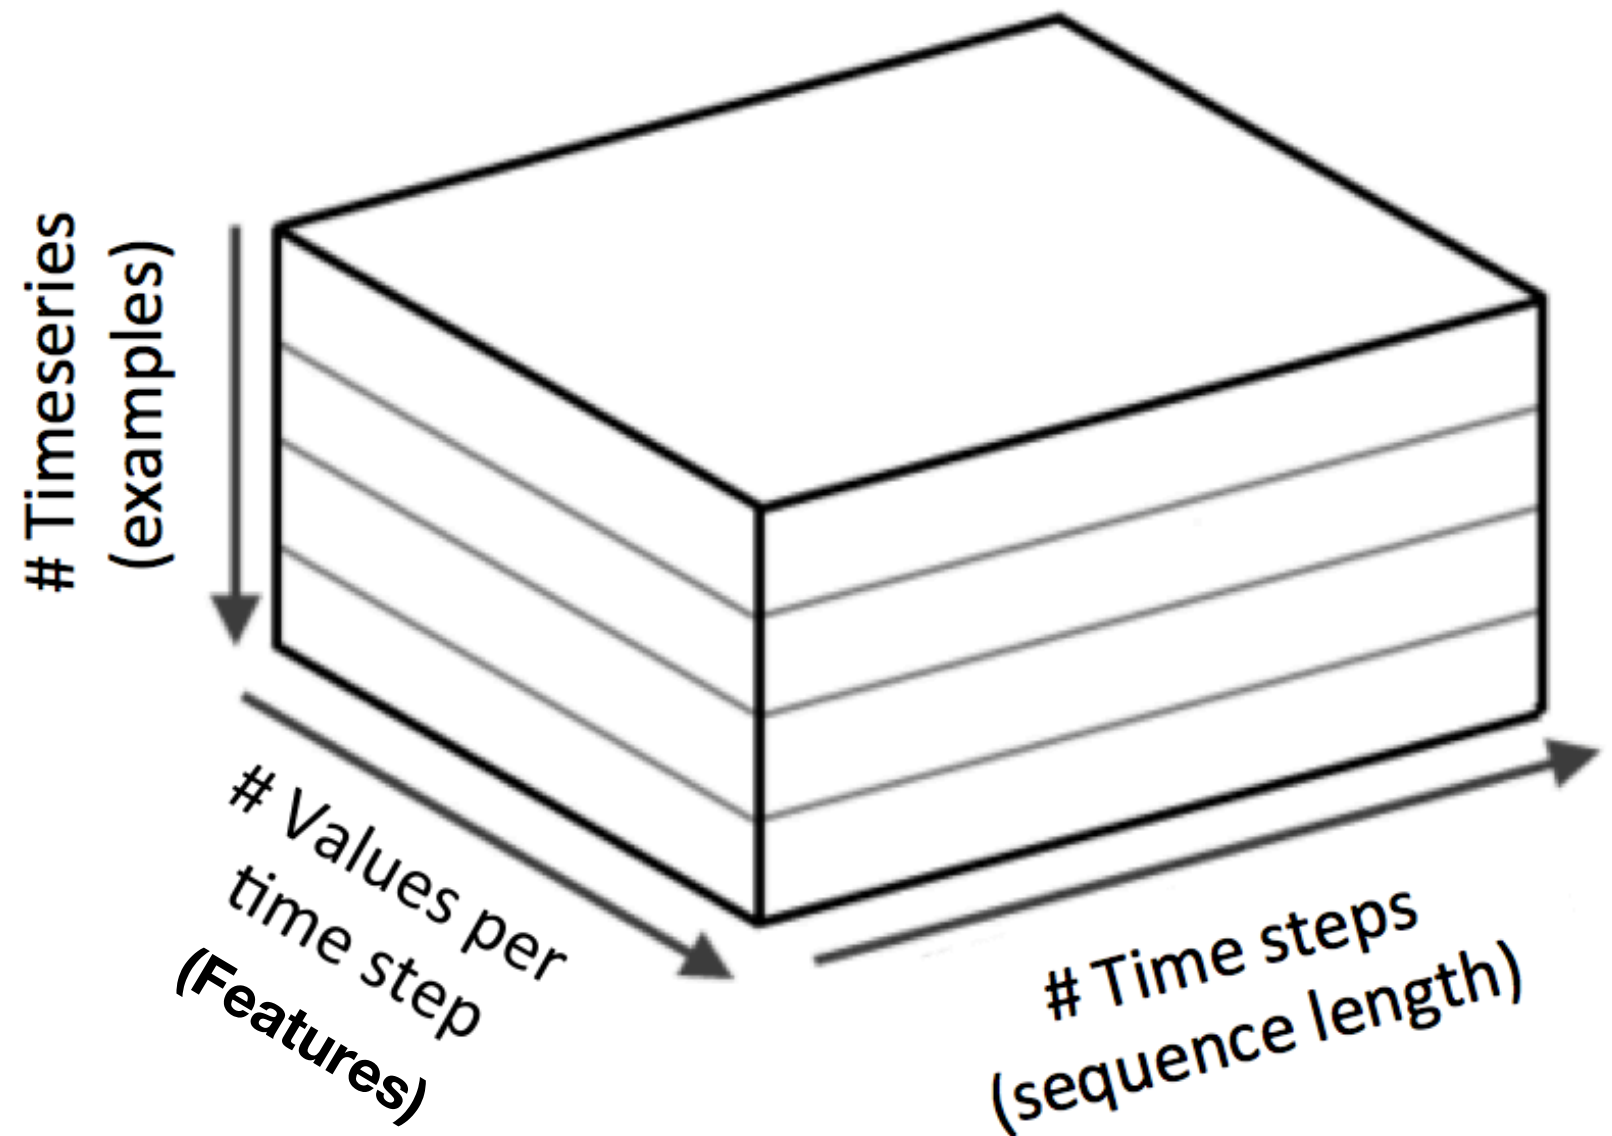


## Feed-Forward Network Data



Data = (examples, features)



## Recurrent Network Data



Data = (examples, sequence\_length, features)

# Which task?


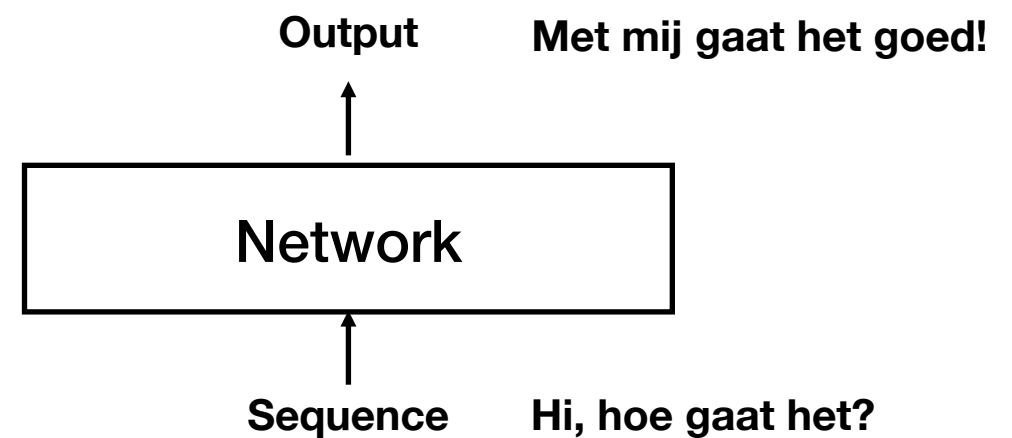
- Many-to-one
  - Regression
    - Tomorrows weather
    - Next word
  - Classification
    - Is this a question?
    - Assign rating based on review






# Which task?

- Many-to-many
  - Regression
    - Daily weather for next week
- Classification
  - Next sentence
  - Take a piece of music and compose the next section
  - Machine translation



# Which task?

- One-to-many
  - Classification
  - Automatic captioning




# How to model?

## Naive model

- Why not just a fully connected network?
- For this we will need VERY many parameters.
- The sequence length might vary between examples.
- We can often process each element independently and equally.
- "gaat" is the same word regardless of position in a sentence.

Naive sequence model




# How to model?

- What about a very simple model that just models the next word?

Dataset:

- The president of the United States is Donald Trump.
- Wherever Donald Trump goes, security is tight.
- A simple network could learn that 'Donald' is always followed by 'Trump'
- Pretty easy to learn, not so realistic...




# How to model?

- What about a very simple model that just models the next word?

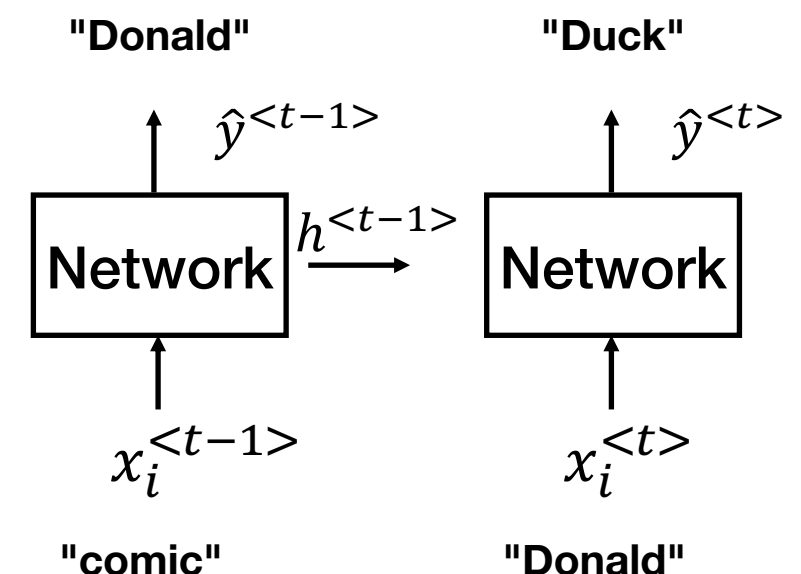
Dataset:

- The president of the United States is Donald Trump.
- Wherever Donald Trump goes, security is tight.
- The comic Donald Duck is very famous.
- Now what?




# How to model?

- Meaning depends on context!
- Recurrent networks have a memory (a 'state'),  $h^{<t-1>}$  in the figure below
- The next prediction  $\hat{y}^{<t>}$  depends not only on the input  $x_i^{<t>}$ , but also on the context 'remembered' by the network,  $h^{<t-1>}$ .



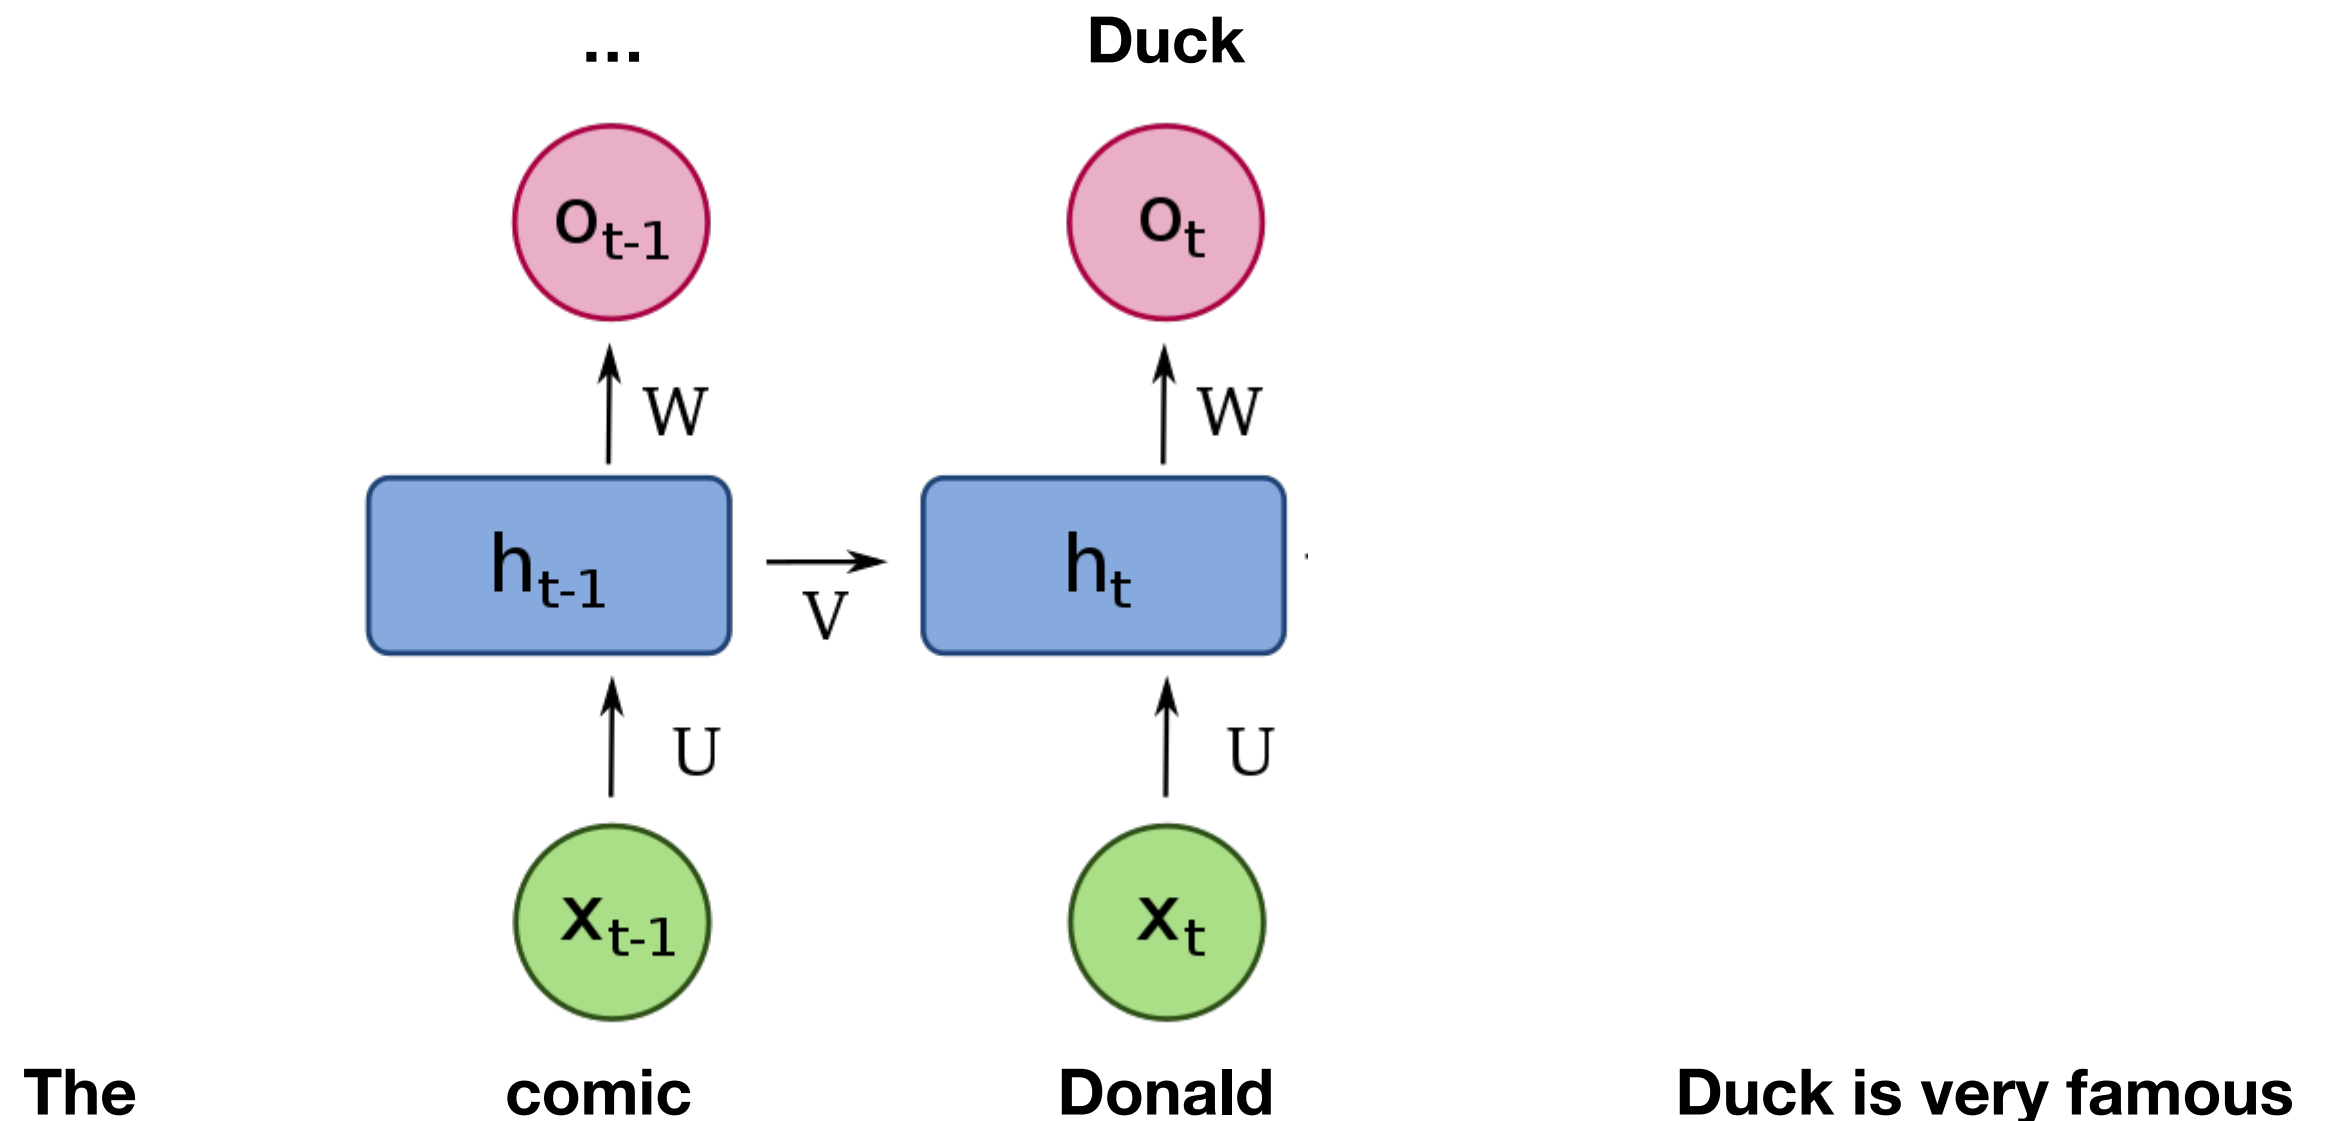
# Simple RNN

- $U$ ,  $V$ ,  $W$  are the weight matrices
- Weights are reused!



# Simple RNN

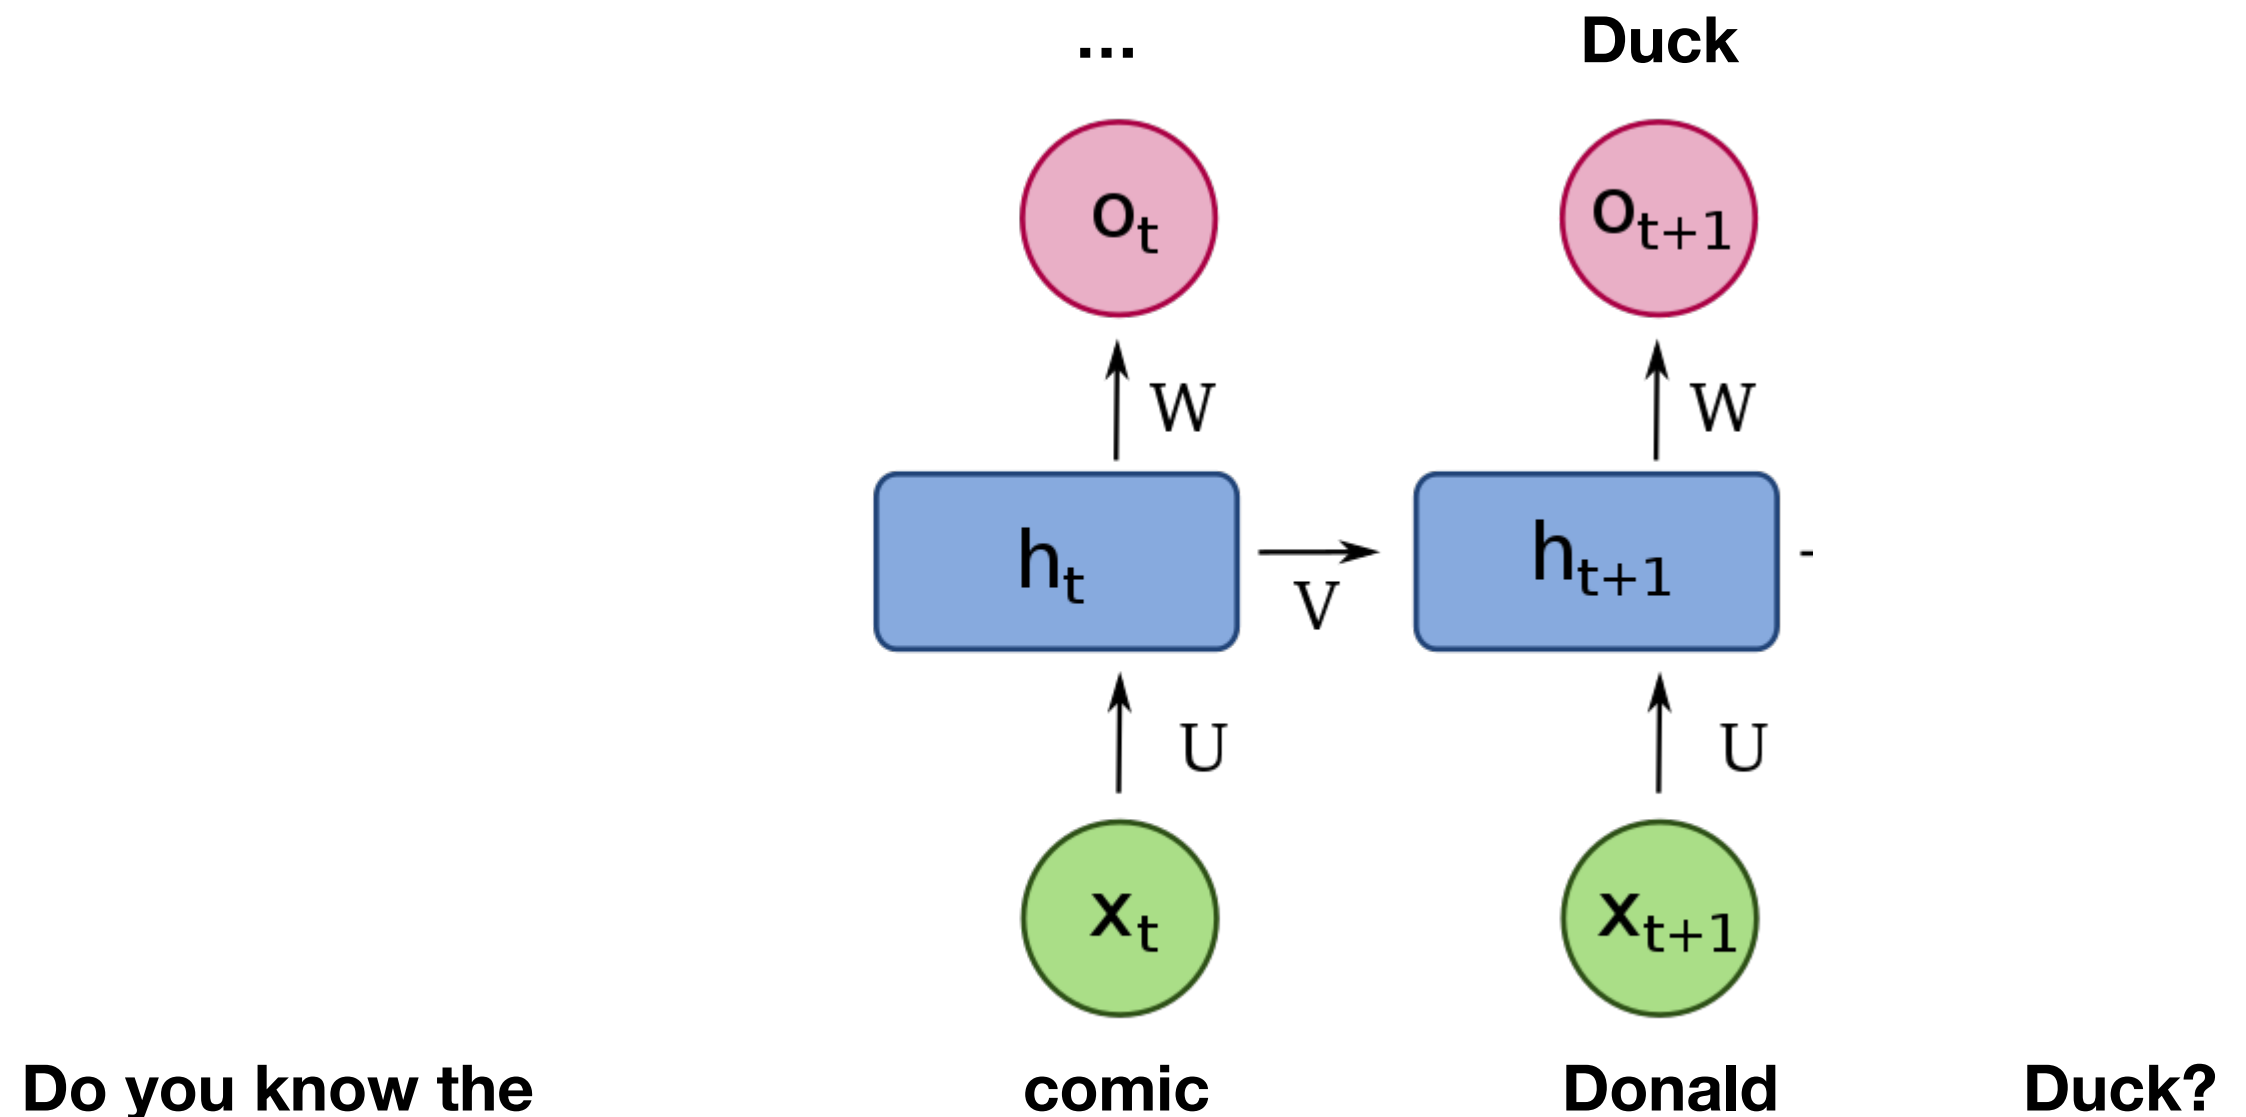
- Weights are reused: predictions the same, irrespective of where “comic Donald” occurs in a sequence





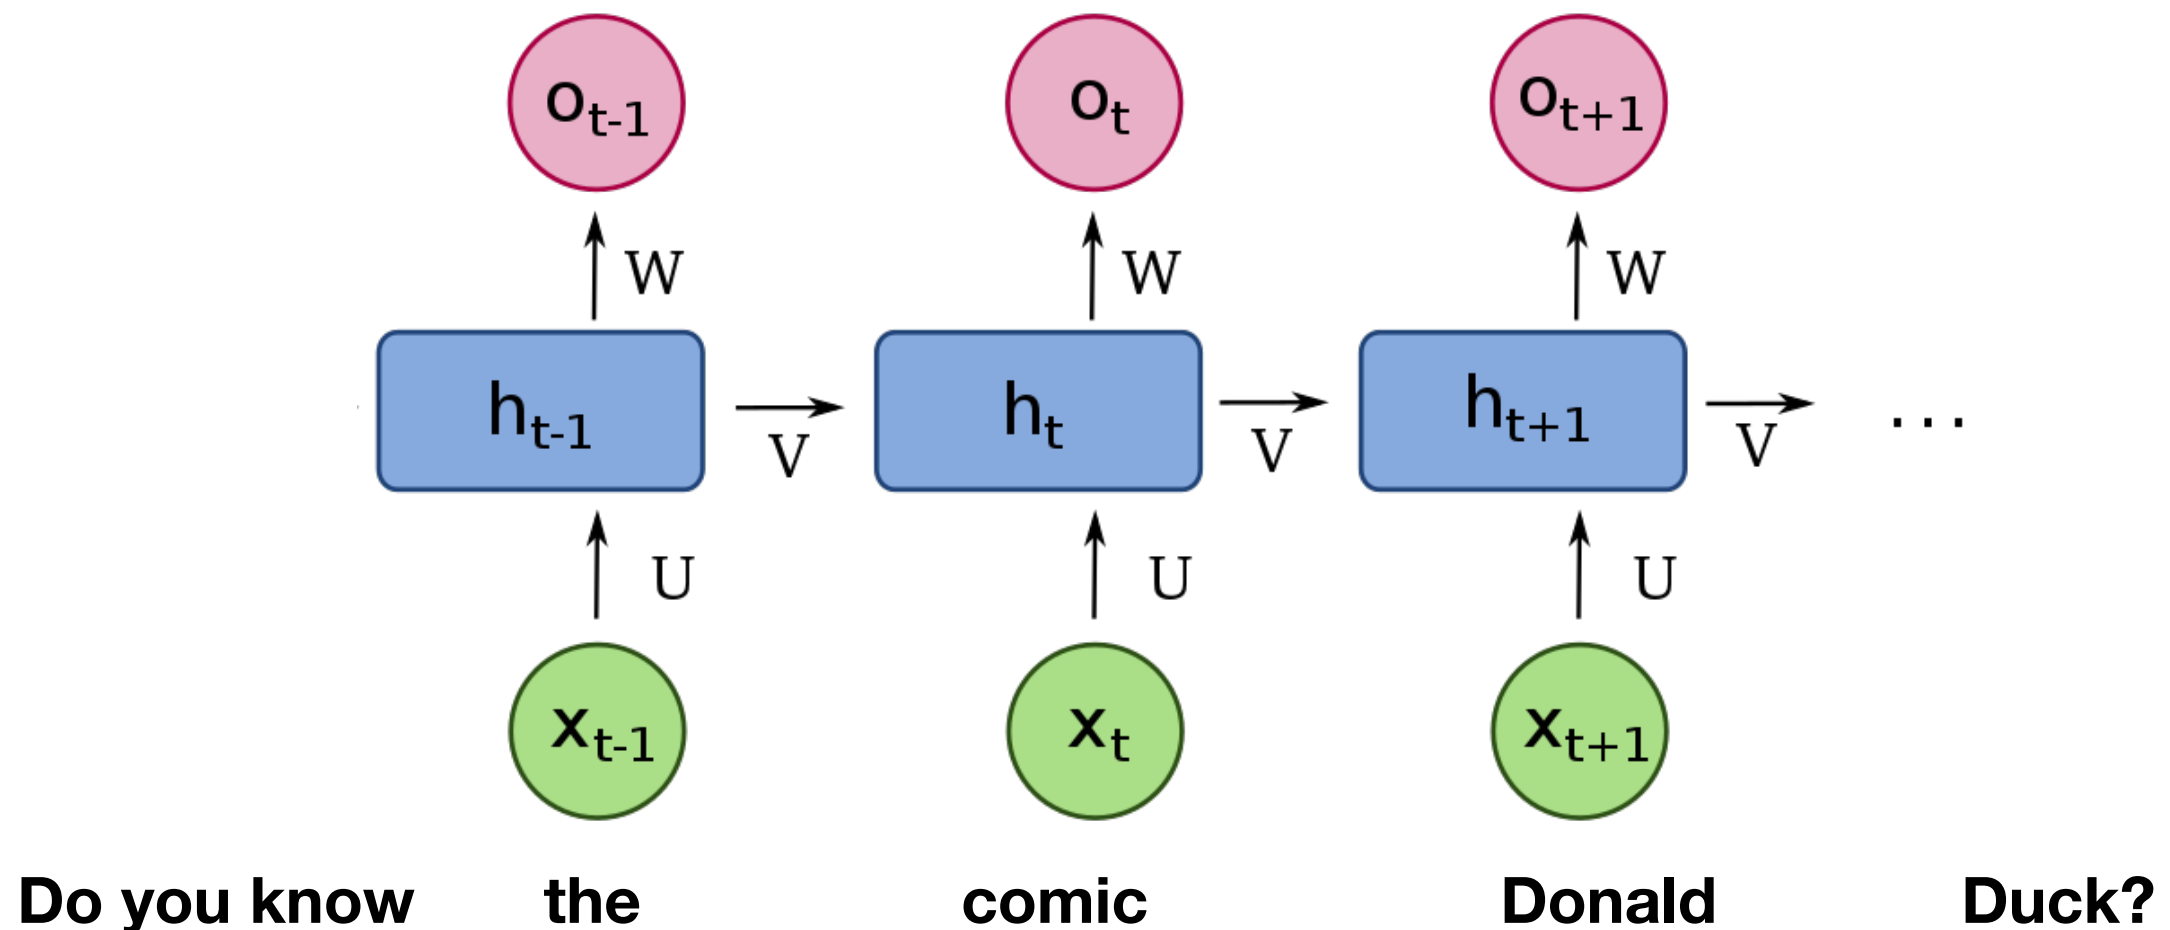
# Simple RNN

- Weights are reused: predictions the same, irrespective of where “comic Donald” occurs in a sequence



# Simple RNN

- Of course, this changes if we use longer sequence length (i.e. look at more context/history)
- Memory  $h_{t-1}$  *could* alter the prediction



# Simple RNN


- The math inside a simple RNN cell...

$$h_t = \sigma(U \cdot x_t + V \cdot h_{t-1} + b_h)$$

$$o_t = \sigma(W \cdot h_t + b_o)$$

←  
This is often  
considered the  
next dense layer

←  
This is often  
considered the  
RNN cell



- Compare to dense layer


$$o_t = \sigma(W \cdot x_t + b)$$

# Simple RNN

- How many parameters do we have?

$(n, m)(m)$	$(n, n)(n)$	$(n)$
$h_t = \sigma(U \cdot x_t + V \cdot h_{t-1} + b_h)$		
$o_t = \sigma(W \cdot h_t + b_o)$		
$(k, n)(n)$	$(k)$	

Layer\_simple\_rnn  
Layer\_dense




- Length of vector  $x_t$ : # features ( $m$ )
- Length of hidden state vector  $h_t$ : we choose it! (based on how complex a memory we think we need) ( $n$ )
- Length of vector  $o_t$ : # outputs (classes / numbers) per timestep ( $k$ )
- It does NOT depend on sequence length! (because weights are 'recycled')

# Simple RNN

- How many parameters do we have?

$(n, m)(m)$	$(n, n)(n)$	$(n)$	
$h_t = \sigma(U \cdot x_t + V \cdot h_{t-1} + b_h)$			Layer_simple_rnn
$o_t = \sigma(W \cdot h_t + b_o)$			Layer_dense
$(k, n)(n)$	$(k)$		



- Sum up all weights and biases:

$$\#params = (n \cdot m + n \cdot n + n) + (k \cdot n + k)$$

# Simple RNN

- How many parameters do we have?

$$params = (n \cdot m + n \cdot n + n) + (k \cdot n + k)$$

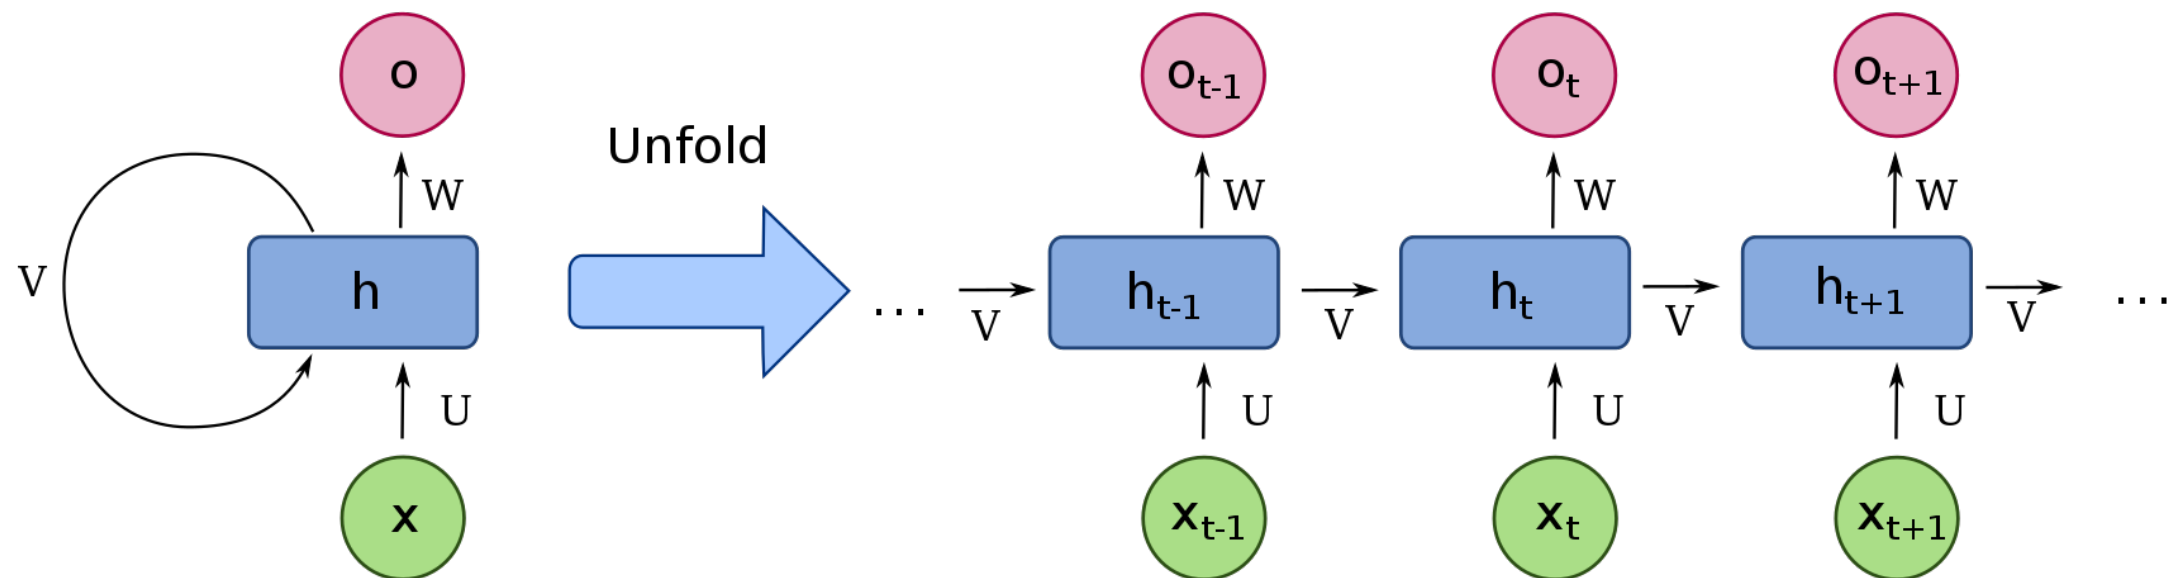
Example: trying to predict temperature for tomorrow

- 14 features (e.g. pressure, temperature, wind velocity, etc)
- 16 hidden states (we choose it!)
- 1 output per timepoint (only temperature)

$$\#params = (16 \cdot 14 + 16 \cdot 16 + 16) + (16 \cdot 1 + 1) = 496 + 17 = 513$$


# Unfolding an RNN

- Sometimes, literature shows RNN as a loop. We have looked at ‘unfolded’ RNNs. Same thing, just different visual representation.



# Notebook

- The Jena weather dataset.
- A long sequence of weather measurements.
- Each measurement consists of 15 variables.
- We see a time-dependent pattern in the data.





# Today's program

- 14:00-14:15 Recap
- 14:15-15:00 Machine learning tasks: regression / classification
- 15:00-15:45 Hands-on: multiclass Fashion MNIST
- 15:45-16:15 Break
- 16:15-16:45 Optimizers, regularization techniques
- 16:45-17:30 Hands-on: Regularization techniques on F-MNIST
- 17:30-18:00 Analyzing sequential data, RNNs
- **18:00-19:00 Diner**
- 19:00-19:45 Hands-on: Predicting future temperatures with an RNN
- 19:45-20:15 Types of RNNs: LSTM, GRU
- 20:15-21:00 Hands-on: creating sequences, temperature prediction with GRU-based RNN
- Time left: Improving RNNs: regularization, stacking, stateful and bi-directional RNNs
- Time left: Hands-on: Improved RNNs on temperature prediction

# Hands-on



Go to <https://jupyter.lisa.surfsara.nl:8000/>

Or <https://dba.projects.sda.surfsara.nl/>

Notebook: `04a-time-series-prediction.ipynb`


19:00-19:45

# Hands-on recap

## sharing parameters

- Each time-step we use the **same network** (same weight matrices), it just gets **different "memory"** (hidden state) passed to it.
- Sharing parameters **reduces** the number of **parameters** in the model.
- Reduces model complexity and risk of overfitting. This is the value of RNNs over fully connected!
- Assumption behind RNN: meaning of an input depends on
  - Input
  - Current hidden state
  - *Not* on location in the sequence

We process "the" in the same way.



```
layer_simple_rnn(units = 12, input_shape = c(sequence_length, features))
```

# Today's program

- 14:00-14:15 Recap
- 14:15-15:00 Machine learning tasks: regression / classification
- 15:00-15:45 Hands-on: multiclass Fashion MNIST
- 15:45-16:15 Break
- 16:15-16:45 Optimizers, regularization techniques
- 16:45-17:30 Hands-on: Regularization techniques on F-MNIST
- 17:30-18:00 Analyzing sequential data, RNNs
- 18:00-19:00 Diner
- 19:00-19:45 Hands-on: Predicting future temperatures with an RNN
- **19:45-20:15 Types of RNNs: LSTM, GRU**
- 20:15-21:00 Hands-on: creating sequences, temperature prediction with GRU-based RNN
- Time left: Improving RNNs: regularization, stacking, stateful and bi-directional RNNs
- Time left: Hands-on: Improved RNNs on temperature prediction

# Sequential data

Long term dependencies

- After each time-step we store some information.
- When we train RNNs we are training them to do two things.
  - Store the correct information between time-steps. This is **hard**.
  - Map the stored information to solve the task. This is **easy**.



# Sequential data

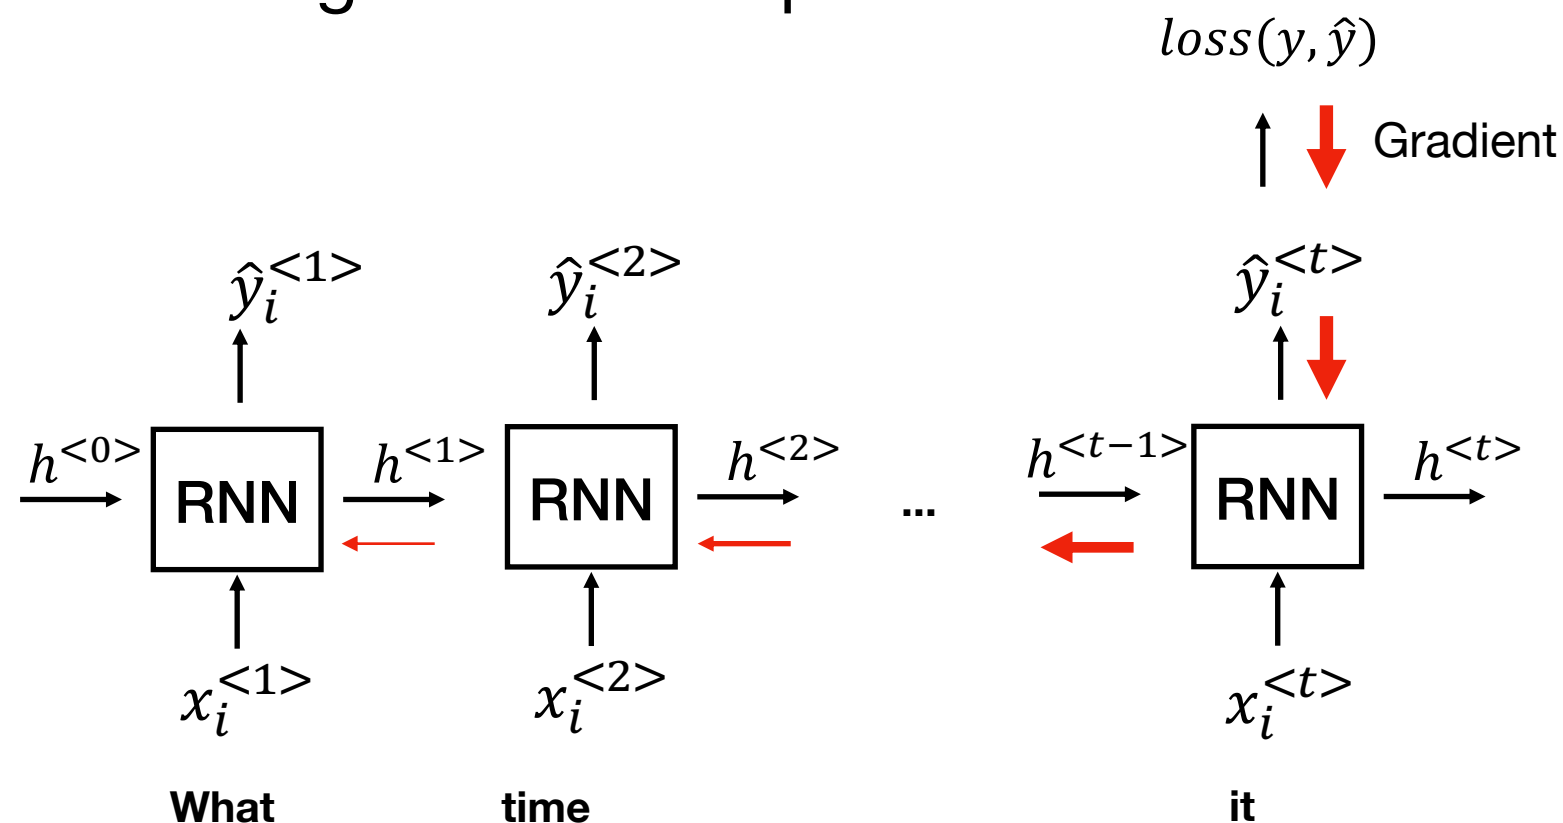
Long term dependencies

- Why is it hard?
- RNNs are trying to learn to **represent sequences by remembering what they contain.**
  - In question detection we want to remember if we saw "what", ... .
- We learn to represent the sequence in order to solve the task at hand.
  - At **start** we are doing **poorly** (random weights) and we see almost no indication that "what" was used previously.
  - We want to update our RNN cell so that next time we remember when we see "what".

# Sequential data

Long term dependencies

- If the sequence is long, little information is passed all the way to the end so a small error signal is sent back for the word "what".
- Small error signal = small updates




# Sequential data

Long term dependencies

- There is the problem of **long term dependencies**, which is because the **gradient vanishes** (the error signal).
- This is a **general problem** in neural networks trained with gradient descent, but very tangible in RNNs due to their depth.

$$\text{loss}(\text{Pred}, \text{Truth}) = \mathbf{E}$$



$$k(x) = f(g(x))$$

$$k'(x) = f'(g(x)) \cdot g'(x) \text{ of}$$


$$\frac{dk}{dx} = \frac{df}{dg} \cdot \frac{dg}{dx}$$



# Residual connections

Long term dependencies

- The solution to this problem in general are **residual connections**, (ResNet, 2015).
- We add connections which **bypass non-linear activations** (or go through fewer).
- This allows the error signal to flow directly to earlier layers.



# LSTMs and GRUs


Long term dependencies

- The solution to the vanishing gradient problem in RNNs was a different implementation of the RNN cell.
  - LSTM (1997)
  - GRU (2014)
- They are more complex and expensive but are able to deal better with long term dependencies.
- LSTM is heavier than GRU.


```
layer_gru(units = 10)
```

```
layer_lstm(units = 10)
```

## LSTM



## GRU



sigmoid




tanh



pointwise  
multiplication



pointwise  
addition



vector  
concatenation

# Generating sequences

We want to break this long sequence into many sequences

(1, ... ) (2, ... ) (3, ... ) (4, ... ) (5, ... ) (6, ... ) (7, ... ) (8, ... ) (9, ... ) (10, ... ) (11, ... ) (12, ... ) (13, ... ) (14, ... )

Reshape approach - sequence length = 7

(1, ... ) (2, ... ) (3, ... ) (4, ... ) (5, ... ) (6, ... ) (7, ... ) (8, ... ) (9, ... ) (10, ... ) (11, ... ) (12, ... ) (13, ... ) (14, ... )



1st example

(1, ... ) (2, ... ) (3, ... ) (4, ... ) (5, ... ) (6, ... ) (7, ... )

2nd example

(8, ... ) (9, ... ) (10, ... ) (11, ... ) (12, ... ) (13, ... ) (14, ... )

# Generating sequences

We want to break this long sequence into many sequences

(1, ... ) (2, ... ) (3, ... ) (4, ... ) (5, ... ) (6, ... ) (7, ... ) (8, ... ) (9, ... ) (10, ... ) (11, ... ) (12, ... ) (13, ... ) (14, ... )

Shift approach, using shift = 2, sequence length = 7

(1, ... ) (2, ... ) (3, ... ) (4, ... ) (5, ... ) (6, ... ) (7, ... ) (8, ... ) (9, ... ) (10, ... ) (11, ... ) (12, ... ) (13, ... ) (14, ... )



1st example

(1, ... ) (2, ... ) (3, ... ) (4, ... ) (5, ... ) (6, ... ) (7, ... )

2nd example

(3, ... ) (4, ... ) (5, ... ) (6, ... ) (7, ... ) (8, ... ) (9, ... )

3rd example

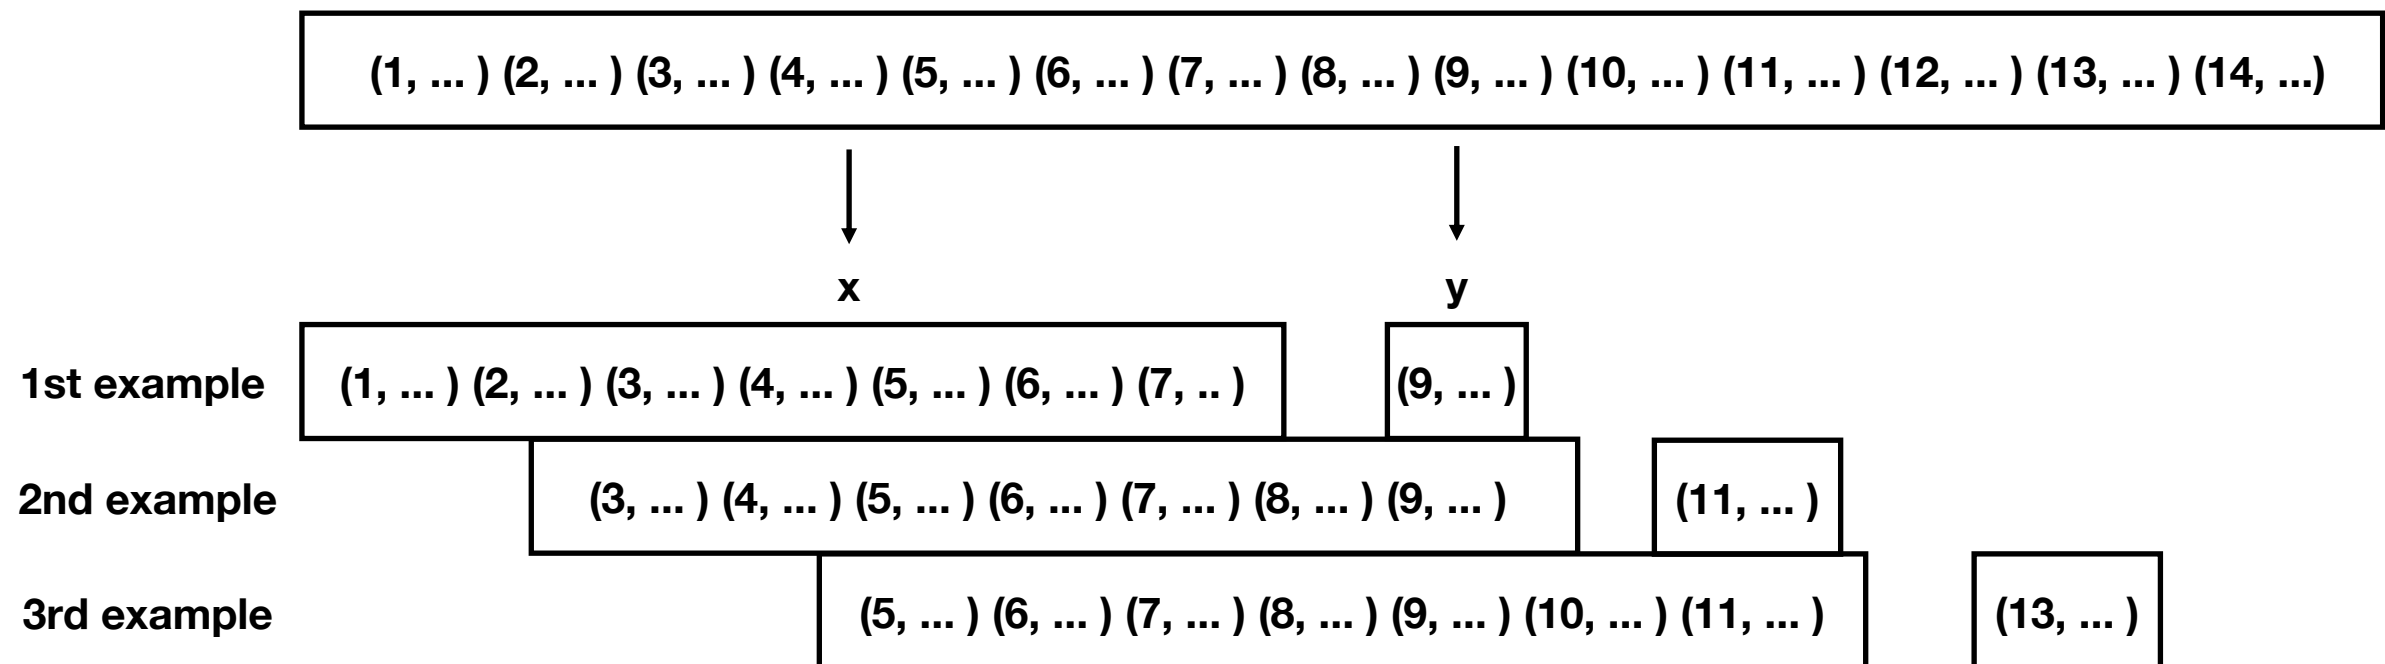
(5, ... ) (6, ... ) (7, ... ) (8, ... ) (9, ... ) (10, ... ) (11, ... )

4th example

(7, ... ) (8, ... ) (9, ... ) (10, ... ) (11, ... ) (12, ... ) (13, ... )

# Generating sequences

Shift approach, using shift = 2, sequence length = 7, target shift = 1



# Hands-on



Go to <https://jupyter.lisa.surfsara.nl:8000/>

Or <https://dba.projects.sda.surfsara.nl/>

Notebook: 05a-rnns.ipynb

20:15-21:00

# Today's program

- 14:00-14:15 Recap
- 14:15-15:00 Machine learning tasks: regression / classification
- 15:00-15:45 Hands-on: multiclass Fashion MNIST
- 15:45-16:15 Break
- 16:15-16:45 Optimizers, regularization techniques
- 16:45-17:30 Hands-on: Regularization techniques on F-MNIST
- 17:30-18:00 Analyzing sequential data, RNNs
- 18:00-19:00 Diner
- 19:00-19:45 Hands-on: Predicting future temperatures with an RNN
- 19:45-20:15 Types of RNNs: LSTM, GRU
- 20:15-21:00 Hands-on: creating sequences, temperature prediction with GRU-based RNN
- **Time left: Improving RNNs: regularization, stacking, stateful and bi-directional RNNs**
- **Time left: Hands-on: Improved RNNs on temperature prediction**



# Improving RNNs

- Regularisation
  - L1/L2
  - Dropout, recurrent dropout
- Improving RNNs
  - Stacking
  - Stateful
  - Bi-directional



# Improving RNNs

## L2/L1 regularisation

- Just like with normal dense layers.
- we add L2/L1 regularisation to the weights learnt in the RNN cell.

```
layer_gru(units = 10, kernel_regularizer = regularizer_l2(l = 0.001))
```

```
layer_gru(units = 10, kernel_regularizer = regularizer_l1(l = 0.001))
```



Red line: identical predictions of the NN (too many degrees of freedom)  
Red point: solution for  $(w_1, w_2)$  preferred by each norm.

# Improving RNNs


## Dropout

- In RNNs we consider dropouts in two locations.

- **Input**

- **Recurrent dropout**

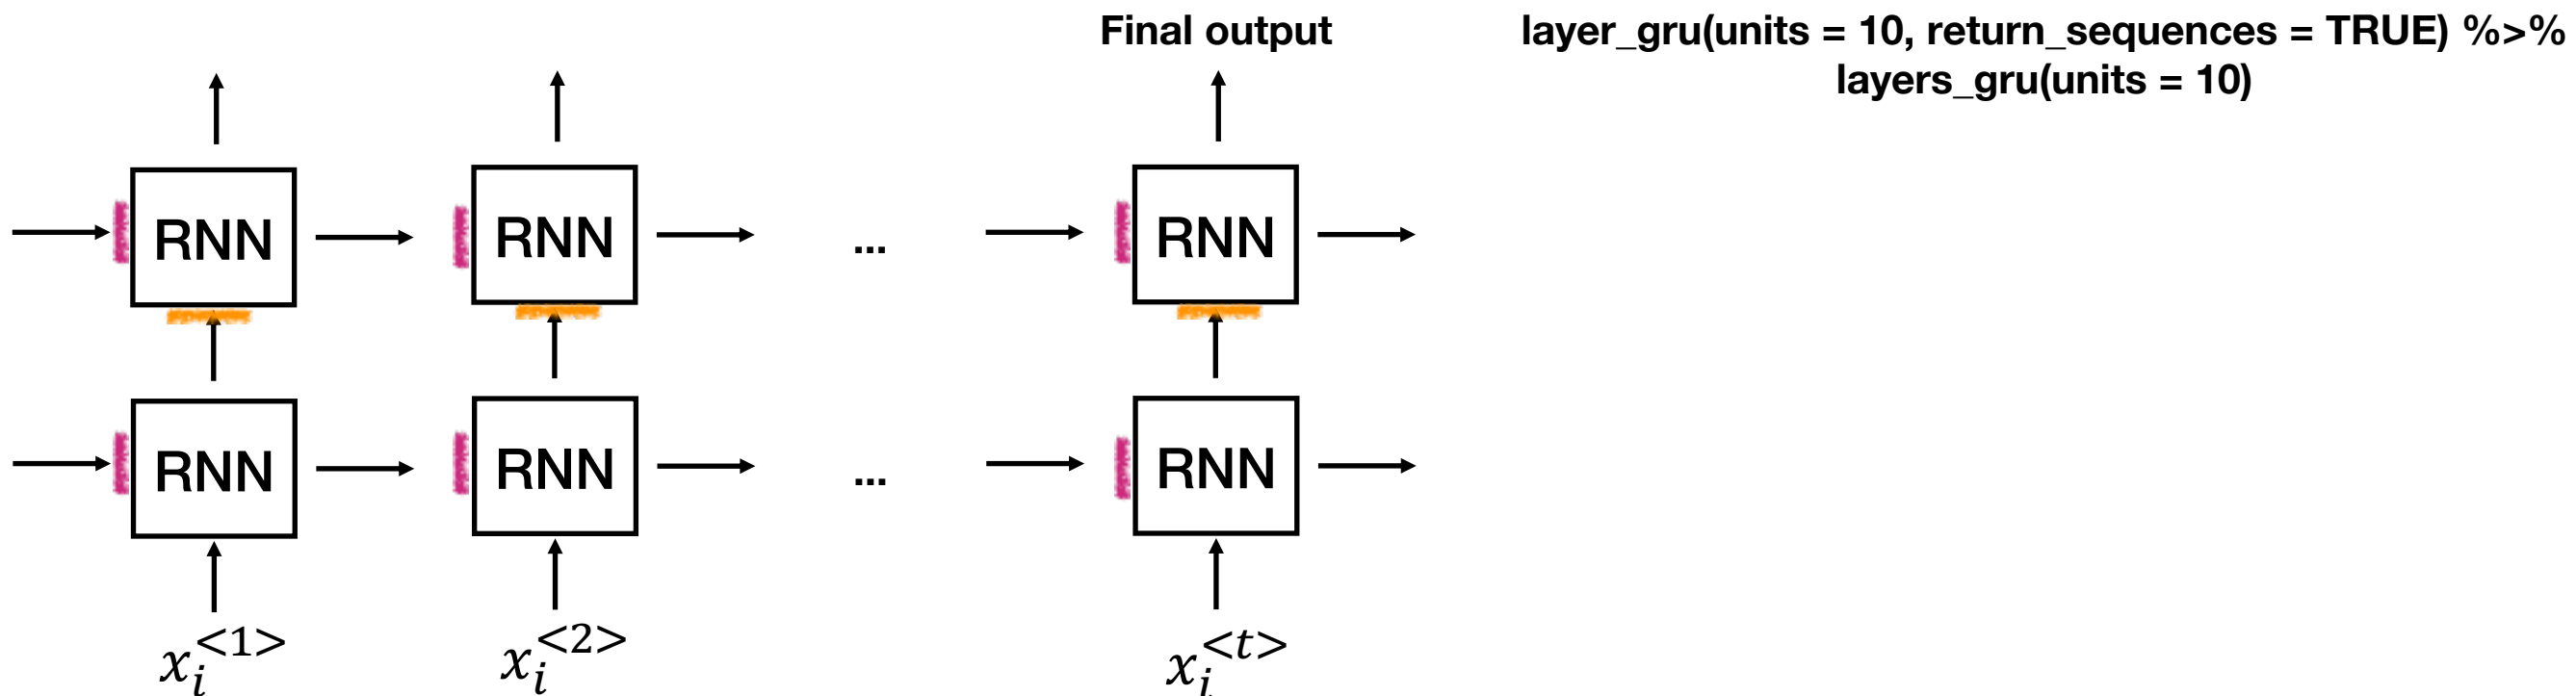
`layer_gru(units = 10, dropout = 0.2, recurrent_dropout = 0.3)`



# Improving RNNs

## Stacking RNNs


- Why would we consider input dropout?
  - Maybe in production we might not always get all inputs.
- More likely, we are stacking RNNs.
- Stacking RNNs is like adding additional layers in a dense network.
  - We never go that deep, 1-6 layers. Long training time.



# Improving RNNs

## Stateful


- A stateful RNN passes the last state of the previous batch to as an initial state to the next batch.
- Otherwise the initial state is "all zeroes".
- This is useful if there is some connection between batches.
- For example, the batches are in sequence.



# Improving RNNs

## Bi-directional

- We process the sequence in both directions.
- Very helpful for example in named entity recognition, in which we classify every word as a "person", "place", ...
  - He said, "Teddy Roosevelt
  - ...
  - He said, "Teddy bears ...



# Summary

- Regularisation
  - L1/L2
  - Dropout, recurrent dropout
- Improving RNNs
  - Stacking
  - Stateful
  - Bi-directional

# Hands-on



Go to <https://jupyter.lisa.surfsara.nl:8000/>

Or <https://dba.projects.sda.surfsara.nl/>

Notebook: `05b-rnns-improved.ipynb`

**Duration: +/- 45 mins**



# Will not cover

- We did not cover any natural language processing (NLP).
- **Word embeddings**, representing words as vectors
- RNNs have been very successful in NLP over the years.
- NLP requires **a lot** of data preprocessing and large models.
- Same models used.
- New Paradigm: Transformers

