

Deep learning


CNNs

Announcements

- Environment
 - I have yet to add "common issues" section for windows.
 - Halfway through updating cluster for more speed.
 - Docker image 1.3
- Assignment 2 set on 9th of April, due 23rd of April.


Recap / Questions?

- Training RNNs is hard
- **Vanishing gradient** problem.



Recap / Questions?

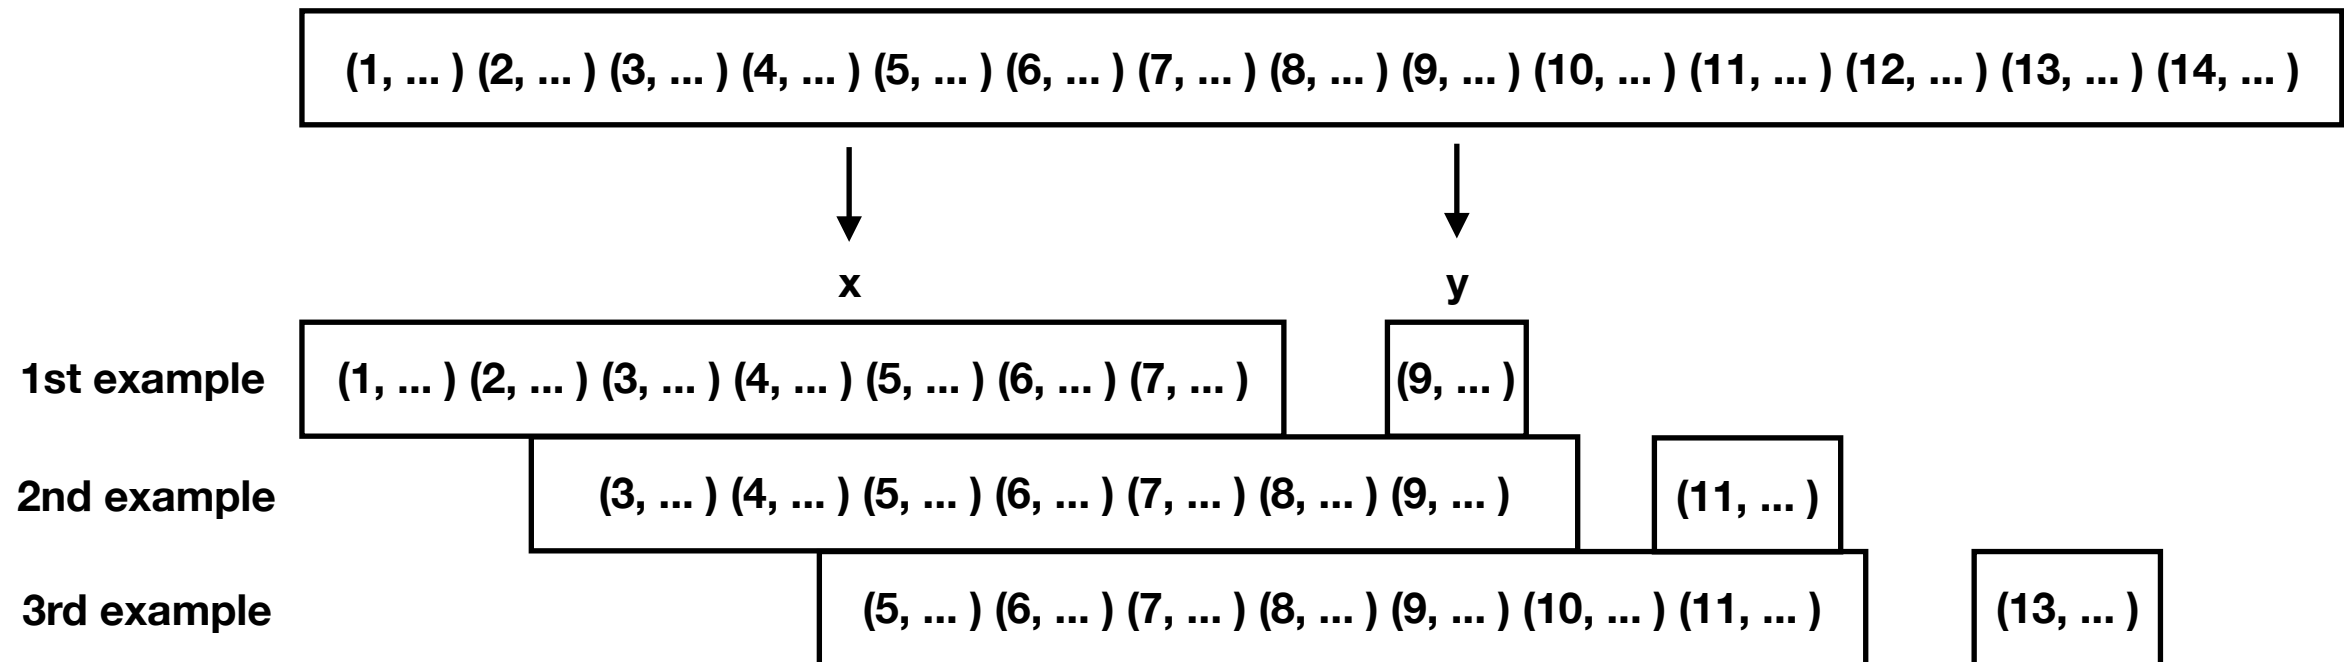
- **GRU** and **LSTM** solve that problem decently.
- Solution was **residual connections**.



Recap / Questions?



- Generating sequences

Shift approach, using shift = 2, sequence length = 7, target shift = 1



Recap / Questions?

- Improving RNNs using L1/L2 regularisation and dropout.
- Stacking RNNs
- Stateful RNNs
- Bi-directional RNNs



Overview

Today we will cover

Topic: CNNs.

- Image data and signal processing.
- Convolutions.
- Convolutional layers.

Notebook: Simple CNN using convolutions on image data.

Topic: CNN architectures

- Max/Avg pool.
- LeNet
- ResNet
- Transfer learning

Notebook: Using a pretrained ResNet50 on image data.

Image data

- Today, a lot about "convolutions" on an image.
- These approaches work well on all data in which the data-points have a "**closeness**" relationship to each other.
 - Sound data
 - Image data
 - Timeseries
- Data can be 1D, 2D, ...
- We talk about picture data.




Image data

- An image = Matrix
- Gray = 256 colors (1 byte)
- White = 255
- Black = 0

Pixel and Digital Number

A photograph could also be represented and displayed in a **digital format** by subdividing the image into small equal-sized and shaped areas, called **picture elements** or **pixels**, and representing the brightness of each area with a numeric value or **digital number**.


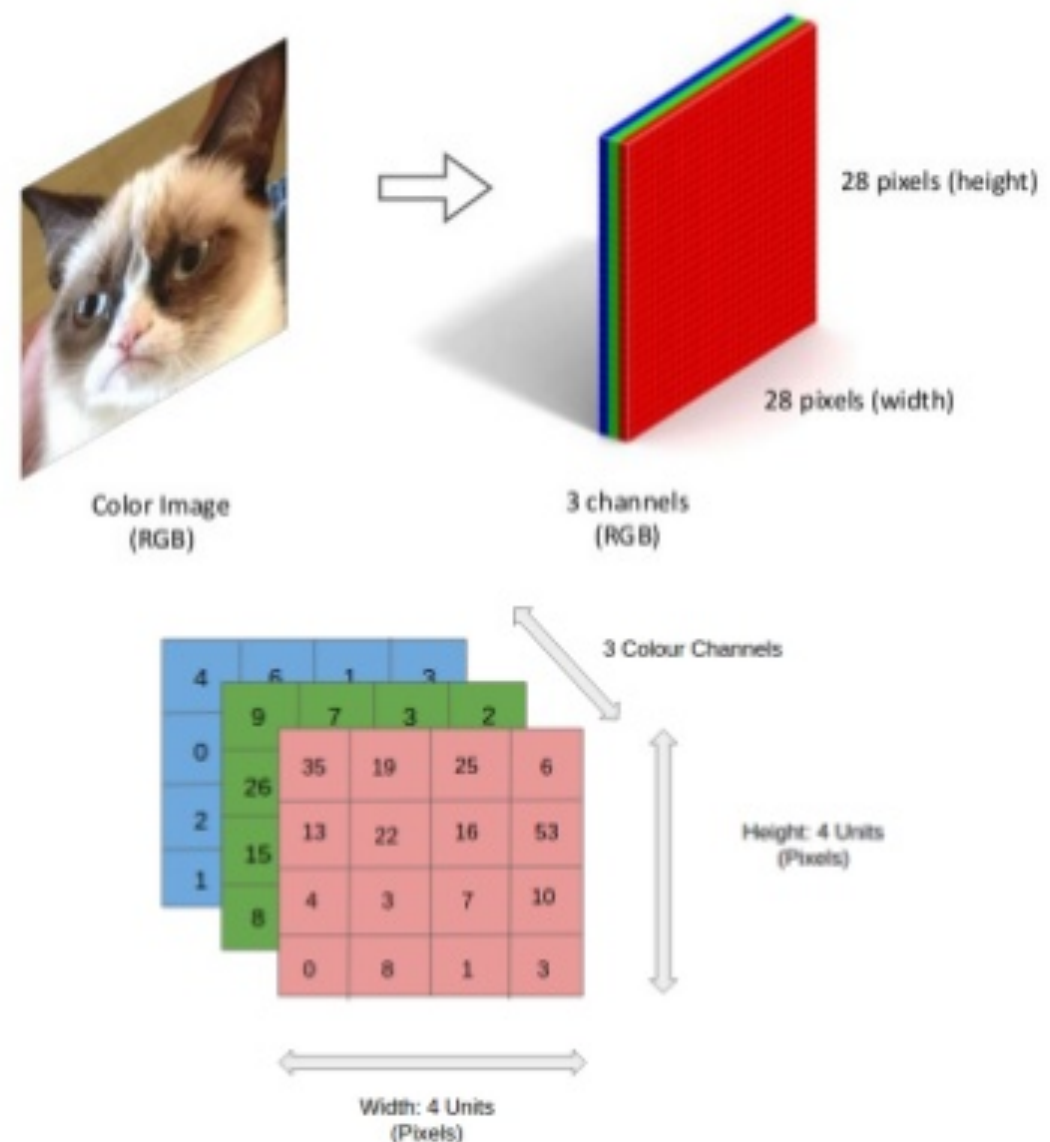


Image data

color image is 3rd-order tensor

- With colors = 3D
- 3 color channels
- Instead of just storing single value between 0-255 we store 3 such values.



Signal processing

- **Signal processing** deals with pattern detection in such data. Very generically.
 - They need to **represent the data** cleverly when transmitting (streaming online, bluetooth, ...) so no/little information is lost.
 - Deep learning is about learning good representations.
- The main tool borrowed from signal processing is the concept of **convolutions**, or **filters** which are applied repeatedly to the signal in attempts to find a specific pattern.

Convolutions



- We can think of convolutions as small filters on the image.
- They process parts of the image and "fire" when they detect the filter pattern.

1x1	1x0	1x1	0	0
0x0	1x1	1x0	1	0
0x1	0x0	1x1	1	1
0	0	1	1	0
0	1	1	0	0

4		


Convolutions

- Applying a filter to an image.




Convolutions

- We then applying a convolution upon a convolution.
- Allows us to create filters from many other filters.
- Then we can detect complex patterns over a larger area.



Convolutions

Technicalities



Convolutions

- We apply a **kernel** with dimensions (3, 3).
- Using **stride** = 2.
- **Padding** = 1, to increase output size
- **Same-padding** = Output is same size as input
- Otherwise **valid-padding**.
- Apply activation function to output.
- We try to **learn the filters!** The filters are the parameters/weights.

0 ₂	0 ₀	0 ₁	0	0	0	0
0 ₁	2 ₀	2 ₀	3	3	3	0
0 ₀	0 ₁	1 ₁	3	0	3	0
0	2	3	0	1	3	0
0	3	3	2	1	2	0
0	3	3	0	2	3	0
0	0	0	0	0	0	0

1	6	5
7	10	9
7	10	8

```
layer_conv_2d(filters = 1,
               kernel_size = c(3, 3),
               stride = 2,
               padding = "valid",
               activation = "",
               kernel_regularizer = "",
               input_shape = c(5, 5))
```


Convolutions





image as a 3D tensor


flatten




each layer is a "channel" represent a color

a 3x3x3 convolution Kernel


"scan"



"scan"



"scan"



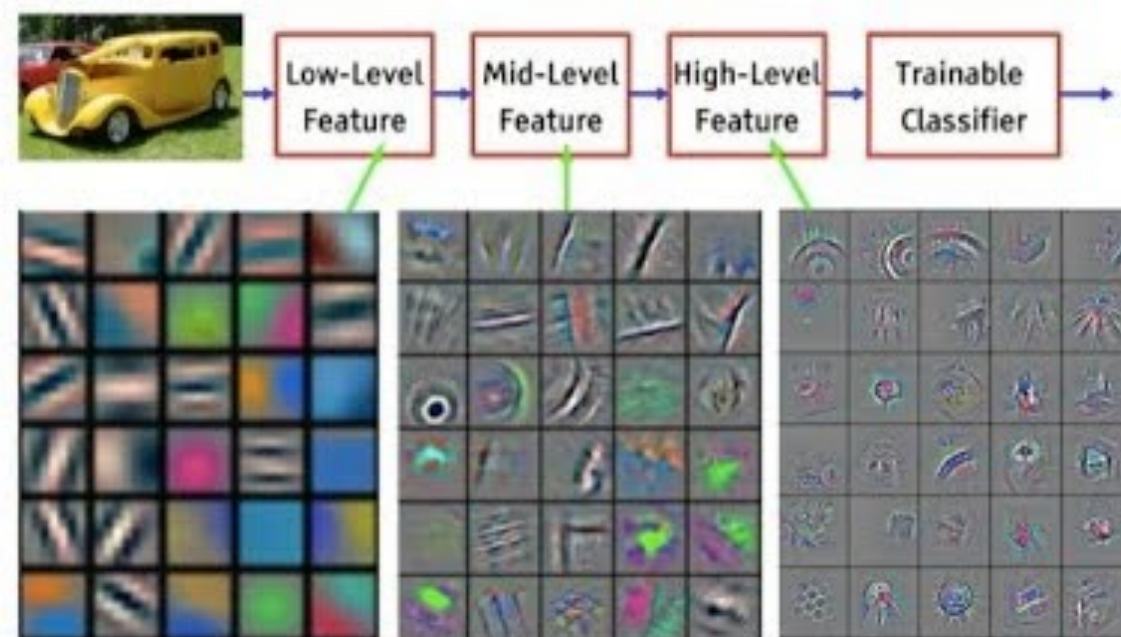
we "scan" the intensity dimension, but not the color dimension!

Convolutions

- We then create many filters/kernels in each layer.

```
layer_conv_2d(filters = 32,  
              kernel_size = c(3, 3),  
              stride = 2,  
              padding = "valid",  
              activation = "",  
              kernel_regularizer = "",  
              input_shape = c(5, 5))
```

Convolutional Neural Network



Feature visualization of convolutional net trained on ImageNet from [Zeiler & Fergus 2013]

Hands-on




Go to <https://dba.projects.sda.surfsara.nl/>

Notebook: `06a-cnns.ipynb`

CNNs

Pooling layers


- Our filter maps are sometimes very large, so we make them **smaller using pooling**.
- **Max pooling**, take the max.
- **Average pooling**, take the average.
- Applied after convolutions



CNNs

Architectures - AlexNet

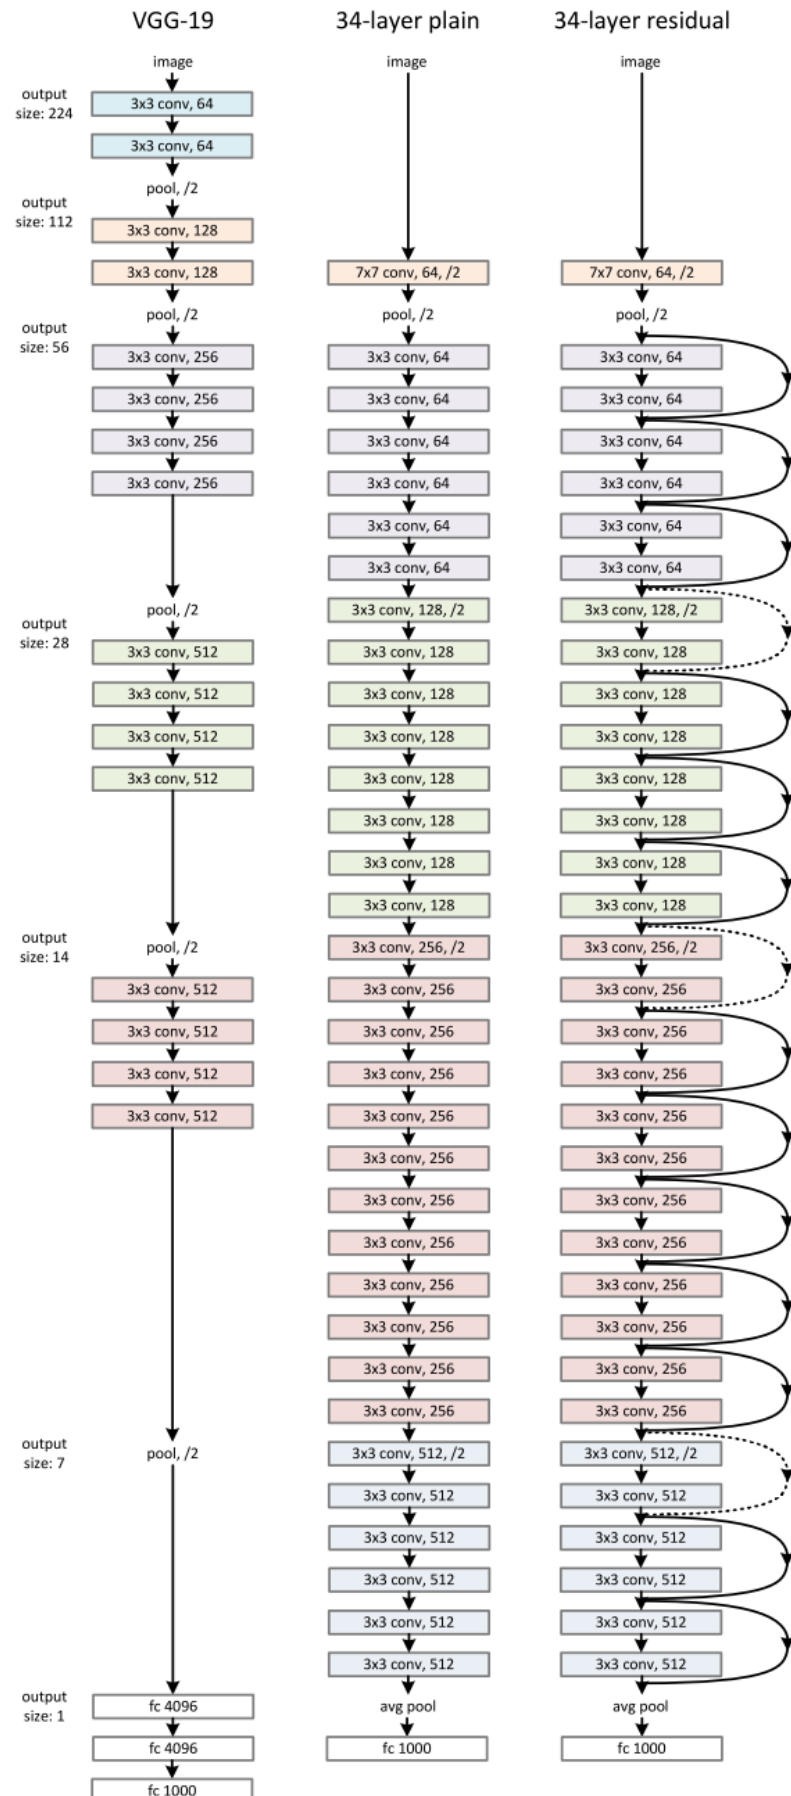
- By chaining different combinations of them we can create Convolutional Neural Networks (CNNs).



Source: <https://www.jeremyjordan.me/convnet-architectures/>

CNNs

Architectures - ResNet (2015)




Source: <https://arxiv.org/abs/1512.03385>

CNNs

Pretrained

- Should we develop these architectures ourselves?
- No, we used them, pre-trained
- This is called **transfer learning**.




Source: <https://www.jeremyjordan.me/convnet-architectures/>

Transfer learning

- We use the fact that someone trained the model on a very large dataset, using **more data and compute** than we can (usually) hope to achieve.
- The last part is always specific to the task at hand, we want the "general" part, so we do not use the whole network.
- We **do not train the network**, just the part we add.
- Our hope is that the network has learnt some generally **good representations** we can use.
- Not training the weights of the network is called **freezing the weights**.


Transfer learning



Transfer learning

Idea: use outputs of one or more layers of a network trained on a different task as generic feature detectors. Train a new shallow model on these features.

Assumes that $D_S = D_T$



Transfer learning

Caveats

- We need to **download** the network and weights, which can be large (still 1000 times faster than training ourselves).
- We use a **smaller learning rate**.
- To increase speed, many people preprocess the input through the static network and **save the representations to disk** and then train a new network separately.

Hands-on



Go to <https://dba.projects.sda.surfsara.nl/>

Notebook: 06b-cnns-transfer.ipynb

Wrap-up at 12:20 / 16:20

Transfer learning

- You also see transfer learning in RNNs, especially with word embeddings.

Summary

Topic: CNNs.

- Image data and signal processing.
- Convolutions.
- Convolutional layers.

Notebook: Simple CNN using convolutions on image data.

Topic: CNN architectures

- Max/Avg pool.
- LeNet
- ResNet
- Transfer learning

Notebook: Using a pretrained ResNet on image data.